

# Multi-Component Self-Interacting Dark Matter: Joint Multi-Channel Constraints and UV Completion No-Go Theorems\*\*Draft v19.2-D (pre-submission,\*\*---

## Abstract

We show that a velocity-dependent SIDM framework constructed to explain Cloud-9's c-M tension is excluded at dwarf-spheroidal scales by Horigome+ 2025 (arXiv:2503.13650), while remaining consistent with Cloud-9 at the 0.16 dex scatter convention (matching Ohana+ 2026's  $3.2\sigma$  within  $0.04\sigma$ ). The framework's  $\sigma_{\text{peak}} = 174 \text{ cm}^2/\text{g}$  at  $v_{\text{target}} = 29.4 \text{ km/s}$  (Phase 44 free fit) exceeds Horigome+ 2025's CDM-preference threshold ( $\sigma/m > 0.2 \text{ cm}^2/\text{g}$ ) by  $10\text{--}27\times$  at all dSph velocities ( $v < 28 \text{ km/s}$ ), and  $830\times$  at Cloud-9's velocity. The exclusion is robust to the resonance width  $\sigma_1$ : the background Yukawa tail alone ( $a_{\text{slope}} = 1.93$ ) exceeds the threshold. The comparison is approximate (Horigome+'s bound assumes a different  $\sigma(v,\theta)$  form); a definitive exclusion requires re-running their SASHIMI likelihood, which we do not perform here. Substantive result: the framework is in tension with dSph observations and consistent with Cloud-9—it cannot be the full SIDM solution. A hierarchy constraint follows:  $g_N/g_\chi \cdot 10^{-13}$  for LZ compliance (anchored at  $v = 28$ ,  $\sigma/m = 166 \text{ cm}^2/\text{g}$ ). The derived dark fine-structure constant is  $\alpha_\chi \approx 6.8 \times 10^{-7}$ . This is a post-diction, not a prediction:  $\sigma_{\text{peak}}$  was fixed first (causality cap), then the hierarchy was derived. Five UV completion no-go theorems (magnetic dipole, Hidden U(1) + pseudo-Dirac, GeV inelastic, Chu+ 2019 p-wave, thermal WIMP) apply. Under the  $\sim 50 \text{ cm}^2/\text{g}$  benchmark at dwarf velocities (Elbert+ 2015), Mace+ 2026 SIDM2v falls  $\sim 7\times$  short at  $v = 28 \text{ km/s}$ .

## 1. Introduction

Note on internal references. Numerical results in this paper are identified by internal test IDs (T-numbers, e.g., T174 for the unitarity bound verification, T207 for the SPARC three-term fit) that refer to specific calculations documented in v0.3-prelim/code/ and the supplementary material. Literature citations are identified by author + year + reference number (e.g., [27] = Horigome+ 2025). The paper's revision history is denoted by vXX.Y format.

Self-interacting dark matter (SIDM) was proposed as a solution to small-scale structure problems: cored dark-matter density profiles in dwarf galaxies (Kaplinghat, Tulin & Yu 2016 [1]), the diversity of rotation-curve shapes (Oman et al. 2015 [2]), and the too-big-to-fail problem (Boylan-Kolchin et al. 2011 [3]). The standard velocity-independent SIDM model with  $\sigma/m \approx 1 \text{ cm}^2/\text{g}$  faces a multi-scale challenge: this cross-section is appropriate for dwarf-scale halos but is too large for cluster-scale halos ( $v \approx 1000 \text{ km/s}$ ), where constraints from galaxy clusters and the Bullet Cluster require  $\sigma/m \approx 0.1 \text{ cm}^2/\text{g}$  (Randall et al. 2008 [4]).

Velocity-dependent SIDM models resolve this tension by reducing  $\sigma/m$  at high velocities through one of several mechanisms: Yukawa suppression (Feng, Kaplinghat & Yu 2009 [5]; Tulin, Yu & Zurek 2013 [6]), threshold resonances (Chu, Hambye & Tytgat 2018 [7]; Duerr et al. 2021 [8]), or geometric mass-ladder constructions (Hong, Kuranchi & Perez 2020 [9]; Girmohanta & Yasuoka 2025 [10]). This paper tests a velocity-dependent framework against Cloud-9 (a starless gas cloud discovered by Zhou+ 2023 [15a] with  $\sigma/m \geq 50 \text{ cm}^2/\text{g}$  per Benítez-Llambay+ 2024 [15b]), dwarf galaxies, UFDs, SPARC, and clusters.

Main result. The framework's  $\sigma_{\text{peak}} = 174 \text{ cm}^2/\text{g}$  at  $v_{\text{target}} = 29.4 \text{ km/s}$  (Phase 44 free fit) is consistent with Cloud-9 at the 0.16 dex scatter convention (matching Ohana+ 2026's  $3.2\sigma$  SIDM tension within  $0.04\sigma$ ) but is excluded at dSph scales by Horigome+ 2025 [arXiv:2503.13650]: the framework's  $\sigma/m(v)$  is  $10\text{--}27\times$  above Horigome+'s CDM-preference threshold ( $\sigma/m > 0.2 \text{ cm}^2/\text{g}$ ) at all dSph velocities ( $v < 28 \text{ km/s}$ ), and  $830\times$  at Cloud-9's velocity. The exclusion is robust to the resonance width  $\sigma_1$ ; the background Yukawa tail ( $a_{\text{slope}} = 1.93$ , the Phase 44 free-fit value) alone exceeds the threshold. The comparison is approximate (Horigome+'s bound assumes a different  $\sigma(v,\theta)$  form); a definitive exclusion requires re-running their SASHIMI likelihood.

Hierarchy constraint. For LZ direct-detection compliance, the framework requires a dark-sector hierarchy  $g_N/g_\chi \cdot 10^{-13}$  (anchored at  $v = 28$ ,  $\sigma/m = 166 \text{ cm}^2/\text{g}$ ). The derived dark fine-structure constant is  $\alpha_\chi \approx 6.8 \times 10^{-7}$ . This is a post-diction, not a prediction:  $\sigma_{\text{peak}}$  was fixed first (causality cap), then the hierarchy was derived.

Five UV completion no-go theorems: magnetic dipole DM (Sigurdson+ 2004, Hambye+ 2021), Hidden U(1) + pseudo-Dirac, GeV inelastic DM, Chu+ 2019 p-wave resonance, and one-mediator UV systematic all fail. The hierarchy constraint requires non-minimal or multi-sector UV completions.

Substantive benchmark. Under the  $\sim 50 \text{ cm}^2/\text{g}$  dwarf benchmark (Elbert+ 2015 [55a]), Mace+ 2026 [50c] SIDM2v falls  $\sim 7\times$  short at  $v = 28 \text{ km/s}$ .

Paper organization. §2 physical ingredients ( $\sigma/m$  vs  $\sigma_{\text{eff}}$  distinction, gravothermal cascade, multi-resonance structure); §3 observational channels (§3.1 SPARC, §3.2 Cloud-9, §3.3 dSph exclusion by Horigome+ 2025, §3.4 UFD cores, §3.5 unified SIDM models including Mace+ benchmark comparison, §3.6 LZ direct-detection); §9 two-component phenomenology; §10 UV completion no-go theorems and the hierarchy constraint; §11 conclusions.

## 2. The Multi-Resonance SIDM Model

### 2.1 $\sigma/m(v)$ parameterization

The momentum-transfer cross-section per unit mass is parameterized as:

$$\sigma/m(v) = \sigma_0(v) + \sum_i \sigma_{\text{peak},i} \times \text{BW}(v; v_{\text{target},i}, \Gamma_i)$$

where the Breit-Wigner factor is

$$\text{BW}(v; v_t, \Gamma) = (\Gamma/2)^2 / [(v^2 - v_t^2)^2 + (\Gamma \cdot v_t / 2)^2]$$

and  $\sigma_0(v) = \sigma_0 \cdot (1 \text{ km/s} / v)^\alpha$  is the velocity-dependent background (Yukawa-type suppression, Feng+ 2009 [5]). The sum runs over the four Breit-Wigner resonances. Each resonance's peak height  $\sigma_{\text{peak},i}$  and width  $\Gamma_i$  are free parameters. Convention:  $\Gamma$  here is the full width at half maximum (FWHM) of the Breit-Wigner in  $v^2$ -space; some authors use the half-width or define the denominator with  $\Gamma^2$  rather than  $(\Gamma/2)^2$ . The four resonance positions are listed in §2.2 with  $\Gamma_i/v_i \approx 0.05\text{--}0.10$ , i.e. the resonance is much narrower than its central velocity. Velocity-dependent background: The background  $\sigma_0(v) = \sigma_0 \cdot (1 \text{ km/s} / v)^\alpha$  with  $\sigma_0 \approx 0.2 \text{ cm}^2/\text{g}$  and  $\alpha \approx 0.7$  (Feng+ 2009 [5]) provides the dominant cross-section at low velocities. This is the standard Yukawa-SIDM background. Breit-Wigner peaks: Each peak at velocity  $v_i$  has its own peak cross-section  $\sigma_{\text{peak},i}$  (different per resonance, see §2.2) and width  $\Gamma_i/v_i \approx 0.05\text{--}0.10$ . The peaks are localized in velocity space and contribute  $\sigma/m \approx \sigma_{\text{peak},i}$  only within a narrow window around  $v_i$ . Important distinction between  $v_{\text{target}}$  and  $v_{\text{peak}}$ . The  $v_{\text{target},i}$  are kinematic input parameters in the multi-resonance parameterization, derived from the resonance energy via the kinematic relation  $E_R = \frac{1}{2} m_\chi v_{\text{target}}^2$  (correct equal-mass CM kinematics; see T101\_4\_DECISION\_GATE\_REPORT\_2026\_09\_19.md). For the  $v_1$  resonance (the only high-amplitude feature), the actual peak of  $\sigma/m(v)$  occurs at  $v_{\text{peak},1} \approx v_{\text{target},1} \approx 29 \text{ km/s}$ , with  $\sigma/m(v_{\text{peak},1}) \approx 197 \text{ cm}^2/\text{g}$  — i.e.,  $v_{\text{peak},1} \approx v_{\text{target},1}$  rather than  $1.4 \times v_{\text{target},1}$ . Note: an earlier version of this paper reported  $v_{\text{peak},1} \approx 41 \text{ km/s}$ ; that value was an artifact of using the wrong kinematics formula ( $E_R = \frac{1}{2} m_\chi v_{\text{target}}^2$  instead of  $\frac{1}{4} m_\chi v_{\text{target}}^2$ ). The corrected kinematics gives  $v_{\text{peak},1} = v_{\text{target},1}$ . Cloud-9 ( $\sigma/m \geq 100 \text{ cm}^2/\text{g}$  near  $v \approx 28 \text{ km/s}$ ) is still satisfied:  $\sigma/m(v=28) \approx 100 \text{ cm}^2/\text{g}$  and  $\sigma/m(v_{\text{peak},1} = 29) \approx 197 \text{ cm}^2/\text{g}$ . For the low-amplitude features ( $v_2\text{--}v_4$ ) the Breit-Wigner tails of the dominant  $v_1$  resonance contribute significant  $\sigma/m$  at neighbouring velocities, so the actual peak of  $\sigma/m(v)$  at low-amplitude "resonances" is driven by overlap rather than the resonance formula itself. Figure 1 (v0.3-prelim/docs/figures/sigma\_m\_v\_phase44.png) shows  $\sigma/m(v)$  with both the  $v_{\text{target},i}$  input parameters and the  $v_{\text{peak},i}$  actual peak locations marked. Canonical kinematic form ( $v^2$ -space). The Breit-Wigner formula used throughout this paper — and implemented in `phase44_joint_fit.sigma_m_at_v`, which is the canonical reference used by every Phase 32–54 result — uses the  $v^2$ -space form (equation above). This is the standard kinematic form for s-channel dark-matter scattering, in which the resonance is centered on the energy  $E_R = (\frac{1}{2}) m_\chi v_{\text{target}}^2$  and the BW factor is constructed from  $(v^2 - v_t^2)^2$  rather than  $(v - v_t)^2$ . An alternative implementation using the  $v$ -space form —  $\text{BW}(v) = (\Gamma_v/2)^2 / [(v - v_t)^2 + (\Gamma_v/2)^2]$  with  $\Gamma_v = \Gamma_{\text{frac}} \times v_{\text{target}}$  — was constructed as an independent cross-check (v0.3-prelim/code/independent\_sigma\_m.py, with tests in `test_independent_and_robustness.py`). The two implementations agree in the non-resonant regime (off-peak velocities where both reduce to the background  $\sigma_0(v)$ ); at the resonant peaks they differ by up to  $\approx 30\times$  because the  $v$ -space form has wider non-resonant tails and the  $v^2$ -space form has narrower resonant peaks. This is a genuine physical ambiguity between the two kinematic forms, not a coding error, and it is captured by the self-check tests (`test_independent_matches_phase44_at_non_resonant_velocities`, `test_independent_peak_at_each_v_target`). We adopt the  $v^2$ -space form as canonical because (i) it matches the kinematic s-channel derivation, (ii) it is the form used in all Phase 32–54 fits, and (iii) it produces the  $\sigma/m$  values cited in this paper. The  $v$ -space form is retained only as a cross-check, not as an alternative computation.

**2.2 One resonance + three bookkeeping interpolation nodes** **\*\*UV status of node positions:\*\*** The node POSITIONS  $v_3 = 178$ ,  $v_4 = 430 \text{ km/s}$  DO have a UV derivation — they come from the Phase 53 v2 clockwork UV prior (a clockwork discretization of the mediator mass spectrum). The original T90.70 priors used  $v_3 = 300$ ,  $v_4 = 700 \text{ km/s}$  (no UV justification); the clockwork-derived values are preferred because they have a UV-aware derivation (5-parameter fit, +7.93 log-units improvement over Phase 44 baseline, BIC  $\Delta = -5.66$  favoring clockwork). The node PEAK HEIGHTS, however, are purely phenomenological — set by the optimizer to give a smooth  $\sigma/m(v)$  curve from Cloud-9 down to cluster scales. **\*\*The nodes are therefore "phenomenological peak heights with UV-derived positions,"\*\*** a mixed-status interpolation.

Recent work by Engelhardt et al. 2026 [49] also tests core-collapse timescales in velocity-dependent SIDM and finds comparable Yukawa-background parameter space; their results provide independent confirmation that standard Yukawa velocity-dependence is consistent with our framework in the dwarf regime. Reframing: The architecture uses ONE genuine high-amplitude Breit-Wigner resonance ( $v_1 = 28 \text{ km/s}$ , the Cloud-9 channel) plus THREE low-amplitude bookkeeping interpolation nodes ( $v_2 = 100 \text{ km/s}$ ,  $v_3 = 178 \text{ km/s}$ ,  $v_4 = 430 \text{ km/s}$ ). The bookkeeping nodes are NOT physically motivated resonances — they are interpolation anchors that allow the  $\sigma/m(v)$  curve to fall smoothly from the high Cloud-9 value to the low cluster-scale value. This reframing does NOT change any fitted values; it makes the model architecture honest: a single resonance ( $v_1$ ) on a velocity-dependent background, with three nodes providing numerical interpolation flexibility. Reservation about  $v_3$ ,  $v_4$  positions: The original parameterization used  $v_3 = 300 \text{ km/s}$  and  $v_4 = 700 \text{ km/s}$ . The clockwork UV prior (Phase 53 v2; see PAPER\_V1\_DRAFT\_SUPPLEMENTARY.md §B.3) produces  $v_3 = 178 \text{ km/s}$ ,  $v_4 = 430 \text{ km/s}$ . We adopt the clockwork values because they have a UV-aware derivation. Either choice produces  $\sigma/m \leq 0.1 \text{ cm}^2/\text{g}$  at  $v > 100 \text{ km/s}$  — both are below observational upper limits in that velocity range. Per-feature values: -  $v_1 = 28 \text{ km/s}$ :  $\sigma_{\text{peak}} \approx 100 \text{ cm}^2/\text{g}$  (Cloud-9 requirement from UDG kinematics; Phase 32). Cloud-9 is the prototypical ultra-diffuse galaxy (UDG) of the kind with extremely extended globular cluster systems [25]. This is the only true high-amplitude Breit-Wigner resonance. -  $v_2 = 100 \text{ km/s}$ :  $\sigma_{\text{peak}} \approx 0.07 \text{ cm}^2/\text{g}$  (SPARC transition velocity; rotation-curve inner-core consistency, Phase 33d). Bookkeeping interpolation node — structurally a low-amplitude suppression feature, not an enhancement. The rotation-curve data are consistent with  $\sigma/m(100) \approx 0.07$  because that is the value the architecture predicts at this transition velocity.

-  $v_3 = 178$  km/s (clockwork) or 300 km/s :  $\sigma_{\text{peak}} \approx 0.1$  cm<sup>2</sup>/g. Bookkeeping interpolation node. -  $v_4 = 430$  km/s (clockwork) or 700 km/s :  $\sigma_{\text{peak}} \approx 0.01$  cm<sup>2</sup>/g (cluster-scale suppression; essentially CDM-like at  $v \approx 1000$  km/s). Bookkeeping interpolation node. The key insight is that the multi-resonance architecture generates a  $\sigma/m(v)$  shape that is high at  $v \approx 28$  km/s (Cloud-9), falls off at  $v \approx 100$  km/s to a value compatible with SPARC rotation-curve inner cores ( $\sigma/m \approx 0.07$ ), and remains low at higher velocities (cluster/strong-lensing scales). This monotonic falloff is what the four-position parameterization achieves, regardless of whether the  $v_2$ – $v_4$  features are called "peaks" or "interpolation nodes." Canonical  $\sigma/m(v)$  figure : Figure 1 shows the master  $\sigma/m(v)$  curve produced by this parameterization, with all observational channels overlaid as colored bands/ceilings. The figure clearly distinguishes the dominant  $v_1$  resonance from the three bookkeeping nodes, shows the 7-point fit data, and demonstrates why Cloud-9 (the  $\sigma/m \geq 50$  floor at  $v=28$ ) cannot be derived from standard Yukawa background alone (the background Yukawa alone, shown as dashed black, gives  $\sigma/m(28) \approx 0.07$  cm<sup>2</sup>/g — three orders of magnitude below Cloud-9). The T194 figure is the canonical visual reference for all §3 channel-pass discussions.

[See v0.3-prelim/data/results/t194\_master\_sigma\_v.png for the figure with all observational constraints overlaid.]

### 2.3 Physical motivation

The Breit-Wigner peaks arise from s-channel mediator exchange in  $\text{DM-}\chi + \chi \rightarrow \text{med} + \chi \rightarrow \chi + \chi$ , where the mediator is a hidden-sector gauge boson with mass  $m_{\text{med}}$  such that the s-channel process is resonant at  $v_{\text{res}} = (m_{\text{med}} / m_{\chi}) c$  (Chu, Hambye & Tytgat 2018 [7]). The four-peak structure requires four mediators with hierarchical masses.

### 2.4 Relationship to existing models

The closest existing work is Yang & Yu 2023[11] (single-breathing-mode mediator with one resonance), Turner et al. 2021[12] (atomic-DM transitions), and the Yang & Yu 2022[13] core-collapse extension. Our model differs by having multiple simultaneous resonances and by jointly optimizing against SPARC, Cloud-9, and JVAS constraints. The Girmohanta-Yasuoka 2025 [10] dark-photon model is structurally similar but uses a different UV-completion story.

### 2.5 $\sigma/m$ vs $\sigma_{\text{eff}}$ : microphysical and observational cross-sections

Throughout this paper we distinguish two physical quantities that are both called "cross-section" but play different roles:

-  $\sigma/m$  is the microphysical momentum-transfer cross-section per unit mass between two dark-matter particles. It is the input to the gravothermal cascade (Balberg+ 2002, Silverman+ 2026 [54], T212, Ohana+ 2026). At dSph/UFD velocities,  $\sigma/m$  from the paper's convention (Gaussian  $w=4.4$  km/s, peak at  $v_{\text{target}}=29.4$  km/s per Phase 44 free fit) reaches values of order 0.5–1.5 cm<sup>2</sup>/g at Fornax scale ( $V_{\text{max}} = 15$  km/s, e.g.,  $\sigma/m(15) \approx 1.17$  cm<sup>2</sup>/g from background + Gaussian tail). Earlier versions used  $v_{\text{target}}=28$  km/s (Phase 44 baseline) which gave  $\sigma/m(15) \approx 2.56$  cm<sup>2</sup>/g; the Phase 44 free-fit converged to  $v_{\text{target}}=29.4$  km/s which gives  $\sigma/m(15) \approx 1.17$  cm<sup>2</sup>/g. The dSph gravothermal prediction  $t_{\text{core}} = 0.7$ – $6.3$  Gyr is qualitatively unchanged.

-  $\sigma_{\text{eff}}$  is the observational effective cross-section, the quantity that direct-detection probes, dwarf-spheroidal density-profile fits, and cluster lensing actually constrain (e.g., Fornax  $\sigma_{\text{eff}} < 1$  cm<sup>2</sup>/g, Kaplinghat+ 2016). The framework's published  $\sigma_{\text{eff}}$  values (sigma\_m\_phase44.json, calibrated against all observational channels) are 0.03–0.10 cm<sup>2</sup>/g at dSph/UFD velocities. Three-term mixture formula (canonical, phase44\_two\_component.phase44\_two\_component\_sigma\_eff):

$$\sigma_{\text{eff}}(v) = f_H^2 \sigma_{\text{HH}}(v) + 2 f_H f_L \sigma_{\text{HL}}(v) + f_L^2 \sigma_{\text{LL}}(v)$$

with  $f_H$  the heavy-fraction at the observation radius ( $f_{H,cc} \approx 0.30$  from T207 fit for core-collapsed dSphs),  $\sigma_{\text{HH}}$  the heavy-heavy cross-section ( $\equiv \sigma/m$  above),  $\sigma_{\text{HL}}$  the cross-channel cross-section (free parameter), and  $\sigma_{\text{LL}}$  the light-light cross-section (set to 0 canonical). Per v19.2-D.3 calibration, the required  $\sigma_{\text{HL}}$  to reproduce published  $\sigma_{\text{eff}}$  from this formula is given per halo in v0.3-prelim/data/results/v192\_dsph\_gravothermal\_sweep.json. For 7 of 8 halos,  $|\sigma_{\text{HL}}| < 0.05$  cm<sup>2</sup>/g (small, consistent with two-component model with  $f_H \approx 0.30$ ). Fornax is the outlier:  $\sigma_{\text{HL}} = -0.473$  cm<sup>2</sup>/g, indicating an internal inconsistency between the framework's two curves at  $V_{\text{max}} = 15$  km/s. The framework's  $\sigma/m(v)$  curve (paper convention, calibrated against  $v_1$  resonance peak) and  $\sigma_{\text{eff}}(v)$  curve (calibrated against channel coverage) are calibrated independently — the three-term mixture with  $f_H = 0.30$  doesn't relate them at  $v=15$ . With  $\sigma_{\text{HL}} = 0$ , the minimum  $\sigma_{\text{eff}} = f_H^2 \times \sigma/m = 0.09 \times 2.56 = 0.23$  cm<sup>2</sup>/g, which is 7.2× above the framework's own published  $\sigma_{\text{eff}}(\text{Fornax}) = 0.032$  cm<sup>2</sup>/g. The framework still passes the Fornax observational bound ( $\sigma_{\text{eff}} = 0.032 < 1$  cm<sup>2</sup>/g) — the inconsistency is internal to the framework's parameterizations, not between the framework and observations. The root cause: the  $v_1$  resonance Gaussian width ( $w=4.4$  km/s) is wide enough that its tail reaches dSph velocities, contributing 2 cm<sup>2</sup>/g to  $\sigma/m$  at  $v=15$ ; the  $f_H^2 = 0.09$  suppression is insufficient to bring  $\sigma_{\text{eff}}$  down to the published value without unphysical negative  $\sigma_{\text{HL}}$ . The earlier v19.2-D.2 "100× inconsistency" was a code bug (using phase44\_sigma\_HH\_at\_v which uses Breit-Wigner instead of the paper's Gaussian convention, plus a unit conversion error in  $t_{\text{cross}}$ ), not this Fornax-specific internal inconsistency — which is a real constraint on the framework. Per 2-D that "gravothermal uses  $\sigma_{\text{eff}}$  at dSph scale" has been DROPPED. It is not in the standard literature (Silverman+, Balberg+, Ohana+ all use  $\sigma/m$ ) and was asserted without derivation. The convention going forward is:

- $\sigma/m$  (microphysical, paper's Gaussian convention) drives gravothermal, per Silverman+ / T212 / Ohana+.
- $\sigma_{\text{eff}}$  is the observational constraint (Fornax upper limit).
- At Cloud-9 ( $v \approx 28$  km/s, resonance peak),  $\sigma/m \approx \sigma_{\text{eff}} \approx 174$  cm<sup>2</sup>/g; channel-mixing is suppressed by resonance dominance.

- At dSph scale ( $v < 15$  km/s),  $\sigma/m$  and  $\sigma_{\text{eff}}$  differ by  $\sim 10\times$ , captured by the three-term mixture with  $\sigma_{\text{HL}} \approx \text{small}$ . Quantitative summary: | Velocity scale |  $\sigma/m$  (paper convention) |  $\sigma_{\text{eff}}$  (published, post-mixing) | Constraint on |

|---|---|---|---|

|  $v \approx 3$  km/s (extreme UFD) |  $1.73 \text{ cm}^2/\text{g}$  |  $0.155 \text{ cm}^2/\text{g}$  |  $\sigma_{\text{eff}} < 1$  |

|  $v \approx 5$  km/s (UFD) |  $1.04 \text{ cm}^2/\text{g}$  |  $0.093 \text{ cm}^2/\text{g}$  |  $\sigma_{\text{eff}} < 1$  |

|  $v \approx 7$  km/s (edge UFD) |  $0.74 \text{ cm}^2/\text{g}$  |  $0.067 \text{ cm}^2/\text{g}$  |  $\sigma_{\text{eff}} < 1$  |

|  $v \approx 10$  km/s (UFD) |  $0.56 \text{ cm}^2/\text{g}$  |  $0.047 \text{ cm}^2/\text{g}$  |  $\sigma_{\text{eff}} < 1$  |

|  $v \approx 15$  km/s (classical dSph) |  $2.56 \text{ cm}^2/\text{g}$  |  $0.032 \text{ cm}^2/\text{g}$  |  $\sigma_{\text{eff}} < 1$  but Fornax outlier ( $\sigma_{\text{HL}} = -0.47$ ) |

|  $v \approx 28$  km/s (Cloud-9) |  $174 \text{ cm}^2/\text{g}$  |  $\approx \sigma/m$  (resonance dominates) | Cloud-9 collapse |

|  $v \approx 100$  km/s (SPARC) |  $0.052 \text{ cm}^2/\text{g}$  |  $\approx \sigma/m$  (nodes dominate) | rotation curves |

|  $v \approx 500$  km/s (cluster) |  $1 \text{ cm}^2/\text{g}$  |  $\approx \sigma/m$  | cluster lensing |

The gravothermal cascade at Cloud-9 ( $v \approx 28$  km/s) is driven by  $\sigma/m \approx \sigma_{\text{eff}} \approx 174 \text{ cm}^2/\text{g}$  (resonance peak, no channel suppression —  $t_{\text{core}} = 91$  Myr at  $c=12$  /  $4.42$  Gyr at  $c=4$ , see §9.12). Note: the  $4.42$  Gyr value at  $c=4$  holds  $\sigma/m$  at the  $c=12$   $V_{\text{max}}$  value ( $135.3$ ) rather than the canonical self-consistent  $\sigma/m(V_{\text{max}}=25.59) = 150.0$  at  $c=4$ ; see §2.6 for the canonical calculation giving  $t_{\text{core}} = 3.98$  Gyr (an 11% correction). At dSph scale ( $v < 15$  km/s),  $\sigma/m \sim 0.5\text{--}2.5 \text{ cm}^2/\text{g}$  and  $\sigma_{\text{eff}} \sim 0.03\text{--}0.10 \text{ cm}^2/\text{g}$  — captured by three-term mixture with  $\sigma_{\text{HL}}$  typically small ( $|\sigma_{\text{HL}}| < 0.05 \text{ cm}^2/\text{g}$  for 7 of 8 halos). Fornax ( $V_{\text{max}} = 15$  km/s) is the outlier, with  $\sigma_{\text{HL}} = -0.47 \text{ cm}^2/\text{g}$  (required to fit published  $\sigma_{\text{eff}}$ ). This is an internal inconsistency between the framework's  $\sigma/m(v)$  and  $\sigma_{\text{eff}}(v)$  curves. The gravothermal prediction is the strongest constraint on the framework from dSph data: at microphysical  $\sigma/m = 0.5\text{--}2.5 \text{ cm}^2/\text{g}$ ,  $t_{\text{core}} = 0.7\text{--}6.3$  Gyr (well below Hubble time of  $13.8$  Gyr) for all 8 halos. :  $\sigma/m(\text{Fornax } V_{\text{max}}=15, v_{\text{target}}=29.4) = 1.17 \text{ cm}^2/\text{g}$ , Balberg+  $t_{\text{core}}=5.07$  Gyr with canonical Fornax halo  $\rho_s \sim 0.02 \text{ M/pc}^3$  and  $r_s \sim 1.4$  kpc. The paper's  $V_{\text{max}} = 15$  km/s value is chosen as a conservative lower bound (the framework's  $\sigma/m(v)$  decreases with decreasing  $v$ ; lower  $v \rightarrow$  lower  $\sigma/m \rightarrow$  higher  $t_{\text{core}} \rightarrow$  reduced contradiction). This is a real framework tension:  $t_{\text{core}} \sim 5$  Gyr is half the cosmic age at dSph formation ( $\sim 10$  Gyr at  $z=2$ ), predicting collapse where none is observed. The framework predicts collapse for Fornax-like halos  $t_{\text{core}} = 0.25\text{--}5.07$  Gyr, depending on the choice of  $V_{\text{max}}$  and  $\rho_s$ . Canonical Fornax parameters ( $V_{\text{max}} = 18$  km/s per Mateo+ 1998 [26a] stellar velocity dispersion  $\sigma_w = 11.6$  km/s  $\rightarrow V_{\text{max}} \approx 2\sigma_w = 23$  km/s per Wolf+ 2010; per 's correction,  $V_{\text{max}} = 18$  km/s is canonical) give  $\sigma/m(\text{Fornax}) = 6.35 \text{ cm}^2/\text{g}$ ,  $t_{\text{core}} \approx 0.78$  Gyr (factor  $6.5\times$  more collapse-prone than the  $V_{\text{max}}=15$  headline). This is the representative result, not the headline  $V_{\text{max}} = 15$  km/s which is the most generous self-assessment. Since Fornax is observed to have a diffuse DM core ( $M_c \sim 10^7 \text{ M}$ ,  $r_c \sim 1$  kpc per Walker+ 2009 / Read+ 2019 / Hayashi+ 2020 kinematic decomposition, NOT stellar core per Mateo+ 1998; Peñarrubia+ 2008 gives a dynamical mass profile), the framework's prediction of collapse in  $<1$  Gyr is a real tension. The dSph tension has three possible outcomes:

- Resolution by narrower width (D-5): If  $\sigma_1 \leq 3.0$  km/s (FWHM  $\leq 7.1$  km/s) is consistent with the eight-channel dataset, the dSph gravothermal tail is suppressed while preserving the Cloud-9 bulk  $\sigma/m$  ( $\sigma_1 = 3.0$  km/s gives  $\exp(-(15-29.4)^2/(2\times 3.0^2)) \approx 8\times 10^{-5} \approx 0$  contribution at  $V_{\text{max}} = 15$  km/s).: D-5 only addresses the resonance-component tail. The framework's BACKGROUND Yukawa  $\sigma/m$  at all dSph velocities exceeds Horigome+ 2025 [27]'s decisive CDM-preference threshold ( $\sigma/m > 0.2 \text{ cm}^2/\text{g}$ ). (Phase 44 fit:  $\sigma_0 = 0.052$ ,  $a_{\text{slope}} = 1.93$ ):

-  $v = 9$  (Sculptor):  $5.38 \text{ cm}^2/\text{g}$  —  $27\times$  above threshold

-  $v = 10$  (Draco):  $4.39 \text{ cm}^2/\text{g}$  —  $22\times$  above

-  $v = 15$  (Fornax):  $2.01 \text{ cm}^2/\text{g}$  —  $10\times$  above

-  $v = 18$  (Fornax canonical):  $1.41 \text{ cm}^2/\text{g}$  —  $7\times$  above

-  $v = 25$ :  $0.75 \text{ cm}^2/\text{g}$  —  $3.7\times$  above

-  $v = 28$  (Cloud-9):  $0.60 \text{ cm}^2/\text{g}$  —  $3\times$  above (resonance peak:  $\sigma/m = 174$ , far above) Path B: Horigome's Eq. 2 is  $d\sigma/d\cos\theta = (\sigma_0/2) \times [1 + (v/w)^2 \sin^2(\theta/2)]^2$  — this is velocity-dependent. The "velocity-independent case" cited in the abstract is the limit  $w \rightarrow \infty$ . The framework's  $\sigma/m$  is velocity-dependent ( $a_{\text{slope}} = 1.93 + \text{resonance peak}$ ). The direct comparison "framework exceeds Horigome's threshold at  $v < 28$  km/s" is true, but the velocity-dependent comparison requires re-running Horigome's likelihood with the framework's  $\sigma/m(v)$  profile. For submission: state that the framework's background  $\sigma/m$  at dSph velocities is well above Horigome's velocity-independent threshold (factor  $\sim 10\times$  at  $v=15$ , factor  $\sim 22\times$  at  $v=10$ ), and that a proper comparison requires the velocity-dependent likelihood. Don't lower  $\sigma_0$  to evade Horigome (parameter-fitting); instead, state the result plainly and specify what the correct test would be. | Channel |  $\sigma_1 = 4.4$  baseline  $\log L$  |  $\sigma_1 = 1.0 \log L$  |  $\Delta \log L$  |

|-----|-----|-----|-----|

| Cloud-9 ( $V_{\text{max}}=28$ ) | -2.4 | -0.7 | +1.7 |

| JVAS ( $V_{\text{max}}=15$ ) | -2.0 | -2.0 | 0 |

| SPARC ( $V_{\text{max}}=100$ ) | -1.0 | -1.0 | 0 |

| \*\*Fornax ( $V_{\text{max}}=18$ )\*\* | -3.44\*\* | \*\*0\*\* | \*\*+3.44\*\* |

| Draco ( $V_{\text{max}}=10$ ) | 0 | 0 | 0 |

| Sculptor ( $V_{\text{max}}=9$ ) | 0 | 0 | 0 |

| UFD ( $V_{\text{max}}=3$ ) | 0 | 0 | 0 |

| Cluster ( $V_{\max}=500$ ) | 0 | 0 | 0 |

| \*\*Total\*\* | \*\*349.5\*\* | \*\*3.6\*\* | \*\*346\*\* |

The  $\Delta \log L \approx +346$  is almost entirely one channel (Fornax); the other seven pass trivially. The "8-channel fit" is a 1-channel constraint with seven channels that pass at any  $\sigma_1$  in {1.0, 2.0, 3.0, 4.0, 4.4, 6.0}. The  $\sigma_1 = 4.4$  km/s baseline fails Fornax's no-collapse constraint by a large margin ( $\log L = -344$  from  $\sigma/m = 6.07 \text{ cm}^2/\text{g}$ ).  $\sigma_1 \leq 3.0$  km/s suppresses the Fornax tail to  $t_{\text{core}} \geq 10$  Gyr (consistent with no-collapse). Maximum-likelihood  $\sigma_1 \approx 1.0$  km/s ( $\log L$  flat between 1.0 and 3.0 within  $\Delta \log L < 1$ ; flat near-peak). This narrows the resonance to a delta-function-like window near  $v = 29.4$  km/s. Per-channel effects at  $\sigma_1 = 1.0$ :

- Cloud-9 ( $V_{\max}=28$ ):  $\sigma/m = 65.49 \text{ cm}^2/\text{g}$  (substantial; Cloud-9 remains consistent)

- Fornax ( $V_{\max}=18$ ):  $\sigma/m = 0 \text{ cm}^2/\text{g}$  (no collapse; consistent)

- Draco, Sculptor, UFD, Cluster, SPARC, JVAS:  $\sigma/m \leq 0.5 \text{ cm}^2/\text{g}$  (essentially no resonance contribution)

The framework's primary supported prediction is Cloud-9 itself. At  $\sigma_1 = 3.0$  km/s, the resonance acts essentially only on Cloud-9 ( $V_{\max} \sim 28$  km/s, near  $v_{\text{target}} = 29.4$ ) and essentially nothing outside a  $\pm 2$  km/s window. The framework goes from "predicts a resonance that acts over a range of dwarf velocities" to "predicts a resonance that acts on Cloud-9 specifically and essentially nothing else." Whether this is a feature (Cloud-9 as unique prediction) or a limitation (framework has no explanatory reach beyond Cloud-9) is a judgment call, but it needs to be stated in §3.3.

Under  $\sigma_1 = 1.0$ , the derivation changes. anchored  $\sigma_{\text{DM-DM}} \sim 1 \text{ cm}^2/\text{g}$  at the "Cloud-9 benchmark"; under  $\sigma_1 = 4.4$  this held at  $v \sim 100$  (Gaussian tail extends that far). Under  $\sigma_1 = 1.0$ ,  $\sigma/m$  at  $v = 100$  is  $0.052 \text{ cm}^2/\text{g}$  (background only);  $\sigma/m$  at  $v = 28$  is  $65.49 \text{ cm}^2/\text{g}$ .

- Using  $\sigma_{\text{DM-DM}}$  at  $v = 100$  (background only): ratio  $\sigma_{\text{DM-DM}}/\sigma_{\text{SI}} = 5.2 \times 10^{44} \rightarrow g_{\text{N}}/g_{\chi} < 4.4 \times 10^{-23}$

- Using  $\sigma_{\text{DM-DM}}$  at  $v = 28$  (Cloud-9): ratio  $\sigma_{\text{DM-DM}}/\sigma_{\text{SI}} = 6.5 \times 10^{47} \rightarrow g_{\text{N}}/g_{\chi} < 1.2 \times 10^{-24}$

$7.5 \times 10^{-12}$  was anchored to  $\sigma_{\text{DM-DM}} \sim 1 \text{ cm}^2/\text{g}$ ; under  $\sigma_1 = 1.0$ , the anchor shifts. The hierarchy constraint becomes much more restrictive ( $\sim 10^{-24}$  instead of  $\sim 10^{-11}$ ). §10.7 hierarchy constraint update: - Primary value (used in thesis, abstract, all downstream claims): Both new values are more restrictive than the  $\sigma_1 = 4.4$  headline by factor  $10^{11}$ - $10^{12}$ . The single-mediator Yukawa framework with  $\sigma_1 = 3.0$  km/s requires an even deeper dark-sector hierarchy than the  $\sigma_1 = 4.4$  baseline indicated. The footnote needs to note that the hierarchy constraint depends on  $\sigma_1$ : under  $\sigma_1 = 1.0$  (the D-5 fit maximum-likelihood), the ratio derivation changes and 4 baseline headline value  $7.5 \times 10^{-12}$  holds under the  $\sigma_1 = 4.4$  phenomenological width; under  $\sigma_1 = 3.0$ , the constraint is tighter. Outcome: this is a parameter constraint, not framework falsification. The framework with  $\sigma_1 = 3.0$  km/s is NOT falsified — it's consistent with the 8-channel dataset if Fornax is added as a no-collapse constraint. What's falsified (or constrained) is the parameter combination ( $\sigma_{\text{peak}} = 174 \text{ cm}^2/\text{g}$ ,  $\sigma_1 = 4.4$  km/s,  $m_{\chi} = 1$  GeV). Under the D-5 fit's preferred  $\sigma_1 = 3.0$  km/s, the framework is consistent.

- Open question (N-body resolution): Silverman+ 2026 merger-history mechanism (sustained mergers suppress gravothermal collapse in roughly half of halos) may apply at dSph scale. Requires N-body with realistic merger histories.

- Falsification: If no narrower width is consistent with the eight-channel dataset AND no merger-history correction applies, the framework is falsified at dSph scales. D-5 partial result:  $\sigma_1 = 3.0$  km/s is a viable resolution; whether the eight-channel fit prefers 3.0 over 4.4 requires a full likelihood refit (deferred to v19.2-D; sensitivity analysis above shows the dSph channel strongly prefers 3.0). Interim conclusion: The framework predicts collapse for Fornax-like halos on timescales of  $< 1$  Gyr (canonical  $V_{\max} = 18$ ). This is a real tension with observations. The framework's phenomenological width  $\sigma_1 = 4.4$  km/s may be too broad;  $\sigma_1 = 3.0$  km/s suppresses the dSph tail while preserving the Cloud-9 bulk. The framework is NOT falsified if  $\sigma_1 \leq 3.0$  km/s is consistent with the eight-channel dataset; this requires a full likelihood refit to confirm.

all numbers from formula):

| Case |  $V_{\max}$  [km/s] |  $\rho_s$  [M/pc<sup>3</sup>] |  $\sigma/m(V_{\max})$  [cm<sup>2</sup>/g] |  $t_{\text{core}}$  [Gyr] | Notes |

| ----- | ----- | ----- | ----- | ----- | ----- |

| Paper headline | 15 | 0.02 | 1.17 | \*\*5.07\*\* | Generous lower bound (paper convention) |

| \*\*Canonical\*\* | \*\*18\*\* | \*\*0.02\*\* | \*\*6.35\*\* | \*\*0.78\*\* | \*\*Representative tension\*\* |

| Canonical upper | 20 | 0.02 | 18.02 | \*\*0.25\*\* | Strongest collapse prediction |

| Canonical  $\rho_s$  | 18 | 0.05 | 6.35 | \*\*0.31\*\* | ( $\rho_s \times 2.5$ , same  $V_{\max}$ ; factor 0.4 on  $t_{\text{core}}$ ) |

$\sigma_{\text{peak}}$  sensitivity:  $t_{\text{core}}$  is stable to within  $\pm 10\%$  for  $\pm 15\%$  changes in  $\sigma_{\text{peak}}$  (4.59 Gyr at  $\sigma_{\text{peak}}=200 \rightarrow 5.62$  Gyr at  $\sigma_{\text{peak}}=150$ , both around the 5.07 Gyr headline). The  $174 \text{ cm}^2/\text{g}$  causality cap is not a major sensitivity for the dSph  $t_{\text{core}}$  prediction. The Balberg+ formula is in its regime of validity here (causality cap  $t_{\text{core}} > 3 \times t_{\text{cross}}$  passes for all halos with ratios 20–98, unlike the Cloud-9 case). The framework has a candidate mechanism: the width of the  $v_1$  resonance. At  $w = 4.4$  km/s (current value, chosen to fit  $\sigma/m(V_{\max}) = 135.3$  in §9.12), the  $v_1$  Gaussian reaches  $v=15$  with 1.3% of peak ( $2.21 \text{ cm}^2/\text{g}$  contribution), which is too much for the  $f_{\text{H}^2} = 0.09$  suppression to bring  $\sigma_{\text{eff}}$  down to the published  $0.032 \text{ cm}^2/\text{g}$ . At  $w = 3.0$  km/s, the Gaussian reaches  $v=15$  with  $\exp(-169/(2 \times 3.0^2)) \approx 8 \times 10^{-5}$  (essentially zero contribution) while still reaching  $V_{\max} = 31.12$  km/s with  $\exp(-(31.12-28)^2/(2 \times 3.0^2)) \approx 0.58$  (contributes  $\sim 100 \text{ cm}^2/\text{g}$  to  $\sigma/m$  at Cloud-9's bulk velocity). A narrower width would suppress the dSph tail while keeping the Cloud-9 bulk contribution intact. The current  $w=4.4$  was chosen to fit §9.12's  $\sigma/m(V_{\max}) = 135.3$ ; a re-fit with the dSph gravothermal constraint included would test whether a narrower width ( $w \lesssim 3$  km/s) is consistent with all 8 channels. Alternatively,

theSilverman+ 2026 merger-history mechanism(sustained mergers suppress gravothermal collapse in roughly half of halos) may apply at dSph scale. Both candidates require N-body with realistic merger histories to test.The gravothermal tension and the Fornax  $\sigma_{\text{HL}}$  outlier are two symptoms of the same underlying question: is the framework's  $\sigma/m$  at dSph velocities correct?The framework is consistent at the observational level ( $\sigma_{\text{eff}} < 1 \text{ cm}^2/\text{g}$ ) but has a real gravothermal tension at dSph scale that requires N-body resolution.

**2.6  $\sigma_{\text{peak\_HH\_1}}$  Sensitivity Sweep****Question:****Different Phase 44 fits give different values for the  $v_1$  resonance peak amplitude. The causality-boundary value is  $\sigma_{\text{peak\_HH\_1}} = 174 \text{ cm}^2/\text{g}$ , the Phase 44 free joint fit prefers  $196.3 \text{ cm}^2/\text{g}$ , and various prescription modes (borrowed, yang, t202) give  $33.9\text{--}104.8 \text{ cm}^2/\text{g}$ .**Is the framework's  $\sigma_{\text{peak}}$  correct, and what range is consistent with Cloud-9 causality + the Cloud-9  $\sigma/m$  floor + dSph non-collapse?****Method:****Sweep  $\sigma_{\text{peak\_HH\_1}}$  {30, 50, 75, 100, 125, 150, 174, 200, 250}  $\text{cm}^2/\text{g}$  using the****paper's  $\sigma/m$  convention****(Gaussian,  $w = 4.4 \text{ km/s}$ , peak at  $v_1 = 29.4 \text{ km/s}$  per Phase 44 free fit `phase44\_joint\_fit.json` best value):****

$$\sigma/m(v) = \sigma_{\text{m\_at\_v}}(0.052, 1.0, v) + \sigma_{\text{peak}} \times \exp(-(v - 29.4)^2 / (2 \times 4.4^2))$$

This is the same parameterization as §2.5 (imported from v192\_dsph\_gravothermal\_sweep.py).Cloud-9 NFW params are canonical and self-consistent: from  $M_{200} = 5 \times 10^9 \text{ M}$  and  $\rho_{\text{crit}} = 1.381 \times 10^{-7} \text{ M/pc}^3$ , derive  $r_{\text{vir}} = 35.09 \text{ kpc}$ ; at each  $c$ , derive  $\rho_s = (200/3) \times c^3 \times \rho_{\text{crit}} / [\ln(1+c) - c/(1+c)]$ ,  $r_s = r_{\text{vir}}/c$ , and  $V_{\text{max}}$  at  $r_{\text{max}} = 2.16 r_s$ (self-consistent at each  $c$ , not held constant). This gives:

-c = 12: $V_{\text{max}} = 31.12 \text{ km/s}$ ,  $\rho_s = 9.69 \times 10^{-3}$ ,  $r_s = 2.92 \text{ kpc}$

-c = 4: $V_{\text{max}} = 25.59 \text{ km/s}$ ,  $\rho_s = 7.28 \times 10^{-4}$ ,  $r_s = 8.77 \text{ kpc}$

Three constraints are checked for each  $\sigma_{\text{peak}}$  value:

- 1.Cloud-9 causality(paper's §9.12):  $t_{\text{core}} > 3 \times t_{\text{cross}}$ . At each  $c$ , use the canonical NFW self-consistent  $V_{\text{max}}$ .
- 2.Cloud-9  $\sigma/m$  floor(BLN24 / Ohana+ published floor): $\sigma/m(v = 28) \geq 50 \text{ cm}^2/\text{g}$ at the resonance peak velocity (BLN24). Floor velocity is  $v = 28$ , distinct from  $V_{\text{max}}$ .
- 3.Fornax  $\sigma_{\text{HL}}$  outlier:  $\sigma_{\text{HL\_required}} = (\sigma_{\text{eff\_published}} - f_{\text{H}}^2 \sigma_{\text{HH}}) / (2 f_{\text{H}} f_{\text{L}})$  at  $f_{\text{H}} = 0.30$ .  $\sigma_{\text{HL}} < 0 = \text{unphysical}$ . Distinguishmarginal( $|\sigma_{\text{HL}}| \leq 0.15$ ) fromsubstantive( $|\sigma_{\text{HL}}| > 0.15$ ).

Code: scripts/v192\_a\_phase44\_sigma\_peak\_sensitivity.py.Findings table:|  $\sigma_{\text{peak}}$  |  $\sigma/m(28)$  |  $\sigma/m(V_{\text{max}} c=12)$  | c=12  $t_{\text{core}}$  | c=12 ratio | c=12 verdict |  $\sigma/m(V_{\text{max}} c=4)$  | c=4  $t_{\text{core}}$  | c=4 ratio | c=4 verdict |  $\sigma_{\text{HL}}$  req |

|---|---|---|---|---|---|---|---|---|---|

| 30 | 30.19 | 23.49 | 0.524 Gyr | 5.71 | \*\*OK\*\* | 26.03 | 22.91 Gyr | 68.35 | OK | -0.08 (marginal) |

| \*\*50\*\* | \*\*50.19\*\* | \*\*39.04\*\* | \*\*0.315 Gyr\*\* | \*\*3.43\*\* | \*\*OK\*\* | \*\*43.25\*\* | \*\*13.79 Gyr\*\* | \*\*41.14\*\* | \*\*OK\*\* | \*\*-0.13 (marginal)\*\* |

| 75 | 75.19 | 58.47 | 0.211 Gyr | 2.29 | below cap | 64.77 | 9.21 Gyr | 27.47 | OK | -0.20 (substantive) |

| 100 | 100.19 | 77.91 | 0.158 Gyr | 1.72 | below cap | 86.29 | 6.91 Gyr | 20.62 | OK | -0.27 (substantive) |

| 125 | 125.19 | 97.34 | 0.127 Gyr | 1.38 | below cap | 107.81 | 5.53 Gyr | 16.50 | OK | -0.34 (substantive) |

| 150 | 150.19 | 116.77 | 0.105 Gyr | 1.15 | below cap | 129.33 | 4.61 Gyr | 13.76 | OK | -0.41 (substantive) |

| \*\*174\*\* | \*\*174.19\*\* | \*\*135.43\*\* | \*\*0.091 Gyr\*\* | \*\*0.99\*\* | \*\*below cap\*\* | \*\*149.99\*\* | \*\*3.98 Gyr\*\* | \*\*11.86\*\* | \*\*OK\*\* | \*\*-0.47 (substantive)\*\* |

| 200 | 200.19 | 155.64 | 0.079 Gyr | 0.86 | below cap | 172.37 | 3.46 Gyr | 10.32 | OK | -0.54 (substantive) |

| 250 | 250.19 | 194.51 | 0.063 Gyr | 0.69 | below cap | 215.42 | 2.77 Gyr | 8.26 | OK | -0.68 (substantive) | (Cloud-9 verdict "below cap" =  $t_{\text{core}}/t_{\text{cross}} < 3$ , fails the paper's §9.12 causality criterion.  $\sigma/m(V_{\text{max}})$  differs at c=12 vs c=4 because  $V_{\text{max}}$  is canonical NFW self-consistent at each  $c$ . Fornax  $\sigma_{\text{HL}}$  req "marginal" =  $|\sigma_{\text{HL}}| \leq 0.15$ ; "substantive" =  $|\sigma_{\text{HL}}| > 0.15$ .)\*\*Three key results:\*\*1.\*\*At  $c = 12$  ( $\Lambda$ CDM-conservative, §9.12 analytical-only):\*\*the continuous viable  $\sigma_{\text{peak}}$  window is\*\*[49.81, 57.24]  $\text{cm}^2/\text{g}$  (width  $\sim 7.43$ )\*\*— a knife-edge. This is computed using the exact  $1/\sigma_{\text{m}}(V_{\text{max}})$  scaling (ratio =  $K/\sigma_{\text{m}}(V_{\text{max}})$  with  $K \approx 134.0$ ). On the swept grid {30, 50, 75, 100, 125, 150, 174, 200, 250},  $\sigma_{\text{peak}} = 50 \text{ cm}^2/\text{g}$  is the only grid value inside this window (ratio = 3.43).  $\sigma_{\text{peak}} = 30$  fails the floor ( $\sigma/m(28) = 30.19 < 50$ );  $\sigma_{\text{peak}} \geq 75$  fail  $c = 12$  causality (ratio  $< 3$ ). The framework's canonical  $\sigma_{\text{peak}} = 174$  fails  $c = 12$  causality (ratio = 0.99).\*\*The grid result is a discretization check; the continuous window [49.81, 57.24] is the physical c-M tension at  $\Lambda$ CDM-standard concentration.\*\*2.\*\*At  $c = 4$  (Ohana+ physical anchor, §9.12):\*\*on the swept grid,\*\* $\sigma_{\text{peak}}$  {50, 75, ..., 250} all pass BOTH constraints.\*\*Continuous intersection: [49.81, 250]  $\text{cm}^2/\text{g}$  (limited by sweep range). At  $\sigma_{\text{peak}} = 174$ :  $\sigma/m(V_{\text{max}} = 25.59) = 149.99$ ,  $t_{\text{core}} = 3.98 \text{ Gyr}$  — see §9.12 reconciliation below.

3.Fornax  $\sigma_{\text{HL}}$  outlier:At  $\sigma_{\text{peak}}$  {30, 50},  $\sigma_{\text{HL\_required}}$  is marginal (-0.08, -0.13 — could be rounding/systematic). At  $\sigma_{\text{peak}} \geq 75$ ,  $\sigma_{\text{HL}}$  is substantive ( $\leq -0.20$ ). At  $\sigma_{\text{peak}} = 174$  (canonical),  $\sigma_{\text{HL}} = -0.47$ .The outlier is structural— consequence of  $v_1$  Gaussian tail reaching dSph velocities.Reconciliation with §9.12:§9.12 reports  $\sigma_{\text{peak}} = 174$ ,  $t_{\text{core}} = 91 \text{ Myr}$  at  $c = 12$  (this sweep reproducesexactly:  $t_{\text{core}} = 0.091 \text{ Gyr}$ ;  $\sigma/m(V_{\text{max}}=31.12) = 135.43$  matches §9.12's 135.3 within 0.1% rounding). At  $c = 4$ , §9.12 reports  $t_{\text{core}} = 4.42 \text{ Gyr}$ . This sweep computest\_core = 3.98 Gyr — an 11% discrepancy from §9.12's 4.42 Gyr. 12's  $\sigma/m = 135.3$  being held constant from the c=12 case when applied to c=4, where the canonical NFW self-consistent  $\sigma/m(V_{\text{max}}=25.59) = 150.0$ .§9.12's "same  $\sigma/m = 135.3$  at  $c = 4$ " is internally inconsistent( $\sigma/m$  depends on  $V_{\text{max}}$  which differs at each  $c$ ). The canonical-NFW sweep gives the self-consistent answer: 0.091 Gyr at c=12 (exact match), 3.98 Gyr at c=4 (corrects §9.12's 11% inconsistency). :§9.12 uses  $V_{\text{max}} = 31.12$  (NFW  $V_{\text{max}}$  at  $r_{\text{max}} = 2.16 r_s$ ). The Cloud-9  $\sigma/m$  floor is defined at  $v = 28$  (resonance peak, BLN24). Two different velocities serving two purposes; both used correctly here. The  $V_{\text{max}}$  now varies at each  $c$  per canonical NFW, rather than held fixed.SPARC consistency: $\sigma/m(v = 100) = 0.052 \text{ cm}^2/\text{g}$  constant across the sweep. SPARC is not a constraint on  $\sigma_{\text{peak}}$  under the paper convention.Parameterization comparison (sanity check):The Breit-Wigner parameterization used in v19.2-A v1 (and the Phase 44 multi-channel fit) gives  $\sigma/m$  values that differ from the paper's Gaussian by up to  $\sim 17\times$  at UFD velocities. The Gaussian (paper convention) is what §2.5 and

§9.12 state. Bottom line: Under the paper's own  $\sigma/m$  convention (Gaussian  $w = 4.4$ ) and causality criterion (ratio  $> 3$ ), with self-consistent canonical NFW at each  $c$ :

-At  $c = 12$ : the continuous viable window is  $[49.81, 57.24] \text{ cm}^2/\text{g}$  (knife-edge, width  $\sim 7.43$ ). On the swept grid, only  $\sigma_{\text{peak}} = 50$  passes both constraints (ratio = 3.43).  $\sigma_{\text{peak}} \geq 75$  fails  $c = 12$  causality.

-At  $c = 4$  (Ohana+ anchor, §9.12 physical anchor): continuous window is  $[49.81, 250] \text{ cm}^2/\text{g}$ . On the swept grid,  $\sigma_{\text{peak}} \in \{50, \dots, 250\}$  all pass both. Framework is consistent at  $c = 4$ . -Fornax  $\sigma_{\text{HL}}$  outlier is marginal for  $\sigma_{\text{peak}} \leq 50$ , substantive for  $\sigma_{\text{peak}} \geq 75$ .

The  $c = 12$  vs  $c = 4$  distinction is itself the  $c$ -M tension against  $\Lambda\text{CDM}$  (Ohana+ 2026). At  $c = 12$ , the Balberg+ analytical formula is unreliable for  $\sigma_{\text{peak}} \geq 75$ ; N-body is required, as §9.12 acknowledges. Synthesis — what  $\sigma_{\text{peak}}$  reconciles Cloud-9 with the framework's causality criterion? The framework is asking: given the published  $\sigma/m \geq 50 \text{ cm}^2/\text{g}$  floor (BLN24) and the paper's own causality criterion (ratio  $> 3$ ), what  $\sigma_{\text{peak}}$  values satisfy both? At  $\Lambda\text{CDM}$ -standard concentration ( $c = 12$ ), the answer is a knife-edge continuous window  $[49.81, 57.24] \text{ cm}^2/\text{g}$  (width  $\sim 7.43$ ) — a narrow viable band centered at  $\sigma_{\text{peak}} \approx 53.5 \text{ cm}^2/\text{g}$ . On the swept grid, only  $\sigma_{\text{peak}} = 50$  falls inside this band (ratio = 3.43, just above the 3.0 cap). At Ohana+ concentration ( $c = 4$ ), the answer is an open interval  $[49.81, 250] \text{ cm}^2/\text{g}$  across the entire swept range. The framework's canonical  $\sigma_{\text{peak}} = 174$  satisfies the floor and passes causality only at  $c = 4$  (Ohana+ physical anchor) — not at  $c = 12$  ( $\Lambda\text{CDM}$ -conservative). The synthesis: the  $\sigma_{\text{peak}}$  is not a free parameter once both the Cloud-9 floor and the causality criterion are imposed; at  $c = 12$ , it is constrained to  $\sim 50 \text{ cm}^2/\text{g}$  (a factor of  $3.5\times$  below the canonical framework value of 174); at  $c = 4$ , it is constrained to  $\sigma_{\text{peak}} \geq \sim 50 \text{ cm}^2/\text{g}$  (consistent with the canonical framework value). Footnote —  $V_{\text{max}} = 31.12$  derivation:  $V_{\text{max}} = 31.12$  at  $c = 12$  is derived from the canonical NFW profile at  $M_{200} = 5 \times 10^9 \text{ M}$ ,  $\rho_{\text{crit}} = 1.381 \times 10^{-7} \text{ M/pc}^3$ , evaluated at  $r_{\text{max}} = 2.1626 r_s$ . This is the same derivation §9.12 used ( $r_{\text{max}} = 2.16 r_s$ , post v18.43 T215 IC correction). The match is a consistency check, not a coincidence. At  $c = 4$ , the same derivation gives  $V_{\text{max}} = 25.59 \text{ km/s}$  — §9.12 implicitly holds  $\sigma/m(V_{\text{max}})$  constant across  $c$  rather than re-deriving  $V_{\text{max}}$ , which is the source of the 11%  $t_{\text{core}}$  discrepancy this sweep corrects.

**2.7 External Consistency Checks** \*\*Question:\*\* Does the framework's Cloud-9 c-M tension and core-radius prediction agree with external observations (Ohana+ 2026 SIDM tension; Nadler+ 2025 SIDM Concerto)? \*\*Method:\*\* Two independent consistency checks, both at Cloud-9 mass scale ( $\sim 5 \times 10^9 \text{ M}$ ). \*\*Check 1 — v19.2-B: Ohana+ 2026 c-M tension reproduction\*\* ('scripts/v192\_b\_ohana3p2sigma\_reproduction.py').

Scatter conventions used in this paper:

- 0.16 dex = DK14 (Diemer & Kravtsov 2014) full-population simulation scatter, Table 1 (verified against arXiv:1407.4730 / ApJ 799, 108). Endorsed by DJ19 §2.1 via "DK15" citation. Citation chain Ohana+  $\rightarrow$  DJ19  $\rightarrow$  DK14/DK15 fully verified.

- 0.085 dex (sensitivity-only): no published source reports exactly this value; closest candidates are Macciò 2008 and Dutton & Macciò 2014 ( $\approx 0.10$  dex). The  $6.18\sigma$  intrinsic tension using this value is illustrative only.

- 0.110 dex (sensitivity-only): no published source. Illustrative only.

Mechanism of the apparent  $6.18\sigma$  vs Ohana+  $3.2\sigma$  gap (resolved): the factor-1.93 ratio is the scatter-convention difference (0.085 dex vs 0.16 dex), not an internal disagreement. With Ohana+'s scatter (0.16 dex, the DK14/DK15 full-population prescription), the simplified pipeline is consistent with Ohana+  $3.2\sigma$  within rounding tolerance ( $3.29\sigma$  MCMC vs  $3.16\sigma$  fiducial). The framework's c-M tension prediction at the Ohana+ scatter convention is correct. The  $6.18\sigma$  illustrative number is secondary; the headline finding is the  $3.16\sigma$  consistency check.

Result. The framework's Cloud-9 c-M tension prediction is consistent with Ohana+ 2026's  $3.2\sigma$  at the fiducial under the 0.16 dex convention (DK14/DK15 scatter, Ohana+ 2026 §3.1 line 29 verbatim). The simplified pipeline gives  $3.16\sigma$  at the fiducial ( $c=4.0$ ,  $M=4.7 \times 10^9 \text{ M}$ ,  $\tau=0.18$ ) and  $3.29\sigma$  at the MCMC best-fit ( $c=3.8171$ ,  $M=4.7571 \times 10^9 \text{ M}$ ) — both within  $0.09\sigma$  of Ohana+'s published  $3.20\sigma$ . The  $6.18\sigma$  at 0.085 dex is footnote-level only (illustrative, source UNVERIFIED for that exact scatter value).

Check 2 — v19.2-C: SIDM Concerto subhalo consistency (scripts/v192\_c\_concerto\_subhalo\_cloud9.py).

-Reference: Nadler+ 2025, arXiv:2503.10748 (SIDM Concerto cosmological N-body simulation; public release at Zenodo 10.5281/zenodo.14933624). Ohana+ 2026 use the Halo004 (GroupSIDM-147 model) subset of Concerto for their analog analysis; our v19.2-C v1 uses the Halo416 (MilkyWaySIDM model) host. This is a real caveat — different hosts have different SIDM models. The framework-level conclusion (Concerto SIDM halos at Cloud-9 mass have  $r_c \approx 0.5\text{--}1 \text{ kpc}$ ) is robust to host choice; the precise median varies by host and model.

-Data: MW-mass host (MW\_Halo416) parametric catalog (2.8 MB). 2489 SIDM subhalos total; 267 in Cloud-9-mass range ( $1 \times 10^9\text{--}1 \times 10^{10} \text{ M}$ ); 264 with valid parametric fits.

-Result: Median SIDM core radius  $rc_1 = 0.82 \text{ kpc}$  [16-84: 0.53–1.20 kpc]; median  $rc_1/R_{\text{max}} = 0.216$ ; median  $R_{\text{max}} = 3.65 \text{ kpc}$ .

-Cloud-9 expectation (Yang+ 2024 SIDM parametric model for  $\tau = 0.18$ ):  $rc \approx 0.5 \pm 0.3 \text{ kpc}$ .

-Match:  $1.07\sigma$  — within  $2\sigma$ . Verdict: What this establishes: At  $M = 1 \times 10^9\text{--}1 \times 10^{10} \text{ M}$  mass scale, SIDM N-body halos from Nadler+ 2025 Concerto have median core radii consistent with the Yang+ 2024 parametric model for  $\tau \approx 0.18$ . What this does not establish: That this is the correct core radius for Cloud-9 specifically, because (a) Concerto subhalos are MW satellites with tidal stripping, (b) Cloud-9 is a RELIC near M94 (different environment), (c) single-host statistics (Halo416 vs Ohana+'s Halo004), (d) single SIDM model (MilkyWaySIDM vs Ohana+'s

GroupSIDM-147). Conclusion: The core-radius scaling is consistent across mass scale, but environmental and model differences prevent cross-validation for Cloud-9 specifically. First, our simplified pipeline is consistent with Ohana+ 2026's  $3.2\sigma$  SIDM tension at the fiducial under the 0.16 dex convention (citation chain Ohana+  $\rightarrow$  DJ19 §2.1  $\rightarrow$  DK14/DK15, chain is valid.  $16\sigma$  at the fiducial ( $c=4.0$ ,  $M=4.7\times 10^9$  M,  $\tau=0.18$  — the canonical consistency-check number) or  $3.29\sigma$  at the MCMC-recovered best-fit ( $c=3.8171$ ,  $M=4.7571\times 10^9$  M — the sampled variant) — both within  $0.09\sigma$  of Ohana+'s published  $3.20\sigma$  ( $|3.16-3.20|=0.04\sigma$  and  $|3.29-3.20|=0.09\sigma$ ,  $.13\sigma$  (driven by  $c_{\text{best-fit}}$  sampling, not  $c_{\text{med}}$  change). This is a consistency check at the fiducial (synthetic-data tautology caveat applies); full reproduction requires real BLN24  $N_{\text{HI}}$  data + hydrostatic. The previous v19.2-B.2-v19.2-B.6  $6.18\sigma$  value at 0.085 dex scatter is a footnote-only sensitivity result. Second, the framework's core-radius prediction ( $r_c \approx 0.5 \pm 0.3$  kpc) is consistent with the median core radius of Cloud-9-mass SIDM subhalos in the Nadler+ 2025 SIDM Concerto (0.82 kpc, 16-84: 0.53–1.20), though environmental differences (tidal MW satellites vs isolated RELHIC) and model differences (MilkyWaySIDM vs Ohana+'s GroupSIDM-147) prevent cross-validation for Cloud-9 specifically. Both checks are consistent with the constraint-map framing — the framework's predictions are testable AND is consistent with Ohana+  $3.2\sigma$  at the fiducial under their 0.16 dex convention.---

### 3. Multi-Channel Observational Constraints

Channel count convention: With the Yang+-derived  $f_{\text{H}} \text{ at } r$ , the multi-resonance profile does not simultaneously satisfy the Cloud-9 and dSph channels at Phase 44 parameters, because (a) gravothermal cascade timescale Hubble time at Phase 44  $\sigma/m$ , and (b) Yang+ Fig. 2 segregation is modest ( $f_{\text{L}}$  0.3-0.6), not extreme ( $f_{\text{H}}$  from 0.95 to 0.10) as the placeholder claimed. The honest verdict is that at Phase 44 parameters, the multi-resonance  $\sigma/m(v)$  profile is a phenomenological interpolation through 8 channels, not a first-principles derivation of the underlying physics. The Cloud-9 vs dSph tension is unresolved at Phase 44. The LZ September 2026 direct-detection event (§3.5a) is a separate falsifiability test — NOT a 9th bulk-halo channel. The Cloud-9 vs dSph tension is not a tension between a Cloud-9 observational requirement and the framework's prediction — it is a tension between two framework-chosen benchmarks (Cloud-9  $\sigma/m \geq 50$  cm<sup>2</sup>/g from Elbert+ 2015's largest simulated dwarf-scale cross-section, and dSph  $\sigma/m \leq 0.8$  cm<sup>2</sup>/g from Horigome+ 2025 upper limit). Direct fetch of BLN24 (Benítez-Llambay+ 2024, ApJ 973, 61, arXiv:2406.18643 [1a]) and Anand+ 2025 (arXiv:2508.20157, HST/ACS) confirms Cloud-9's observational data:  $M_{200} = 3.9\text{--}4.3 \times 10^9$  M (BLN24 VLA hydrostatic equilibrium of isothermal gas sphere),  $W_{50} = 12 \pm 1$  km/s (BLN24 VLA-D),  $M_{\text{c}} < 10^{3.5}$  M (Anand+ 2025 HST, 99.5% confidence). Cloud-9's data do not place a  $\sigma/m \geq 50$  cm<sup>2</sup>/g lower bound — that value is Elbert+ 2015's largest simulation cross section for  $v_{\text{rms}} \sim 40$  km/s, adopted by the framework as a benchmark. Conversely, Cloud-9's diffuse non-collapsed structure places an upper bound on  $\sigma/m$  at  $v \sim 20\text{--}30$  km/s: if  $\sigma/m$  were very large at this velocity scale, Cloud-9's halo would have undergone gravothermal collapse on a Hubble time, producing a denser and more thermally-supported gas distribution than is observed. The correct upper-bound value requires N-body modeling of Cloud-9's halo + gas, which this paper does not perform. The two constraints (Elbert+ 2015 simulation lower-end benchmark  $\sim 50$  cm<sup>2</sup>/g vs Cloud-9 collapse-history upper bound  $\sim$  unknown) point in opposite directions and are not directly comparable. The framework's  $\sigma_{\text{peak}} = 174$  cm<sup>2</sup>/g satisfies the Elbert+ 2015 benchmark by construction but does not address the collapse-history upper bound, which is currently unquantified. This is the honest status. Honest status of the framework: The virial theorem calculation  $\sigma^{\text{3D}} = \sqrt{G M_{\text{c}} / (3 r_{\text{c}})} = 12.81$  km/s does NOT depend on  $\sigma_{\text{peak}}$ . It only depends on  $M_{200}$ ,  $c$ ,  $r_{\text{s}}$ ,  $\rho_{\text{s}}$  (NFW parameters from BLN24) and the Sánchez Almeida scalings  $r_{\text{c}} = 0.45 r_{\text{s}}$ ,  $\rho_{\text{c}} = 2.4 \rho_{\text{s}}$ . Counterfactual: change  $\sigma_{\text{peak}}$  from 174 to 50 or 500, the virial velocity is still 12.81 km/s (as long as collapse has not completed). The "prediction" is generic to ANY NFW halo with  $M = 5 \times 10^9$  M,  $c = 4$ ; it does not test  $\sigma_{\text{peak}}$  specifically, and does not distinguish the SIDM framework from  $\Lambda$ CDM or other SIDM models. gas  $W_{50} = 12$  km/s reflects thermal broadening from UVB equilibrium at  $T \sim 10^4$  K ( $c_{\text{s}} \sim 12$  km/s per BLN24 verbatim, giving thermal  $W_{50} = 2.355 \times c_{\text{s}}/\sqrt{3} = 16.3$  km/s for pure isothermal), NOT a combination of thermal + DM velocity dispersion. The correct interpretation is:  $W_{50}$  is set by gas temperature; DM velocity dispersion affects halo structure (core formation), not gas line width directly. The framework's  $\sigma_{\text{peak}} = 174$  cm<sup>2</sup>/g is a phenomenological parameter that does NOT currently make a prediction that distinguishes it from  $\Lambda$ CDM or other SIDM models at Cloud-9 or any other specific observation. The  $\sigma_{\text{peak}}$  does determine the gravothermal phase, and so predicts whether core formation/collapse has occurred — but the resulting core radius and density are Sánchez Almeida scalings of the underlying NFW, not distinguishing features of  $\sigma_{\text{peak}} = 174$  specifically. Position vs unified SIDM models: We adopt  $\sigma/m \sim 50$  cm<sup>2</sup>/g at  $v \sim 30\text{--}40$  km/s as a working benchmark motivated by large cross-sections explored in dwarf-scale SIDM simulations (Elbert, Bullock, Garrison-Kimmel, Rocha, Oñorbe, Boylan-Kolchin 2015 [55a], MNRAS 453, 29; arXiv:1412.1477). Per Elbert+ 2015's abstract, quoted verbatim: "Our work suggests that SIDM cross-sections as large or larger than 50 cm<sup>2</sup>/g remain viable on velocity scales of dwarf galaxies ( $v_{\text{rms}} \sim 40$  km/s)." Elbert+ 2015 used 50 cm<sup>2</sup>/g as the largest SIDM cross-section in their dwarf-galaxy simulation suite, showing it remains viable. It is NOT a direct Cloud-9 observational lower limit (BLN24 [1a] constrains mass/structure from VLA observations, not  $\sigma/m$ ). Under that benchmark, Mace+ 2026 SIDM2v yields  $\sigma_{\text{eff}}(v=28) = 7$  cm<sup>2</sup>/g.18643 (same paper throughout). The reference key "BLN24" has remained stable; We computed  $\sigma_{\text{eff}}$  at  $v=28$  km/s for the SIDM2v model of Mace, Yang, Zeng+ 2026 [50c] (arXiv:2506.14898, two-component SIDM with inter-species mass segregation). Using their published parameters ( $\sigma_{\text{H}}/m_{\text{H}} = 6.89$  cm<sup>2</sup>/g,  $w_{\text{H}} = 275$  km/s;  $\sigma_{\text{x}}/m_{\text{H}} = 1.125$  cm<sup>2</sup>/g,  $w_{\text{x}} = 2200$  km/s;  $m_{\text{H}}/m_{\text{L}} = 3$ ), the maximum  $\sigma_{\text{eff}}$  ( $f_{\text{H}} = 1$ ) at  $v=28$  km/s is  $\sigma_{\text{eff}} \approx 6.89$  cm<sup>2</sup>/g, a factor of  $\sim 7\times$  short of the Elbert+ 2015 working benchmark (50 cm<sup>2</sup>/g). For mass-weighted  $f_{\text{H}} = 0.75$  (assuming equal number density,  $m_{\text{H}} = 3 m_{\text{L}}$ ),  $\sigma_{\text{eff}} \approx 4.4$  cm<sup>2</sup>/g — still  $\sim 11\times$  short. The original Kaplinghat, Tulin, Yu 2016 PRL 116, 041302 [50d] (arXiv:1508.03339) "Dark Matter Halos as Particle Colliders" found  $\sigma/m \approx 2$  cm<sup>2</sup>/g on galaxy scales to  $\sigma/m \approx 0.1$  cm<sup>2</sup>/g on cluster scales from a unified fit to 12 dwarfs/LSBs and 6 clusters. Scope of this comparison: " $7\times$  short" compares  $\sigma_{\text{eff}}$  at exactly  $v=28$  km/s (Cloud-9 velocity) to the working benchmark of 50 cm<sup>2</sup>/g. The " $3.6\times$  short" number mentioned in earlier versions of this paragraph compared 50

to Mace+'s stated " $\sigma_{\text{eff}} \approx 14 \text{ cm}^2/\text{g}$ " equivalent regime — these are two different comparisons:  $\sigma_{\text{eff}}$  at  $v=28$  vs  $\sigma/m$  at the gravothermal-collapse timescale. The  $\sigma_{\text{eff}}$  at  $v=28$  is the apples-to-apples comparison; the  $14 \text{ cm}^2/\text{g}$  number is the equivalent one-component  $\sigma/m$  that reproduces Mace+'s observed core-collapse time (at much larger  $V_{\text{max}}$ ). Under the Mace+ SIDM2v parameters,  $\sigma_{\text{eff}}(v=28) \approx 7 \text{ cm}^2/\text{g}$ ,  $\sim 7\times$  below the Elbert+ 2015 working benchmark (50). Reaching that benchmark needs a much larger low- $v$  peak than this particular unified fit, OR resonance structure beyond the published SIDM2v parameter space. Mace+ 2026 note that "core collapse time is close to a one-component  $\sigma/m \approx 14 \text{ cm}^2/\text{g}$ " — this is a separate number at different kinematics. The  $\sigma_{\text{peak}}$  fit remains a phenomenological parameterization, NOT a first-principles prediction.  $v_{\text{trans}}$  for our  $\sigma/m$  parameterization: Standard Yukawa SIDM fits to multi-channel data (Cloud-9  $\sigma/m \geq 50$  at  $v=28$  vs dSph  $\sigma/m \leq 0.8$  at  $v=100$ ) imply  $v_{\text{trans}} \sim 30\text{--}50 \text{ km/s}$  for our framework (canonical SIDM Yukawa estimate; Phase 44 free-fit  $\sigma_{\text{peak}_R0} = 196.3$  vs causality threshold 174 at  $v=29.4$  gives  $v_{\text{trans}} \sim 28.5 \text{ km/s}$ ). The implied mediator mass is  $m_{\phi} \sim 150\text{--}250 \text{ eV}$  for  $m_{\chi} \sim 1 \text{ GeV}$  WIMP, or  $m_{\phi} \sim 1.5\text{--}2.5 \text{ keV}$  for  $m_{\chi} \sim 10 \text{ GeV}$ . This is the missing parameter (path) needed to convert the  $\sigma_{\text{eff}}$  constraint map into a falsifiable joint SIDM+LZ prediction. Caveat: the  $\sigma_{\text{peak}}(v) = 196.3 \cdot \exp(-(v-29.4)^2/(2 \cdot 4.4^2))$  form is the best-fit Gaussian Breit-Wigner peak shape from our framework's parameterization, NOT a first-principles prediction. The peak height  $196.3 \text{ cm}^2/\text{g}$ , center  $29.4 \text{ km/s}$ , and width  $4.4 \text{ km/s}$  are phenomenological fit parameters; whether this peak shape can be realized by a specific UV completion (specific  $m_{\phi}$ , mediator coupling  $g_{\chi}$ , resonance structure) remains an open question that the framework does not currently address. Option 1:  $\sigma_{\text{peak}}$  UV motivation. Standard non-relativistic Yukawa gives  $\sigma/m = g_{\chi}^4 m_{\chi}^2 / (32 \pi m_{\phi}^4) \times (c)^2$  at low velocity ( $v \ll v_{\text{trans}}$ ). For our framework's  $m_{\phi} \sim 200 \text{ eV}$  ( $m_{\chi} \sim 1 \text{ GeV}$ ), this requires  $g_{\chi} \sim 2 \times 10^{-5}$  to reach  $\sigma_{\text{peak}} = 196.3 \text{ cm}^2/\text{g}$ . The coupling is well within perturbativity ( $g_{\chi} / 4\pi \sim 12.6$ ), so the cross-section MAGNITUDE is not the obstacle. However, standard Yukawa is MONOTONICALLY DECREASING with velocity — it cannot peak at  $v = 29.4 \text{ km/s}$ . Our framework's Gaussian Breit-Wigner resonance requires a UV mechanism ON TOP OF the Yukawa background. What UV physics produces a resonance of FWHM  $\approx 4.4 \text{ km/s}$  centered at  $v = 29.4 \text{ km/s}$  with peak amplitude  $196.3 \text{ cm}^2/\text{g}$ ? Possible mechanisms: (a) a second mediator with  $m_{\phi'} \sim m_{\chi} v_{\text{res}}^2 / 2 \sim 50\text{--}100 \text{ eV}$  and coupling tuned for resonance; (c) bound-state formation giving discrete resonances. None is implemented in our framework. Status: amplitude feasible in Yukawa; position/width require additional UV physics not yet specified. Option 2: Joint SIDM+LZ signal under our  $v_{\text{trans}} \sim 30\text{--}50 \text{ km/s}$ . Standard light-mediator DD gives  $\sigma_{\text{SI}}(v) \sim (g_{\chi} g_{\text{portal}})^2 \mu^2 / (\pi m_{\phi}^4 (1 + 2 \mu^2 v^2 / m_{\phi}^2)) \times (c)^2$ . For  $m_{\phi} \sim 200 \text{ eV}$ :  $\sigma_{\text{SI}}(v=28) = 1.74 \times 10^{-12} \times (g_{\chi} g_{\text{portal}})^{-2} \text{ cm}^2$  and  $\sigma_{\text{SI}}(v=250) = 2.73 \times 10^{-16} \times (g_{\chi} g_{\text{portal}})^{-2} \text{ cm}^2$  — a suppression factor  $(28/250)^4 \approx 1.57 \times 10^{-4}$  between Cloud-9 and LZ recoil velocities. LZ's current bound at  $m_{\chi} \sim 1 \text{ GeV}$  is  $\sim 10^{-44} \text{ cm}^2$ , requiring  $g_{\chi} g_{\text{portal}} \sim 6 \times 10^{-15}$ . This is much smaller than  $g_{\chi} \sim 2 \times 10^{-5}$  for  $\sigma_{\text{DM-DM}}$ . Under shared-mediator hypothesis,  $g_{\text{portal}} \sim 10^{-10}$ . This is theoretically possible (kinetic mixing  $\epsilon \sim 10^{-10}$  in dark photon models) and makes the shared-mediator hypothesis testable: a future XENONnT/LZ signal at  $\sigma_{\text{SI}} \sim 10^{-44} \text{ cm}^2$  with  $m_{\chi} \sim 1\text{--}10 \text{ GeV}$  would require  $g_{\chi} g_{\text{portal}} \sim 6 \times 10^{-15}$  — compatible with our framework's Yukawa requirement only if  $g_{\text{portal}} \sim 10^{-10}$ . The LZ230616  $248 \text{ keV}$  event (if real, not  $2.6\sigma$  fluctuation) requires  $m_{\phi} \sim 100 \text{ MeV}$  per Das+ 2026 analysis — incompatible with our framework's  $m_{\phi} \sim 200 \text{ eV}$  range, so the shared-mediator hypothesis fails under our framework. Option 3:  $\sigma_{\text{eff}}$  sensitivity study. Under Mace+ 2026 SIDM2v parameters,  $\sigma_{\text{eff}}(v=28)$  is robustly  $< 10 \text{ cm}^2/\text{g}$  across realistic parameter variations:

-  $f_{\text{H}} [0.50, 1.00]$ :  $\sigma_{\text{eff}} = 2.86\text{--}6.89 \text{ cm}^2/\text{g}$  ( $17\times$  to  $7\times$  short of Elbert+ 2015's 50)

-  $v [20, 40] \text{ km/s}$ :  $\sigma_{\text{eff}} = 6.89 \text{ cm}^2/\text{g}$  (saturated at low- $v$  plateau;  $7.26\times$  short)

-  $w_{\text{H}} [200, 350] \text{ km/s}$ :  $\sigma_{\text{HH}} = 6.4\text{--}19.5 \text{ cm}^2/\text{g}$  (note: published Mace+ value is 6.89;  $>10$  requires different  $\sigma_{\text{H}_0}$ )

-  $\sigma_{\text{H}_0} [5, 20] \text{ cm}^2/\text{g}$ :  $\sigma_{\text{HH}} = 5.0\text{--}20.0 \text{ cm}^2/\text{g}$  (at  $w_{\text{H}}=275$ ) Even at  $\sigma_{\text{H}_0} = 20 \text{ cm}^2/\text{g}$  ( $3\times$  the Mace+ value),  $\sigma_{\text{eff}}$  is  $2.6\times$  short of 50. The  $\sigma_{\text{peak}}$  gap to Elbert+ 2015's benchmark is structural to multi-component Yukawa fits, not a fine-tuned artifact of Mace+ parameters. Reaching the benchmark needs  $\sigma_{\text{H}_0} > 50 \text{ cm}^2/\text{g}$  (single-component heavy species with  $\sigma/m \sim 50$  at  $v \sim 0$ ), not multi-component SIDM2v. Synthesis: All three options converge. Option 1:  $\sigma_{\text{peak}}$  amplitude is Yukawa-feasible; resonance position/width needs additional UV physics. Option 2: shared-mediator requires  $g_{\text{portal}} \sim 10^{-10}$  — testable. Option 3:  $\sigma_{\text{eff}}$  gap to Elbert+ 2015 is robust. Open question for v19.2-D future work: what UV physics produces the  $v = 29.4 \text{ km/s}$  Breit-Wigner resonance AND admits a portal coupling  $g_{\text{portal}}$  satisfying LZ bounds?

**3.1 SPARC rotation curves** \*\*Data:\*\* 127 galaxies from the Spitzer Photometry and Accurate Rotation Curves (SPARC) sample [14]. \*\*Constraint:\*\*  $\sigma/m$  at  $v \approx 100 \text{ km/s}$  should be  $\approx 0.07 \text{ cm}^2/\text{g}$  for the rotation curves to be consistent with the observed  $V_{\text{flat}}$  in the inner core. Phase 33d tested all 127 SPARC galaxies;  $115/127 = 90.6\%$  pass the  $V_{\text{flat}}$  test with the multi-resonance  $\sigma/m(v)$  architecture. This is consistent with, but not better than, single-Yukawa SIDM. \*\*Result:\*\* Multi-resonance architecture is consistent with SPARC.

### 3.2 Cloud-9 ultra-diffuse galaxies

Note on  $\sigma/m$  values used in this paper: The Cloud-9  $\sigma/m(v=28) \approx 100 \text{ cm}^2/\text{g}$  is the framework's evaluation at the kinematic input velocity ( $v=28 \text{ km/s}$ ). The  $\sigma_{\text{peak}} = 174 \text{ cm}^2/\text{g}$  is the Gaussian resonance peak height (the maximum of  $\sigma/m(v)$  over velocity, occurring at  $v_{\text{target}} = 29.4 \text{ km/s}$ ). These are two different physical quantities: the  $174 \text{ cm}^2/\text{g}$  is the peak amplitude of the  $v_1$  Breit-Wigner feature; the  $100 \text{ cm}^2/\text{g}$  is the value at  $v=28 \text{ km/s}$ , where  $\sigma/m(v)$  is offset from the peak by the resonance width. Both are used consistently throughout this paper.

Data: Cloud-9, a Reionization-Limited H I Cloud (RELHIC) candidate near M94, discovered by Zhou+ 2023 [15a] (FAST H I detection,  $M_{\text{HI}} \approx 1.4 \times 10^6 M_{\odot}$ ,  $W50 \approx 20 \text{ km s}^{-1}$ ). The hydrostatic-equilibrium analysis of Benítez-Llambay, Dutta, Fumagalli & Navarro 2024 [15b] (ApJ 973, 61) yields a  $\sigma/m \approx 50 \text{ cm}^2/\text{g}$  floor at  $v \approx 28 \text{ km s}^{-1}$ , with a halo mass  $M_{200} \approx 5 \times 10^9 M_{\odot}$  (consistent with  $M_{\text{crit}}$ ). Stellar-mass upper limits on any luminous

counterpart have been refined by Anand+ 2025 [15c] (HST/ACS star-counts,  $M < 10^{3.5} M$ , 99.5% CL; baseline comparison with Leo T,  $\mu_{0,V} \approx 27$  mag arcsec $^{-2}$ ) and Trujillo+ 2026 [15d] (GTC/HIPERCAM integrated light at surface-brightness limits 31.4 mag arcsec $^{-2}$  in g, 31.0 in r  $\sim 10\times$  deeper than previous DESI Legacy / HST searches,  $M < 1.6\times 10^4 M$  assuming old, metal-poor population; surface mass density  $< 0.01 M/pc^2$ ). Trujillo+ 2026 [15d] is the strongest stellar-mass bound on Cloud-9 to date, with the Leo T comparison caveat noted in Anand+ 2025 [15c] potentially underestimating the bound by 2-3 mag for galaxies with the diffuse extended morphology of Cloud-9. The  $\sigma/m(28) \approx 100 \text{ cm}^2/\text{g}$  value we adopt as the multi-resonance working anchor is an internal derivation, consistent with the Benítez-Llambay+ 2024 published floor and chosen to provide a concrete quantitative target. Constraint: At  $v \approx 28 \text{ km/s}$ ,  $\sigma/m$  should be high ( $50 \text{ cm}^2/\text{g}$  published; we use  $\approx 100 \text{ cm}^2/\text{g}$  as the internal target). Result: Multi-resonance architecture satisfies this via the  $v_i = 29 \text{ km/s}$  Breit-Wigner peak. With the corrected kinematics, the actual peak of  $\sigma/m(v)$  occurs at  $v_{\text{peak},1} = v_{\text{target},1} \approx 29 \text{ km/s}$  (NOT at  $1.4\times v_{\text{target}} \approx 41 \text{ km/s}$  as previously stated in v1.6–v1.8);  $\sigma/m(v_{\text{peak},1}) \approx 197 \text{ cm}^2/\text{g}$ , comfortably above the Cloud-9 target of  $\sigma/m \approx 100 \text{ cm}^2/\text{g}$ . At  $v = 28 \text{ km/s}$  (the Cloud-9 kinematic  $v$ ),  $\sigma/m \approx 100 \text{ cm}^2/\text{g}$ . At the kinematic input velocity  $v_{\text{target},1} = 29 \text{ km/s}$ ,  $\sigma/m(v_{\text{target},1}) \approx 197 \text{ cm}^2/\text{g}$  — i.e.,  $v_{\text{target},1}$  IS the location of the maximum of  $\sigma/m(v)$  under the corrected kinematics.

**3.2c Cloud-9 as a concentration-mass tension (Ohana, Zhang & Yu 2026 [15e])\*\*Reframing:\*\*The  $\sigma/m \geq 50 \text{ cm}^2/\text{g}$  floor at  $v \approx 28 \text{ km/s}$  is a\*\*necessary\*\*constraint, but Ohana, Zhang & Yu 2026 [15e] (arXiv:2608.04362, Aug 2026) show that Cloud-9's \*primary\* tension is with the cosmological concentration–mass (c-M) relation, not solely with  $\sigma/m$  magnitude. Their MCMC analysis finds:**

-Best-fit SIDM:  $\sigma/m = 483 \text{ cm}^2/\text{g}$ ,  $M_{200} = 4.7\times 10^9 M$ ,  $c_{200} = 4.0 - 3.2\sigma$  below median of the c-M relation.

-Extreme case:  $\sigma/m = 2.1\times 10^4 \text{ cm}^2/\text{g}$  (gravothermal core-collapse phase).

-CDM alternative: requires  $7\sigma$  below the c-M median —strongly disfavored versus SIDM. Cross-link to §10.4d (Cloud-9 systematic upper bound): the c-M reframing here (a cosmological tension with the c-M relation per Ohana+ 2026) is complementary to the environmental-systematic reframing in §10.4d (a systematic-uncertainty upper bound per Turini & Benítez-Llambay 2026). Together, these two reframings frame Cloud-9 as a ( $\sigma/m$ ,  $c_{200}$ , environment) joint tension rather than a  $\sigma/m$ -only constraint. Neither reframing alone is sufficient; both are honest characterizations of the current state of the data. Implication for our framework: Our adopted  $\sigma/m \approx 100 \text{ cm}^2/\text{g}$  at  $v = 28 \text{ km/s}$  sits below the Ohana+ best-fit ( $483 \text{ cm}^2/\text{g}$ ) and well below the extreme ( $2.1\times 10^4 \text{ cm}^2/\text{g}$ ). The factor-of-5 gap between our value and the Ohana+ best-fit does NOT necessarily mean our phenomenology is wrong — it means our phenomenology under-shoots Cloud-9's gas profile by a factor of  $5\times$  in  $\sigma/m$ , which Ohana+ resolves by allocating additional tension to the c-M relation. A full c-M tension analysis is out of scope for v19.0 (it requires marginalizing over  $c_{200}$ ,  $M_{200}$ ,  $\sigma/m$  jointly, which we do not do) but is the natural next step in §11 future-work. Concentration–mass corroboration from Silverman+ 2026 [54]: Silverman+ 2026 (arXiv:2606.02566, "Mergers Matter") shows via N-body that 3 of 6 host halos at  $M = 10^{10} M$  with  $\sigma/m = 70 \text{ cm}^2/\text{g}$  collapse within a Hubble time. This is consistent with the Ohana+ picture: a moderately-elevated  $\sigma/m$  (70 vs our 100) can drive gravothermal core-formation that brings  $c_{200}$  closer to the observed value ( $3.2\sigma$  tension is reasonable for a host-halo on the cusp of collapse). Honest framing: the  $\sigma/m \geq 50$  floor remains our primary Cloud-9 constraint. The Ohana+ c-M reframing is complementary: it suggests that the next paper version should treat Cloud-9 as a joint ( $\sigma/m$ ,  $c_{200}$ ) tension rather than a  $\sigma/m$ -only constraint. This is not an automatic upgrade — it requires implementing the c-M likelihood alongside the  $\sigma/m$  likelihood, which is non-trivial and deferred to v19.1.

**3.3 JVAS B1938+666 strong-lensing perturber\*\*Data:\*\*Vegetti et al. 2010 [16] observed a small-density perturbation in the JVAS B1938+666 strong-lensing system that has been \*interpreted\* (in subsequent lensing-modelling literature) as requiring  $\sigma/m(15) \approx 100 \text{ cm}^2/\text{g}$ . Note: this constraint is a derived interpretation of the lensing-perturbation signal rather than a direct cross-section measurement, and it carries substantial modelling uncertainty. The perturber mass is  $(1.13\pm 0.04)\times 10^6 M$  within a projected radius of 80 pc at  $z = 0.881$  [23].\*\*Excluded from the 4-channel fit:\*\*JVAS is a single  $\sim 10^6 M$  substructure measurement,\*\*not a measurement of the bulk halo  $\sigma/m(v)$ \*\*. Including it in the same 4-channel table as Cloud-9, dSph, SPARC, Cluster conflates two different physics regimes. The reframing below (Yu 2026 [23]) confirms JVAS is complementary substructure physics, not an additional bulk channel.\*\*Constraint (as commonly stated):\*\* $\sigma/m \approx 100 \text{ cm}^2/\text{g}$  at  $v \approx 15 \text{ km/s}$ .\*\*Result:\*\*Tension with the multi-resonance architecture at face value. The model achieves  $\sigma/m(15) \approx 4.2 \text{ cm}^2/\text{g}$ , a factor of  $\sim 24\times$  below the JVAS target of  $\sigma/m(15) \approx 100 \text{ cm}^2/\text{g}$ . (Earlier reports in this paper sometimes quote a factor of  $\sim 84\times$ , which refers to a different reference velocity —  $v=15$  is the canonical JVAS velocity used here.)\*\*Reframing:\*\*The JVAS perturber is\*\*a single dense  $\sim 10^6 M$  substructure\*\*, not a measurement of the bulk  $\sigma/m$  of the host halo. Yu (2026) [23] shows via N-body simulation that core-collapsed SIDM halos of mass  $\sim 10^6 M$  naturally produce the JVAS perturber density profile — this is\*\*gravothermal core-collapse physics\*\*at the subhalo mass scale, not the bulk cross-section at  $v \approx 15 \text{ km/s}$ . Our phenomenology at  $v \approx 15 \text{ km/s}$  applies to the\*\*host halo\*\*( $M_{\text{halo}} \sim 10^9 M$ ), where the core-collapse enhancement does NOT apply. The structural limit is therefore\*\*not a failure of our  $\sigma/m(v)$  parameterization\*\*but rather a statement that the JVAS perturber requires substructure physics outside our bulk-phenomenology scope. This is consistent with the §10.4c.A5 verdict (gravothermal enhancement  $\sim 100\times$ , needs  $3125\times$ ) — the missing factor is from the substructure being in core-collapse state, not from our cross-section being wrong.\*\*Cross-confirmation (Fornax 6, Yu 2026 [23]):\*\*Yu (2026) [23] further shows that the same  $\sim 10^6 M$  core-collapsed SIDM halo density profile simultaneously explains (a) the JVAS B1938+666 perturber, (b) the GD-1**

stellar stream perturber, and (c) the Fornax 6 stellar cluster in the Fornax dwarf spheroidal ( $M \approx 7.2 \times 10^3 M_\odot$ ,  $r_h \approx 11$  pc,  $\sigma \approx 5.6$  km/s, anomalously high  $M/L \approx 15\text{--}258$ ; Pace et al. 2021, Peñarrubia et al. 2024). Fornax 6 is therefore an **independent observational anchor** for the same  $\sim 10^6 M_\odot$  core-collapsed SIDM halo physics, at a different cosmic location. Our phenomenology covers dSph ( $M_{\text{halo}} \sim 10^9 M_\odot$ ) and cluster ( $M_{\text{halo}} \sim 10^{14} M_\odot$ ) scales; the JVAS / GD-1 / Fornax 6 anchors at  $10^6 M_\odot$  are **complementary** substructure physics, not in tension with our bulk  $\sigma/m(v)$ . **T204 substructure test:** A  $10^6 M_\odot$  subhalo at the JVAS perturber location has  $r_{\text{vir}} = 1.5$  kpc ( $c = 15$ ),  $V_{\text{max}} = 1.69$  km/s, NFW scale density  $\rho_s = 4.34 \times 10^{-2} M/\text{pc}^3$ . At Phase 44 canonical parameters ( $\sigma_0 = 0.052 \text{ cm}^2/\text{g}$  at  $v_{\text{ref}} = 100$  km/s,  $a_{\text{slope}} = 1.93$ ),  $\sigma/m(v_{\text{max}}) = 0.052 \times (100/1.69)^{1.93} \approx 137 \text{ cm}^2/\text{g}$  at the subhalo's virial velocity. Using the Balberg+ 2002 Eq. 22 normalization,  $t_{\text{core}} \approx 13$  Myr. The halo crossing time is  $t_{\text{cross}} = r_s/V_{\text{max}} \approx 58$  Myr.  **$t_{\text{core}} < t_{\text{cross}}$  — causality violation.** The substructure core-collapse prediction is unphysical at  $a_{\text{slope}} = 1.93$ .

The historical Option A flattening ( $a_{\text{slope}} = 1.0$ ) gives  $\sigma/m(v_{\text{max}}) = 3.08 \text{ cm}^2/\text{g}$ ,  $t_{\text{core}} \approx 560$  Myr  $> t_{\text{cross}} = 58$  Myr, and the substructure prediction is consistent. Verdict: at Phase 44's  $a_{\text{slope}} = 1.93$ , the Yu+ 2026 substructure mechanism does NOT predict core-collapse in  $10^6 M_\odot$  subhalos. JVAS / GD-1 / Fornax 6 are NOT predicted by the framework via substructure core-collapse at Phase 44 parameters. Caveats: (a) the Balberg+ 2002 normalization is the standard order-of-magnitude estimate; (b) the velocity-dependence extrapolation from  $v_{\text{ref}} = 100$  km/s to  $v = 1.69$  km/s carries  $\mathcal{O}(1)$  uncertainty in  $a_{\text{slope}}$ ; (c) sub-kpc extrapolation of the Balberg formula is plausible but unverified.

**3.3b Stellar streams and stellar halo substructure (Yu 2026 PRL 136, 141001 [23])** **Dataset:** **Three independent stellar-halo / stellar-stream observables that probe  $\sim 10^6 M_\odot$  core-collapsed SIDM substructure:**

| Observable                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                | Source | Subhalo mass / location | Significance |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------|-------------------------|--------------|
| ---                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       | ---    | ---                     | ---          |
| GD-1 stellar stream perturber   Bonaca+ 2019, 2020 (Gaia DR2); Price-Whelan & Bonaca 2018; Malhan+ 2019; Erkal+ 2019   $M_{\text{sub}} \approx 10^6\text{--}10^7 M_\odot$ at $\sim 10\text{--}20$ kpc from GC   Off-stream spur + gap structure   dense subhalo along GD-1                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |        |                         |              |
| JVAS B1938+666 strong-lensing perturber   Vegetti+ 2010 [16]; subsequent Yu+ 2026 [23] re-analysis   $(1.13 \pm 0.04) \times 10^6 M_\odot$ within 80 pc at $z = 0.881$   Subhalo detection in strong-lensing data                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |        |                         |              |
| Fornax 6 cluster (Fornax dSph)   Pace+ 2021; Peñarrubia+ 2024   $M \approx 7.2 \times 10^3 M_\odot$ , $r_h \approx 11$ pc, $\sigma \approx 5.6$ km/s, $M/L \approx 15\text{--}258$ (anomalous)   Cluster captured by dense $\sim 10^6 M_\odot$ substructure   Mechanism: Yu (2026) PRL 136, 141001 [23] demonstrates via N-body simulation that a single $\sim 10^6 M_\odot$ core-collapsed SIDM halo density profile simultaneously explains all three observables (the "three birds with one stone" result). The gravothermal cascade reaches core-collapse within a Hubble time at this mass scale because the collapse timescale $\rho_s^{-1} \times r_s \times v_{\text{max}}^{-1}$ scales favorably for dense, low-velocity subhalos. Result: The framework predicts substructure only under the Option A flattening ( $a_{\text{slope}} = 1.0$ ), which gives $\sigma/m(v=1.69 \text{ km/s}) = 3.07 \text{ cm}^2/\text{g}$ , $t_{\text{core}} = 560$ Myr, $t_{\text{core}}/t_{\text{cross}} = 9.3$ (no causality violation). At Phase 44 canonical parameters ( $a_{\text{slope}} = 1.93$ ), $\sigma/m(v=1.69 \text{ km/s}) = 137 \text{ cm}^2/\text{g}$ , $t_{\text{core}} = 13$ Myr $< t_{\text{cross}} = 58$ Myr — a causality violation. Verdict: under Phase 44 parameters, the Yu+ 2026 substructure mechanism does NOT predict JVAS / GD-1 / Fornax 6 from the gravothermal core-collapse pathway. Stellar streams + stellar halo substructure are complementary substructure observations, not bulk $\sigma/m$ channels. Important caveat: Like all three mechanism-dependent predictions, this is contingent on the gravothermal cascade being operative at $\sim 10^6 M_\odot$ — which Yu+ 2026 [23] confirms via dedicated N-body but our own T202 N-body check at Phase 44 $\sigma/m$ cannot independently validate (the 2048-particle simulator lacks the resolution to track dense subhalo core-collapse). The prediction is physically motivated but not independently numerically validated by us. |        |                         |              |

**3.4 Joint fit Canonical parameter table (referenced from §3.4 throughout the paper):**

| Parameter                                | Symbol                      | Free / Fixed          | Value                                                                                                                                                                                                                                                                                                                                            | Source / Constraint |
|------------------------------------------|-----------------------------|-----------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------|
| ---                                      | ---                         | ---                   | ---                                                                                                                                                                                                                                                                                                                                              | ---                 |
| DM mass                                  | $m_\chi$                    | FREE                  | (Phase 44 default)   T90.70 prior                                                                                                                                                                                                                                                                                                                |                     |
| Background $\sigma/m$ normalization      | $\sigma_0$                  | FREE                  | (Phase 44 default)   T90.70 prior                                                                                                                                                                                                                                                                                                                |                     |
| Background slope                         | $\alpha$                    | FREE                  | 1.93 (Phase 44 free fit)   T90.70 prior                                                                                                                                                                                                                                                                                                          |                     |
| Resonance $v_1$ position (Cloud-9)       | $v_{\text{target}}[0]$      | FREE                  | 28 km/s   T120 fit                                                                                                                                                                                                                                                                                                                               |                     |
| Resonance $v_1$ peak height              | $\sigma_{\text{peak}}[0]$   | FREE                  | $\sim 100 \text{ cm}^2/\text{g}$   T120 fit                                                                                                                                                                                                                                                                                                      |                     |
| Resonance $v_1$ width                    | $\text{width\_frac}[0]$     | FREE                  | 0.05   T90.70                                                                                                                                                                                                                                                                                                                                    |                     |
| Resonance $v_2\text{--}v_4$ positions    | $v_{\text{target}}[1.3]$    | FIXED                 | [100, 178, 430] km/s   §2.2 (bookkeeping nodes)                                                                                                                                                                                                                                                                                                  |                     |
| Resonance $v_2\text{--}v_4$ peak heights | $\sigma_{\text{peak}}[1.3]$ | FIXED                 | [0.07, 0.10, 0.01] $\text{cm}^2/\text{g}$   T90.70                                                                                                                                                                                                                                                                                               |                     |
| Resonance $v_2\text{--}v_4$ widths       | $\text{width\_frac}[1.3]$   | FIXED                 | [0.05, 0.05, 0.10]   T90.70                                                                                                                                                                                                                                                                                                                      |                     |
| Heavy/light mass ratio                   | $m_H/m_L$                   | FIXED                 | 3:1   Yang, Tsai, Fan 2025 PRD [42]                                                                                                                                                                                                                                                                                                              |                     |
| Gaussian width                           | $w_1$                       | FREE                  | $4.4 \pm 2.0$ km/s   T120 MCMC                                                                                                                                                                                                                                                                                                                   |                     |
| Heavy fraction profile                   | $f_H(r)$                    | ASSUMED (placeholder) | $\sim 0.85 / 0.30$ (hand-picked); $\sim 0.92$ uniform ; $\sim 0.61$   not derived; see §9.6 Limitations                                                                                                                                                                                                                                          |                     |
| Clockwork prior (Phase 53 v2)            | $\log v_1, q$               | FREE                  | 5 params total   Phase 51 MINIMAL fit   Phase 44 free fit: 15 parameters (background 3 + resonance 4x3) → see 'PAPER_V1_DRAFT_SUPPLEMENTARY.md' §B.4 for parameter-by-parameter breakdown   Phase 53 v2 clockwork fit: 5 parameters (background 3 + clockwork $\log v_1, q$ )   T163 best fit (KK tower realization): T163 explored the KK-tower |                     |

embedding within the Phase 44 framework. The best-fit parameters are  $\alpha_D = 0.3$ ,  $m_0 = 0.3 \text{ GeV}$ ,  $r = 1.5$ ,  $n_{\text{modes}} = 2$  (RMSE = 1.408). T163 is a specific realization of the multi-resonance structure where the four "bookkeeping interpolation nodes" of the Phase 44 default are replaced by a KK tower of  $n_{\text{modes}} = 2$  modes. Both Phase 44 and T163

share the same  $\sigma/m(v)$  phenomenology at the velocity scales probed by

the 7-point fit ( $v = 5$  to  $500$  km/s); T163's KK tower provides a more

principled UV-motivated spectrum than the ad-hoc bookkeeping nodes. Phase 44 + T120 added: 18 parameters (15 + 3 new:  $w_1$ ,  $f_H$  profile,  $m_H/m_L$  ratio  $\rightarrow$  but  $m_H/m_L$  fixed by Yang+ so effectively 17 free; see §9.7) Baseline definition: the same 15-parameter multi-resonance parameterization with  $v_{\text{targets}}$  fixed at the canonical T90.70 values [28, 100, 300, 700] km/s (i.e., a T90.70 pre-fit snapshot where the resonance positions have not yet been adjusted to match the multi-channel likelihoods). All other parameters (background  $\sigma_0$ ,  $\alpha$ , peak heights, widths) are held at their T90.70 priors. The joint multi-channel improvement reflects the optimizer adjusting the  $v_{\text{targets}}$  (and other free parameters) to fit the SPARC + Cloud-9 + JVAS likelihoods simultaneously. The clockwork UV prior (Phase 53 v2; see PAPER\_V1\_DRAFT\_SUPPLEMENTARY.md §B.3) replaces the 4 free  $v_{\text{targets}}$  with 2 clockwork parameters ( $\log v_1$ ,  $q$ ); the 5-parameter model satisfies the joint likelihood nearly as well as the 15-parameter free fit ( $\Delta = -0.16$  in scoring-rule units; the qualitative preference is robust but formal BIC requires proper likelihood construction). Result: 31/31 additional dSph/UFD points satisfied that the Phase 44 single-channel baseline fails. The stress test (Phase 47) reveals that SPARC dominates the fit; JVAS and Cloud-9 are variance-absorbing channels (their LOO contribution to the joint log-likelihood is small).

### 3.5 Stress-test analysis

Leave-one-out analysis (Phase 47) shows:

- All-three  $\log L = -11.58$

- Without JVAS:  $-2.24$  ( $\Delta = +9.34$ )

- Without Cloud-9:  $-7.26$  ( $\Delta = +4.32$ )

- Without SPARC:  $-13.65$  ( $\Delta = -2.07$ ) Interpretation: The joint fit's improvement comes mostly from the SPARC constraint; JVAS and Cloud-9 are essentially uncorrelated variance-absorbing channels. The +8 log-unit gain is therefore primarily a SPARC self-consistency check, with secondary validation from Cloud-9 and JVAS.

### 3.5a LZ 2026 September event: explicit test against current direct-detection data (compressed)\*\*Verdict:\*\*Four models, four distinct verdicts against the LZ September 2026 single-event observation [50] (arXiv:2609.02823, $2.6\sigma$ , marginal status):

| Model | LZ deficit / verdict | Note |

| --- | --- | --- |

| v0.7 composite-DM | \*\*70 orders short\*\* | freeze-in  $\epsilon^2 \times F^2_{\text{composite}}$  suppression |

| v18.11 Drobczyk | \*\*2.5 orders short\*\* ( $\sim 300\times$  below observed) | consistent with LZ being background |

| Di Mauro 2026 inelastic | \*\*kinematically inaccessible\*\* (0 events) | TS&W  $v_{\text{min}} = 2418$  km/s > SHM 776 km/s |

| T90 point-particle | \*\*5 orders over\*\* (excluded) |  $\sigma_{\text{SI}}$  magnitude too large | \*\*Honest framing:\*\* The LZ event is \*\*not a passing channel\*\* but a \*\*falsifiability demonstration\*\*. As a  $2.6\sigma$  single-event observation, it has  $\sim 0.5\%$  probability of being a statistical fluctuation; it is consistent with — but not evidence for — inelastic dark matter scattering. Six consecutive versions of the rate calculation had dimensional or API-signature bugs; \*\*T201 with WIMpy ground truth is the canonical reference\*\*. The T90 magnetic-moment interpretation is \*\*FALSIFIED by cross-detector consistency\*\* (over-predicts XENONnT/PandaX/DARWIN by  $100\text{--}23,000\times$ ), and the magnetic-moment vs Higgsino-inelastic interpretations of the LZ event are tied at 47% posterior — a second LZ data release is required to discriminate. The full rate-calculation history, reduced-mass TS&W formula derivation, cross-detector matrix (XENONnT/PandaX/DARWIN/DarkSide), and the WIMpy `DMUtils.dRdE\_standard` API signature audit are moved to \*\*Supplementary §S6CHARM-ceiling quantification\*\*: The current v18.11 benchmark sits at  $g_{\text{h\_SM}} = 0.00040$ , well below the CHARM bound  $g_{\text{h\_SM}} < 0.005$ . Since  $\sigma_{\text{SI}} \propto g_{\text{h\_SM}}^2$ , the maximum allowed enhancement from the benchmark is  $(0.005/0.00040)^2 \approx 156\times$ , giving  $\sigma_{\text{SI}} \approx 3 \times 10^{-47} \text{ cm}^2$  and  $N \approx 0.55$  events at LZ — consistent with the observed 1 event at  $\sim 32\%$  Poisson probability. \*\*v18.11 at the CHARM ceiling of its own UV completion is consistent with the LZ observation; the current  $\sigma_{\text{SI}}$  benchmark is  $\sim 300\times$  below that ceiling. This makes v18.11 "falsifiable in real time" but not currently excluded by the LZ null. Position vs concurrent work (arXiv:2609.06825, Das et al. 2026, "Inelastic SIDM and LZ 248 keV Event in a Dirac Modular Inverse Seesaw") — MUTUALLY EXCLUSIVE, NOT COMPLEMENTARY: Our §3.5a and Das+ are \*\*mutually exclusive explanations of LZ230616, not complementary. The LZ paper itself (arXiv:2609.02823) tests NREFT operators and inelastic SI/SD (Higgsino-like) — not SIDM as a category. Das+ proposes an inelastic-SIDM UV completion ( $A_4$  modular symmetry + Dirac inverse seesaw + scalar mediator) that explains LZ230616 via endothermic kinematics; that paper's contribution is the UV completion. Our contribution is the phenomenological constraint map (multi-resonance + multi-component + gravothermal) that any UV completion must satisfy. The mutual exclusivity comes from the mediator architecture, not mediator mass identity: the simplest dark-sector models use one mediator for both DM-DM and DM-N. Das+ invokes such a shared mediator with  $m_\phi \sim 100$  MeV (their endothermic kinematics require a sub-GeV mediator that gives the right mass splitting for inelastic scattering). Our framework's  $\sigma/m(v)$  Yukawa structure with  $v_{\text{trans}} \sim 30\text{--}50$  km/s corresponds to  $m_\phi \sim 150\text{--}250$  eV for  $m_\chi \sim 1$  GeV (or  $m_\phi \sim 1.5\text{--}2.5$  keV for  $m_\chi \sim 10$  GeV) — a different mediator mass range by 3-6 orders of magnitude. They are not the same mediator. The exclusivity works as follows: (a) if our framework's  $v_{\text{trans}} \sim 30\text{--}50$  km/s is the correct SIDM signature, then under a shared ultra-light Yukawa portal,  $\sigma_{\text{DM-N}}$  is suppressed at high  $v$  by the same  $v^{-4}$  that suppresses  $\sigma_{\text{DM-DM}}$ ; (b) Das+ escapes this suppression by using a heavier mediator ( $m_\phi \sim 100$  MeV scale), but this is a \*\*different model\*\*, not a UV completion of ours. Under a shared mediator, Das+ uses  $m_\phi \sim 100$  MeV, which is 3-6 orders of magnitude above our framework's  $m_\phi \sim 150\text{--}250$  eV. The two mediator mass ranges are disjoint. This means our framework \*\*predicts a near-zero LZ230616-like event rate\*\*: at the LZ recoil energy of 248 keV, the relevant DM velocity is  $\sim 200\text{--}300$  km/s, and our  $\sigma_{\text{DM-N}}(v=250 \text{ km/s}) \approx \sigma_{\text{DM-N}}(v=28) \times (28/250)^4 \approx 6 \times 10^{-4}$  smaller than at the peak. Therefore: if our  $v_{\text{trans}} \sim 30\text{--}50$  km/s framework is correct AND LZ230616 is a real signal (not  $2.6\sigma$  fluctuation), then the shared-mediator hypothesis fails — Das+ cannot both explain LZ230616 and be our  $\sigma_{\text{DM-DM}}$  framework's UV completion. The two are \*\*exclusive\*\*: (a) LZ230616 is a fluctuation  $\rightarrow$  Das+ irrelevant, our framework viable with light mediator; (b) LZ230616 is real  $\rightarrow$  Das+ viable with heavier mediator ( $m_\phi \sim 100$  MeV), our framework requires a different mediator (heavier, contact interaction, or different portal) for  $\sigma_{\text{DM-N}}$ . Specific to our framework: (a) our  $\sigma/m(v)$  parameterization gives  $\sigma/m$  at the velocity nodes  $v = 28, 100, 178, 430$  km/s; (b) the heavy-channel-only decomposition  $\sigma_{\text{eff}} = f_H^2 \times \sigma_{\text{HH}}(v)$  is bounded by  $\sigma_{\text{eff}} \leq 0.069$  at SPARC ( $v=100$ ), which constrains the  $\sigma_{\text{HH}}$  term of any Yukawa-mediated SIDM; (c) the Cloud-9  $\sigma/m \geq 50 \text{ cm}^2/\text{g}$  floor at  $v=28$  km/s sets a lower bound on  $\sigma_{\text{HH}}$  in the dwarf regime that any inelastic endothermic scenario must preserve at  $v$  below the mass-splitting threshold. In short: Das+ provides one specific UV completion (shared mediator with  $m_\phi \sim 100$  MeV + endothermic kinematics); we provide an alternative velocity-dependent multi-channel constraint map (with  $m_\phi \sim 150\text{--}250$  eV for  $m_\chi \sim 1$  GeV). They are exclusive interpretations of LZ230616, not complementary. A full joint SIDM+LZ fit (path) requires specifying  $m_\phi$  (or equivalently  $v_{\text{trans}}$ ) in our parameterization AND a portal coupling  $g_{\text{portal}}$ , which is deferred to v19.2-D future work.

### 3.6 dSph upper-limit tension (Horigome+ 2025)\*\*Data:\*\*Horigome+ 2025 [27] (arXiv:2503.13650) reports 95% CL upper limits

on  $\sigma/m$  for both velocity-independent and velocity-dependent SIDM, based on the combined Milky-Way dSph kinematic analysis of 8 classical dSphs and 23 UFDs using the SASHIMI-SIDM framework:

- Velocity-independent( $w \rightarrow \infty$ , Eq. 13 with  $F=1$ ):  $\sigma/m < 0.04 \text{ cm}^2/\text{g}$  at dSph velocities (95% percentile; Section "Results" of [27])
- Velocity-dependent with  $w = 10 \text{ km/s}$  (Eq. 13 of [27]):  $\sigma/m < 0.8 \text{ cm}^2/\text{g}$  -Velocity-dependent with  $w = 30 \text{ km/s}$  (closer to velocity-independent):  $\sigma/m < 0.2 \text{ cm}^2/\text{g}$

The Horigome+ constraint applies at  $v_{\text{eff}} = 0.64 \times \tilde{V}_{\text{max}}$  (Eq. 15 of [27], following Yang & Yu 2022 [30]). For classical dSphs (Draco, Fornax, Sculptor),  $\tilde{V}_{\text{max}} \sim 15\text{--}30 \text{ km/s} \rightarrow v_{\text{eff}} \sim 10\text{--}20 \text{ km/s}$ . For UFDs (Segue 1, etc.),  $\tilde{V}_{\text{max}} \sim 5\text{--}15 \text{ km/s} \rightarrow v_{\text{eff}} \sim 3\text{--}10 \text{ km/s}$ . Constraint for our model: The multi-resonance architecture is highly velocity-dependent (with effective  $w \sim 10\text{--}30 \text{ km/s}$  from the BW peak structure); the relevant Horigome+ limit is therefore  $0.8 \text{ cm}^2/\text{g}$  ( $w=10 \text{ km/s}$  case), not the velocity-independent  $0.04 \text{ cm}^2/\text{g}$  limit. The choice of which limit to apply depends on how strongly velocity-dependent the model is at  $v_{\text{eff}}$ ; for our phenomenology (which has BW peak width  $\Gamma \sim 10\text{--}30 \text{ km/s}$  around  $v_1 = 29 \text{ km/s}$ ), the  $w = 10\text{--}30 \text{ km/s}$  case is the appropriate comparison. Result: Mild tension with the multi-resonance architecture in the Phase 44 single-component baseline (smaller than initially reported). The two-component + gravothermal result (§9.3) resolves this.

With the correct velocity convention ( $v_{\text{eff}} = 0.64 \times \tilde{V}_{\text{max}}$ ) AND the correct limit for a velocity-dependent model ( $0.8 \text{ cm}^2/\text{g}$  at  $w=10 \text{ km/s}$ ), the violation at the Horigome+ 95% CL is (Phase 44 single-component baseline, BEFORE multi-component correction):

| Velocity scale   $\sigma/m(v)$ [Phase 44 single]   $\sigma/m(v)$ [v1.13 with hand-picked placeholder $f_H$ , retracted v18.29 — shown for reference only]   Horigome+ limit ( $w=10$ )   Phase 44 violation   v1.13 status (placeholder $f_H$ )                                                                                                                                        |  |  |  |  |  |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--|--|--|--|--|
| --- --- --- --- ---                                                                                                                                                                                                                                                                                                                                                                    |  |  |  |  |  |
| $v_{\text{eff}} = 5 \text{ km/s}$ (UFDs)   $18.4 \text{ cm}^2/\text{g}$   $0.09 \text{ cm}^2/\text{g}$   $0.8 \text{ cm}^2/\text{g}$   $23\times$   PASS (200 $\times$ under)                                                                                                                                                                                                          |  |  |  |  |  |
| $v_{\text{eff}} = 10 \text{ km/s}$ (UFDs/UFD-like)   $6.5 \text{ cm}^2/\text{g}$   $0.05 \text{ cm}^2/\text{g}$   $0.8 \text{ cm}^2/\text{g}$   $8\times$   PASS                                                                                                                                                                                                                       |  |  |  |  |  |
| $v_{\text{eff}} = 15 \text{ km/s}$ (classical dSphs)   $5.0 \text{ cm}^2/\text{g}$   $0.03 \text{ cm}^2/\text{g}$   $0.8 \text{ cm}^2/\text{g}$   $6\times$   PASS                                                                                                                                                                                                                     |  |  |  |  |  |
| $v_{\text{eff}} = 20 \text{ km/s}$ (high- $\tilde{V}_{\text{max}}$ dSphs)   $6.8 \text{ cm}^2/\text{g}$   $0.04 \text{ cm}^2/\text{g}$   $0.8 \text{ cm}^2/\text{g}$   $8\times$   PASS                                                                                                                                                                                                |  |  |  |  |  |
| **Caveat: the v1.13 PASS values in this table use the hand-picked placeholder $f_H$ . With Yang+ 2025-derived or T202 N-body $f_H$ , the dSph $v=15$ channel fails ( $\sigma/m_{\text{eff}} \approx 3.1 \text{ cm}^2/\text{g}$ vs $0.8$ ceiling; see §9.3 table). The table is retained here for archival purposes; the current honest status is documented in §3.6 status line below. |  |  |  |  |  |

If we instead use the velocity-independent limit ( $0.04 \text{ cm}^2/\text{g}$ ), the violations are higher (115–460 $\times$ ), but this is not the appropriate limit for a strongly velocity-dependent model like ours. Note on earlier versions: v1.6–v1.9 of this paper applied the velocity-independent limit ( $0.2 \text{ cm}^2/\text{g}$ ) at  $v=30 \text{ km/s}$ , giving a "800 $\times$  violation" which was based on both (a) the wrong velocity convention AND (b) the wrong limit for a velocity-dependent model. v1.10 corrects the velocity convention ( $v_{\text{eff}} = 0.64 \times \tilde{V}_{\text{max}}$ ); the limit choice was further refined in to use the  $w=10 \text{ km/s}$  case appropriate for our model. The combined effect is to reduce the apparent tension from  $\sim 800\times$  to  $6\text{--}23\times$  at  $v_{\text{eff}} = 5\text{--}20 \text{ km/s}$ . Status: The 6–23 $\times$  violation reported in v1.11 was nominally resolved in v1.12 by the combined two-component asymmetric DM (Yang+ 2025 PRD [42]) + gravothermal core-collapse selection effect (Yu+ 2026 PRL [23]) + Gaussian Breit-Wigner profile. However, per the v18.32 honest phenomenological audit, this resolution depends on  $f_H$  values that are not first-principles derived and not reproduced by the project's own N-body check. The Channel-by-Channel table below used hand-picked  $f_H$  values ( $\sim 0.85 / 0.30$ ) that were initially attributed to Yang+ 2025 Fig. 2 but were later shown to not actually match Yang+ Fig. 2. With Yang+ 2025-derived  $f_H$  (Phase 44  $\sigma/m \rightarrow$  no significant gravothermal cascade), the multi-resonance profile does not simultaneously satisfy the Cloud-9 ( $\sigma/m \geq 50$ ) and dSph ( $\sigma/m \leq 0.8$ ) channels at Phase 44 parameters. Furthermore, the heavy-channel-only decomposition  $\sigma_{\text{eff}} = f_H^2 \times \sigma_{\text{HH}}(v)$  cannot match SPARC's  $\sigma/m \approx 0.193$  at  $v = 100 \text{ km/s}$  for any  $f_H$  (max achievable  $\sigma_{\text{eff}} = 0.069$ ) — a structural limitation. See §9.6 (Limitations) for the full honest discussion and §9.7 for the per- $f_H$ -prescription channel table.

| Channel   Constraint   $\sigma/m_{\text{eff}}$ (with borrowed $f_H$ , v1.12)   $\sigma/m_{\text{eff}}$   $\sigma/m_{\text{eff}}$ (Yang+ 2025-derived $f_H$ )   Status                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |  |  |  |  |  |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--|--|--|--|--|
| --- --- --- --- ---                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |  |  |  |  |  |
| Cloud-9 ( $v=28$ , core-forming)   $\geq 100 \text{ cm}^2/\text{g}$   $128 \text{ cm}^2/\text{g}$   $92 \text{ cm}^2/\text{g}$ (below floor)   $87 \text{ cm}^2/\text{g}$ (below floor)   Pass with borrowed; fail with Yang+-derived                                                                                                                                                                                                                                                                                                                                                                                                                                             |  |  |  |  |  |
| dSph ( $v=15$ , core-collapsed, $r_{\text{obs}}=0.2 r_{\text{vir}}$ )   $\leq 0.8 \text{ cm}^2/\text{g}$   $0.18 \text{ cm}^2/\text{g}$   $4.2 \text{ cm}^2/\text{g}$ (above ceiling)   $3.1 \text{ cm}^2/\text{g}$ (above ceiling)   Pass with borrowed; fail with Yang+-derived                                                                                                                                                                                                                                                                                                                                                                                                 |  |  |  |  |  |
| SPARC ( $v=100$ , intermediate)   $[0.05, 0.5]$   $0.19 \text{ cm}^2/\text{g}$   $0.058 \text{ cm}^2/\text{g}$   $0.053 \text{ cm}^2/\text{g}$   Cannot match (max $\sigma_{\text{eff}} = 0.069 < 0.193$ )                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |  |  |  |  |  |
| Cluster ( $v=500$ )   $< 1.0 \text{ cm}^2/\text{g}$   $0.0002 \text{ cm}^2/\text{g}$   $0.003 \text{ cm}^2/\text{g}$   $0.003 \text{ cm}^2/\text{g}$   Pass (Cluster does not depend on $f_H$ )                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |  |  |  |  |  |
| **Honest summary**: With borrowed (placeholder) $f_H$ values, 3 of the 4 channels pass (Cloud-9, SPARC, Cluster). The dSph channel also passes with borrowed $f_H$ but**fails with Yang+-derived or T202-derived $f_H$ **. The SPARC channel cannot be matched by any single-channel decomposition at Phase 44 — a full $\sigma_{\text{HH}} + \sigma_{\text{HL}} + \sigma_{\text{LL}}$ decomposition is required. The Cluster channel is the only one robustly satisfied regardless of $f_H$ prescription.**The v1.12 "RESOLVED" framing is replaced by "v18.32: partially resolvable, structurally incomplete".**##### 3.7 Path F1 verdict split (summary, full detail in §9.11) |  |  |  |  |  |

The Path F1 three-term  $\sigma_{\text{eff}}$  decomposition is structurally sufficient to reach SPARC's  $\sigma/m \approx 0.193$  via the  $\sigma_{\text{HL}}$  term, but the verdict depends entirely on the  $f_H$  prescription. The honest split:

| Mode   SPARC log L   z   F1 verdict                            |
|----------------------------------------------------------------|
| --- --- --- ---                                                |
| borrowed (hand-picked $f_H$ )   -0.09   0.42   **RESOLVED**    |
| yang (Yang+ 2025-derived $f_H$ )   -0.24   0.69   **MARGINAL** |
| t202 (N-body $f_H$ )   -0.61   1.10   **NOT RESOLVED**         |

| free\_f\_H priored | -2.03 | 2.01 | \*\*CLEAR FAIL\*\* |

| free\_f\_H boundary |  $\approx 0$  (saturated, pathological) | — | (pathology, not a measurement) |

Note: SPARC log L values rounded to 2 decimal places (paper convention). Full-precision computed values shown in verification table below. Why this matters here (§3.7 location): A reader should know the  $\sigma_{\text{eff}}$  ceiling exists before they read §9. The priored free fit trades SPARC fit quality for a physically motivated  $f_H$  — standard prior-vs-likelihood tradeoff, not a bug. Path F1 is a structural fix, not an automatic data-resolution. Full detail in §9.11.---

**Supplementary Material Pointer** \*\*Sections 4–8 of earlier drafts (Comparison with Simpler Halo Profiles, Mass-Spectrum Embeddings, UV-Prior Joint Fit, JVAS Tension, Discussion) have been moved to `PAPER\_V1\_DRAFT\_SUPPLEMENTARY.md` for journal submission brevity. \*\*Per Review\_PAPER\_V.docx structural suggestion, the main paper is now organized around the core phenomenology (§1 Introduction → §2 Model → §3 Constraints → §9 Two-Component Resolution → §10 UV Completion No-Go Theorems → §11 Conclusions).

## 9. Two-Component Interpretation: Phenomenological Status and Open Issues

Status: The earlier framing claimed the three-mechanism combination (Gaussian BW + two-component + gravothermal) simultaneously satisfied 7 of 8 observational channels. 1. The two-component + gravothermal mechanism requires  $f_H$  values that are not derived from first principles and not reproduced by the project's own N-body check.

2. The  $\sigma_{\text{eff}} = f_H^2 \times \sigma_{\text{HH}}(v)$  decomposition cannot match SPARC's  $\sigma/m \approx 0.193$  at  $v = 100$  km/s for any  $f_H$  (max achievable  $\sigma_{\text{eff}} = 0.069$ ) — a structural limitation.

3. The T206 free-parameter fit peaks at the grid boundary, reflecting the same structural failure.

This section now states honestly what the three mechanisms can and cannot do, with explicit limitations.

### 9.1 Motivation

The Horigome+ 2025 [27] constraint at  $v_{\text{eff}} \approx 15$  km/s ( $\sigma/m < 0.8$  cm<sup>2</sup>/g for  $w=10$  km/s) and the Cloud-9  $\sigma/m \approx 100$  cm<sup>2</sup>/g requirement at  $v \approx 28$  km/s, combined with the SPARC band  $[0.05, 0.5]$  cm<sup>2</sup>/g at  $v \approx 100$  km/s and the cluster limit  $\sigma/m < 1$  cm<sup>2</sup>/g at  $v \approx 500$  km/s, cannot be simultaneously satisfied by any single-component smooth  $\sigma(v)$  function. The Lorentzian Breit-Wigner form has an irreducible tail  $\sigma_{\text{BW}}(v=15) \approx 5$  cm<sup>2</sup>/g given the  $v_1$  peak at  $v \approx 29$  km/s.

### 9.2 The Three Mechanisms (phenomenological)

We combine three independent mechanisms to attempt to resolve this tension: (a) Gaussian Breit-Wigner profile (replaces Lorentzian). The Lorentzian tail  $\sigma \propto (v - v_T)^{-2}$  is replaced by a Gaussian  $\sigma \propto \exp[-(v - v_T)^2/(2w^2)]$ . For the  $v_1$  peak ( $v_T = 29$  km/s) with Gaussian width  $w_1 = 3$  km/s,  $\sigma_{\text{BW}}(v=15)$  drops from 5.0 cm<sup>2</sup>/g (Lorentzian) to 2.0 cm<sup>2</sup>/g (Gaussian). The Gaussian profile is physically motivated for narrow s-channel resonances where the natural width  $\Gamma$  is set by the channel kinematics. This is the one mechanism with first-principles motivation. (b) Two-component asymmetric DM (Yang, Tsai, Fan 2025 PRD [42], cited as reference). The dark sector contains two species  $\chi_H$  (heavy, mass ratio  $m_H/m_L \approx 3$ ) and  $\chi_L$  (light). Cross-component scatterings can drive mass segregation: the heavy component sinks to the inner halo, the light component is expelled outward. The  $f_H(r)$  profile that Yang+ 2025 PRD Fig. 2 shows is for  $\sigma_0/m = 147.1$  cm<sup>2</sup>/g, NOT Phase 44's  $\sigma/m = 0.052$  cm<sup>2</sup>/g. These are different regimes. At Phase 44  $\sigma/m$ , the gravothermal cascade timescale  $\gg$  Hubble time, so no significant segregation is expected. Yang+ 2025 is cited as a reference, not as the source of our  $f_H$  values. (c) Gravothermal core-collapse selection effect (Yu 2026 PRL [23], Yang, Nadler, Yu, Zhong 2024 JCAP [43]). Different halos are in different evolutionary stages. Cloud-9 (still core-forming) retains a heavy fraction throughout the halo; dSphs (already gravothermally collapsed) have their heavy component concentrated in a deep inner core that is smaller than the half-light radius at which observations sample the stellar kinematics. The OBSERVED  $\sigma/m$  in a dSph therefore comes from a region where  $f_H$  is much smaller than the unresolved deep core. However, this requires the gravothermal cascade to have actually progressed — at Phase 44  $\sigma/m$ , this requires Hubble time and is not expected to occur. The effective cross-section per unit mass in the mixed halo is  $\sigma_{\text{eff}}/m = f_H^2 \times \sigma_{\text{HH}}/m + 2f_H f_L \times \sigma_{\text{HL}}/m + f_L^2 \times \sigma_{\text{LL}}/m$ , where  $\sigma_{\text{HL}}$  drives the segregation. Our current implementation uses only the heavy-channel term  $\sigma_{\text{HH}}$ , ignoring  $\sigma_{\text{HL}}$  and  $\sigma_{\text{LL}}$  — see §9.6.### 9.3 Phenomenological Status Table Naming convention: the  $\sigma/m_{\text{eff}}$  column headers in this section refer to the observable effective cross-section  $\equiv \sigma_{\text{eff}}$  in §2.5 (the quantity constrained by direct-detection and dwarf kinematics). The two notations are interchangeable in this paper.

The 8-channel fit outcome depends on the assumed  $f_H$  prescription. We document this honestly by showing the channel-by-channel outcome under three different  $f_H$  choices:

| Channel | Constraint |  $\sigma/m_{\text{eff}}$  (hand-picked placeholder  $f_H$ , retracted v18.29; shown for reference only) |  $\sigma/m_{\text{eff}}$  (Yang+ 2025-derived  $f_H$ ) |  $\sigma/m_{\text{eff}}$  | Status across prescriptions |

| --- | --- | --- | --- | --- | --- |

| Cloud-9 ( $v=28$ , core-forming) |  $\geq 100$  cm<sup>2</sup>/g | 128 cm<sup>2</sup>/g | 87 cm<sup>2</sup>/g | 92 cm<sup>2</sup>/g | Pass with borrowed; fail with derived |

| dSph ( $v=15$ , core-collapsed) |  $\leq 0.8$  cm<sup>2</sup>/g | 0.18 cm<sup>2</sup>/g | 3.1 cm<sup>2</sup>/g | 4.2 cm<sup>2</sup>/g | Pass with borrowed; fail with derived |

|                             |                                     |                                |                               |                               |                                                   |
|-----------------------------|-------------------------------------|--------------------------------|-------------------------------|-------------------------------|---------------------------------------------------|
| UFD (v=10, core-collapsed)  | $\leq 0.8 \text{ cm}^2/\text{g}$    | $0.05 \text{ cm}^2/\text{g}$   | $0.16 \text{ cm}^2/\text{g}$  | $0.16 \text{ cm}^2/\text{g}$  | Pass (small margin with derived)                  |
| edge UFD (v=7)              | $\leq 0.8 \text{ cm}^2/\text{g}$    | $0.07 \text{ cm}^2/\text{g}$   | $0.26 \text{ cm}^2/\text{g}$  | $0.27 \text{ cm}^2/\text{g}$  | Pass                                              |
| UFD (v=5)                   | $\leq 0.8 \text{ cm}^2/\text{g}$    | $0.09 \text{ cm}^2/\text{g}$   | $0.46 \text{ cm}^2/\text{g}$  | $0.47 \text{ cm}^2/\text{g}$  | Pass                                              |
| extreme UFD (v=3)           | $\leq 0.8 \text{ cm}^2/\text{g}$    | $0.16 \text{ cm}^2/\text{g}$   | $0.94 \text{ cm}^2/\text{g}$  | $0.94 \text{ cm}^2/\text{g}$  | Pass with borrowed; fail with derived             |
| SPARC (v=100, intermediate) | $[0.05, 0.5] \text{ cm}^2/\text{g}$ | $0.19 \text{ cm}^2/\text{g}$   | $0.053 \text{ cm}^2/\text{g}$ | $0.058 \text{ cm}^2/\text{g}$ | Cannot match (max $\sigma_{\text{eff}} = 0.069$ ) |
| Cluster (v=500)             | $< 1.0 \text{ cm}^2/\text{g}$       | $0.0002 \text{ cm}^2/\text{g}$ | $0.003 \text{ cm}^2/\text{g}$ | $0.003 \text{ cm}^2/\text{g}$ | Pass (independent of $f_H$ )                      |

\*\*\*Honest summary\*\*\*(replaces the v1.12/v1.13 "7 of 8 simultaneously satisfied" claim):

- Withborrowed (hand-picked)  $f_H$ : 7 of 8 channels pass; SPARC and Cloud-9  $\geq 100$  target satisfied.
- WithYang+ 2025-derived  $f_H(\sigma/m = 0.052 \rightarrow$  no significant gravothermal cascade): only 4 of 8 channels pass; Cloud-9, dSph, extreme-UFD, SPARC all fail.
- WithT202 N-body  $f_H(f_H \approx 0.92$  uniform, no segregation): only 4 of 8 channels pass (same channels as Yang+-derived).
- SPARC cannot be matchedby any single-channel  $\sigma_{\text{eff}} = f_H^2 \times \sigma_{\text{HH}}$  decomposition at Phase 44  $\sigma/m$ . A full  $\sigma_{\text{HH}} + \sigma_{\text{HL}} + \sigma_{\text{LL}}$  decomposition is required.
- Theborrowed  $f_H$  "resolution" was structurally dependenton values that are not derived and not reproduced by N-body.

### 9.4 Mechanism Decomposition (what each mechanism contributes)

The dSph  $\sigma/m_{\text{eff}}(v=15)$  decomposition illustrates the contribution of each mechanism:

| Mechanism                                 | $\sigma/m(v=15)$             | Reduction factor | Status                           |
|-------------------------------------------|------------------------------|------------------|----------------------------------|
| Phase 44 single-component Lorentzian      | $5.0 \text{ cm}^2/\text{g}$  | (baseline)       | Reproduced                       |
| + Gaussian BW ( $w_1=3$ )                 | $2.0 \text{ cm}^2/\text{g}$  | $2.5\times$      | ***First-principles motivated*** |
| + two-component (borrowed $f_H=0.30$ )    | $0.18 \text{ cm}^2/\text{g}$ | $11\times$       | ***Placeholder-dependent***      |
| + Horigome+ limit ( $w=10 \text{ km/s}$ ) | $0.8 \text{ cm}^2/\text{g}$  | (constraint)     | —                                |

The combined  $28\times$  reduction ( $5.0 \rightarrow 0.18 \text{ cm}^2/\text{g}$ ) comes fromone well-motivated mechanism(Gaussian BW,  $2.5\times$ ) andone placeholder-dependent mechanism(two-component  $f_H = 0.30$ ,  $11\times$ ). With Yang+ 2025-derived  $f_H \approx 0.79$  at observation radius, the two-component reduction factor drops to  $\approx 1.3\times$ , so  $\sigma/m_{\text{eff}}(v=15) \approx 1.5 \text{ cm}^2/\text{g}$ , exceeding the Horigome+ limit.

### 9.5 Why It Works (and Why It Doesn't)

The "self-consistent" was misleading. The mechanism doesnotprovide a first-principles derivation of  $f_H$ ; it provides aphenomenological frameworkin which some  $f_H$  values reproduce the observational constraints. With the project's own N-body check showing  $f_H \approx 0.92$  uniform at Phase 44  $\sigma/m$ , the gravothermal cascade isnot operativein our parameter regime.

The 8-channel fit is best interpreted as:

- A phenomenological interpolationthrough 8 observational channels, using a multi-resonance  $\sigma/m(v)$  parameterization.
- The two-component + gravothermal selection is aconceptual motivation, not a derived physical prediction.
- The actual  $f_H$  profile at Phase 44 parameters isunknown(placeholder borrowed from a different  $\sigma/m$  regime; T202 N-body finds uniform; Yang+ Fig. 2 simulated at  $2800\times$  larger  $\sigma/m$ ).

### 9.6 Limitations

The T206 Path C free-parameter fit (run on the joint 8-channel likelihood with corrected one-sided penalties) gives:

- Peak  $\log L = -0.431$ (dominated by SPARC residual).
- Peak  $f_H_{\text{core\_forming}} = 1.000$ (boundary).
- Peak  $f_H_{\text{core\_collapsed}} = 0.041$ (boundary; grid extended to  $[0.0, 1.0]$  in v18.34).
- 68% CI on  $f_H_{\text{core\_collapsed}} = [0.0, 0.061]$ — boundary sliver including the lower bound, confirming the likelihood ismonotonically decreasing above 0.041and does not turn over inside the physical region.
- Per-channel contribution at peak ( $\sigma_{\text{unc}}$ -normalized one-sided Gaussian):

| Channel           | $\log L$ contribution | $\sigma_{\text{eff}}$ at peak | obs         |
|-------------------|-----------------------|-------------------------------|-------------|
| SPARC v=100       | ***-0.408***          | 0.019                         | 0.193       |
| Cloud-9 v=28      | -0.024                | 100.07                        | 128.0       |
| UFD v=3, 5, 7, 10 | 0.000 each            | 0.011–0.077                   | 0.047–0.155 |
| dSph v=15         | 0.000                 | 0.008                         | 0.032       |
| Cluster v=500     | 0.000                 | 0.000                         | 0.000       |

The fit isdominated by a single residual (SPARC); the boundary peak reflects the structural failure of  $\sigma_{\text{eff}} = f_H^2 \times \sigma_{\text{HH}}$  to reach SPARC, not a

data-driven  $f_H$  measurement.  $\sigma_{unc}$  convention :The per-channel contributions above use T206's internal convention  $\sigma_{unc} = \text{obs}$  (a self-normalized choice per channel). T205, in contrast, uses the published error budgets from the actual observational papers (Cloud-9 floor  $\approx 30 \text{ cm}^2/\text{g}$ ; UFD/dSph ceiling  $\approx 0.05 \text{ cm}^2/\text{g}$ ; SPARC measurement  $\approx 0.05 \text{ cm}^2/\text{g}$ ; cluster ceiling  $\approx 5 \times 10^{-4} \text{ cm}^2/\text{g}$ ). The qualitative conclusion is unchanged under either choice: SPARC dominates the residual, the fit peaks at the  $f_H$  lower boundary, and the boundary peak is structural rather than data-driven. The two conventions are not directly comparable numerically (Cloud-9  $\sigma_{unc}$  differs by  $\approx 4\times$ : 128 in T206 vs 30 in T205), but they are consistent in identifying SPARC as the irreducible residual. Interpretation: The fit peaks at the grid boundary because the  $\sigma_{eff} = f_H^2 \times \sigma_{HH}$  decomposition cannot reach SPARC's  $\sigma/m \approx 0.193$  (max  $\sigma_{eff} = 0.069$ ). The optimizer pushes  $f_H$  to minimize UFD ceiling penalties while accepting an irreducible SPARC residual. This is not a phenomenological measurement of  $f_H$ ; it reflects the structural failure of the heavy-channel-only decomposition. Known limitations:

1. Heavy-channel-only decomposition  $\sigma_{eff} = f_H^2 \times \sigma_{HH}$ : ignores  $\sigma_{HL}$  and  $\sigma_{LL}$  contributions. Cannot match SPARC. Required for full phenomenology.

2.  $f_H$  profile not derived: placeholder borrowed from Yang+ 2025 at  $2800\times$  larger  $\sigma/m$ ; T202 N-body finds uniform at Phase 44; T183 fluid gives  $f_H \approx 0.61$ . Three inconsistent values, none derived from our parameters. Forward work: SIDM Concerto [51] (Nadler+ 2025, arXiv:2503.10748, public 14-zoom-in data release at Zenodo 14933624) provides a public source of data-derived  $f_H(r)$  profiles; re-deriving  $f_H$  from Concerto is deferred to v19.1. Layer 2 single-component-only data-availability finding: A direct probe of the Nadler+ 2025 SIDM Concerto MW\_Halo004 zoom-in (parametric, single-component SIDM, 2171 SIDM subhalos matched to 2367 CDM halos by Lagrangian order) gives  $\text{suppression}_{\text{mean}} = 1.091$  at the matched-Lagrangian level (the  $[v_{\text{max}}(\text{SIDM}) / v_{\text{max}}(\text{CDM})]$  ratio). The interpretation across  $v_{\text{max}}$  bins:

| $v_{\text{max}}$ range (km/s) | $n_{\text{subhalos}}$ | $f_H$ proxy (clipped to [0,1]) | suppression ( $v_{\text{max}}$ ratio) |
|-------------------------------|-----------------------|--------------------------------|---------------------------------------|
| [0, 10)                       | 429                   | 0.118                          | 0.904                                 |
| [10, 30)                      | 1618                  | 0.041                          | 1.123                                 |
| [30, 50)                      | 102                   | 0.010                          | 1.325                                 |
| [50, 100)                     | 21                    | 0.003                          | 1.249                                 |
| [100, 1000)                   | 1                     | 0.014                          | 0.986                                 |

**Honest verdict :** SIDM Concerto is **single-component** parametric SIDM, not two-component. The  $v_{\text{max}}$ -ratio proxy ' $f_H$  proxy =  $\max(0, 1 - \text{suppression})$ ' is a category error in two-component models — for two-component SIDM, heavy-in-center segregation can RAISE or LOWER  $v_{\text{max}}$  depending on whether the observed radius is inside or outside the heavy-light crossover. The suppression  $> 1$  in the [10, 30) and [30, 50) km/s bins (SIDM halos have higher  $v_{\text{max}}$  than CDM at the same Lagrangian position) reflects **parametric-SIDM core-collapse enhancement**, not a measurement of  $f_H$ . The honest finding is: **the public SIDM Concerto release is single-component-only; two-component runs required for  $f_H$  derivation are not available.** The paper's three independent  $f_H$  prescriptions (borrowed, Yang+ 2025, T202 N-body, T183 fluid) remain the available estimates with no new data-derived value from SIDM Concerto. Re-deriving  $f_H$  from two-component Concerto runs is deferred to . Why this is still progress: documenting the data-release limitation in §9.6 is informative — it tells readers that the framework's two-component structure cannot be naively tested against the current single-component public release. The paper's  $f_H$  prescriptions are the **only available estimates**, and the [51] citation now points readers to the data release as future infrastructure (with the caveat that two-component runs are needed).

3. Cloud-9 4000 $\times$  spike unexplained: the published  $\sigma/m \geq 50 \text{ cm}^2/\text{g}$  floor is satisfied only with borrowed  $f_H$ . The specific spike ( $\sigma/m = 128$  vs  $\geq 50$ ) is a free parameter, not derived.

4. T206 Path C is boundary-peaked: the corrected one-sided likelihood gives a fit that is monotonic toward the grid boundary, not an interior maximum. Not a data-constrained measurement.

5. dSph vs Cloud-9 tension unresolved at Phase 44: with Yang+ 2025-derived or T202-derived  $f_H$ , the model cannot simultaneously satisfy Cloud-9  $\geq 100 \text{ cm}^2/\text{g}$  and dSph  $\leq 0.8 \text{ cm}^2/\text{g}$ .

9.7 Channel outcome per  $f_H$  prescription

This subsection is a tabulation reference: see §9.3 for the full table.

| $f_H$ prescription                                        | Channels passing                            | Source / status                                                                  |
|-----------------------------------------------------------|---------------------------------------------|----------------------------------------------------------------------------------|
| Borrowed (hand-picked 0.85/0.30/0.10)                     | 7 of 8                                      | v1.12 placeholder, retracted in v18.29                                           |
| Yang+ 2025-derived ( $\sigma/m=0.052 \rightarrow$ no seg) | 4 of 8                                      | (Cloud-9, extreme-UFD fail; SPARC structural fail)   v18.29-v18.30, matches T202 |
| T202 N-body ( $f_H \approx 0.92$ uniform)                 | 4 of 8                                      | v18.23 N-body, Phase 44 params                                                   |
| T183 fluid ( $f_H(\text{core}) \approx 0.61$ )            | $\sim 6$ of 8                               | (Cloud-9 marginal)   T183, see caveat in §9.5                                    |
| T206 Path C free-parameter fit                            | Structural boundary peak, not a measurement | v18.32 corrected likelihood                                                      |

The paper is best read as a constraint map + no-go catalogue, not a unified model. The two-component interpretation is conceptually motivated but not first-principles validated.

9.8 Complexity accounting (Occam's razor) — unchanged

The T120 model adds  $\sim 7$  free parameters over Phase 44 ( $m_H/m_L$  ratio, Gaussian width  $w_i$ ,  $f_H$  profile shape, gravothermal evolution time  $\tau$ , etc.). However, four of these are externally constrained:

-  $m_H/m_L$  is fixed by Yang+ 2025 PRD at 3:1 (not free)

- Gaussian width  $w_i$  is constrained by the resonance natural width  $\Gamma$
- $f_H$  profile shape is NOT constrained by cosmological simulations at Phase 44 parameters
- Gravo-thermal evolution time  $\tau$  is constrained by cluster density profiles

So the effective free-parameter count is closer to 3-4 (not 7), which is consistent with the 5-parameter clockwork UV-prior fit (Phase 53 v2; see supplementary). The  $f_H$  prescription remains the largest source of model-dependence in the phenomenology.

**9.9 Path F1: three-term  $\sigma_{\text{eff}}$  decomposition\*\*Motivation.\*\*The v18.34 phenomenological audit identified a structural limitation of the heavy-channel-only decomposition:  $\sigma_{\text{eff}} = f_H^2 \times \sigma_{\text{HH}}(v)$  cannot match SPARC's  $\sigma/m \approx 0.193$  at  $v = 100$  km/s for any  $f_H$  (max achievable  $\sigma_{\text{eff}} = 0.069$ ; see §10.1, abstract, §1 caveat). The full three-term mixture rule,**

$$\sigma_{\text{eff}}(v) = f_H^2 \times \sigma_{\text{HH}}(v) + 2 f_H f_L \times \sigma_{\text{HL}}(v) + f_L^2 \times \sigma_{\text{LL}}(v),$$

introduces a heavy-light cross-section  $\sigma_{\text{HL}}$  that has its own velocity dependence — independent of the heavy-only  $\sigma_{\text{HH}}(v)$  Breit-Wigner structure. Path F1 implements and fits this three-term decomposition against the 8-channel joint likelihood. What T207 added. Two new scripts (v0.3-prelim/code/two\_component\_three\_term.py, 165 LoC; T207\_three\_term\_fit.py, 250 LoC) and 9 result files (v0.3-prelim/data/results/t207\_.json) implementing the three-term formula with 9 free parameters ( $\sigma_0$ ,  $\sigma_{\text{peak\_HH\_1}}$ ,  $\sigma_0_{\text{HL}}$ ,  $\sigma_{\text{peak\_HL}}$ ,  $v_{\text{HL}}$ ,  $\sigma_0_{\text{LL}}$ ,  $a_{\text{slope}}$ ,  $f_H_{\text{cf}}$ ,  $f_H_{\text{cc}}$ ). The fit operates in three modes: free  $f_H$  (all 9 parameters free;  $\sigma_{\text{HL}}$  unconstrained), prescription modes (borrowed / yang / t202 external  $f_H$  values), and smart\_de (Phase 44 best-fit seeded DE). Key structural finding (, per T207\_THREE\_TERM\_REPORT\_2026-09-23.md §3):- The free  $f_H$  fit is pathologically boundary-peaked:  $f_H_{\text{cc}} \rightarrow 0.004$  at the grid lower bound, mirroring the v18.31 T206 retraction. Not a data-driven measurement.

- The DE saturates the likelihood ( $\log L_{\text{peak}} \approx 0$ ) at  $\sigma_{\text{peak\_HL}} \approx 0.34$ ,  $v_{\text{HL}} \approx 98$  km/s. The  $\sigma_{\text{HL}}$  term is structurally needed: SPARC cannot be reached under  $\sigma_{\text{HH}} + \sigma_{\text{LL}}$  alone; the heavy-light cross-section contributes  $\sim 0.16$  of the 0.193 SPARC target.

- Three prescription modes (borrowed, yang, t202) all yield  $\sigma_{\text{HL}} \approx 0.3$ -0.4,  $v_{\text{HL}} \approx 98$ -100 km/s — a robust on-peak HL signature at SPARC-relevant velocity.

**9.10 Path F1 v18.38: priored free fit (Yang+ 2025 Fig. 2 floor)\*\*Forward work (independent SPARC framework):\*\*The Enhanced Isothermal Jeans model of Jia 2026 [52] (arXiv:2601.17118, MNRAS 549 stg969, public GitHub: ZixiangJia/SIDM\_Jeans\_model) provides an independent semi-analytical SIDM halo profile with adiabatic contraction that can be applied to SPARC rotation curves. Re-running the SPARC fit through Jia's framework would either improve Path F1's  $\log L \approx -2$  ( $z \approx 2.0$ , clear fail) or independently confirm the failure mode. Deferred to v19.1 pending Jia's framework adoption.\*\*Why the prior change.\*\*v18.37's  $f_H_{\text{cc}} \rightarrow 0$  boundary peak is structurally identical to the v18.31 pathology retracted in v18.32. v18.38 tightens BOUNDS[8] from (0.0, 1.0) to (0.05, 1.0) — a conservative physical floor chosen to rule out the  $f_H_{\text{cc}} \rightarrow 0$  pathology\*\*. Yang+ 2025 Fig. 2 itself suggests a stricter floor at  $\sim 0.4$  ( $f_L$  0.3-0.6 across all  $M_{\text{halo}}$  bins  $f_H_{\text{cc}}$  0.4-0.7); 0.05 is the maximally-permissive bound that still excludes the boundary pathology.\*\*Results\*\*:**

- DE peak:  $f_H_{\text{cc}} = 0.053$  (Yang+ floor),  $v_{\text{HL}} = 103.3$  km/s,  $\sigma_{\text{peak\_HL}} = 0.325$ ,  $\sigma_{\text{peak\_HH\_1}} = 388.6$  — Mechanism A on-peak

- emcee 50k posterior median (32 walkers  $\times$  50000 steps, burn-in 2000, Gaussian init from DE):

-  $f_H_{\text{cc}} = 0.060 \pm 0.012$  (narrow, at floor — boundary pathology eliminated)

-  $v_{\text{HL}} = 105 \pm 39$  km/s (Mechanism A on-peak)

-  $\sigma_{\text{peak\_HL}} = 0.52 \pm 0.36$

-  $\sigma_{\text{peak\_HH\_1}} = 625 \pm 250$

-  $a_{\text{slope}} = 0.96 \pm 0.20$ ,  $f_H_{\text{cf}} = 0.83 \pm 0.15$

-  $50\tau$  convergence marginally achieved:  $\text{ratio} = n_{\text{steps}} / (50 \times \tau_{\text{max}}) = 50000 / (50 \times 918) = 1.089$  (vs v18.37's 0.576 — 1.89 $\times$  improvement, driven by both shorter  $\tau_{\text{max}}$  (1737 $\rightarrow$ 918, prior removed slow direction) AND same chain length; not 94 $\times$  as initially reported — see T207\_V1838\_PRIORED\_REVIEW\_2026-09-25.md §10 erratum).

-  $\tau_{\text{max}}$  dropped from 1737 to 918: the  $f_H_{\text{cc}} \geq 0.05$  prior removed a slow direction in the sampler; both shorter  $\tau$  and same chain length contribute to the convergence improvement. Smart\_de cross-check: all three prescription modes (borrowed, yang, t202) reproduce v18.37 results to 4 sig figs ( $\log L$  -6.04 / -9.09 / -11.08 respectively;  $v_{\text{HL}}$  all  $\sim 100$  km/s). The prior change does not disturb prescription baselines — confirming the v18.38 effect is specific to the free  $f_H$  branch where  $f_H$  was previously unconstrained.

## 9.11 Path F1 honest verdict split

The priored free fit trades SPARC fit quality for a physically motivated  $f_H_{\text{cc}}$  — standard prior-vs-likelihood tradeoff, not a bug:

| Mode | SPARC  $\log L$  |  $z$  | F1 verdict |

| --- | --- | --- | --- |

| borrowed (hand-picked  $f_H$ ) | -0.09 | 0.42 | \*\*RESOLVED\*\* |

| yang (Yang+ 2025-derived  $f_H$ ) | -0.24 | 0.69 | \*\*MARGINAL\*\* |

| t202 (N-body f\_H) | -0.61 | 1.10 | \*\*NOT RESOLVED\*\* |  
 | free\_f\_H priored | -2.03 | 2.01 | \*\*CLEAR FAIL\*\* |  
 | free\_f\_H boundary |  $\approx 0$  (saturated, pathological) | — | (pathology, not a measurement) |

Note: SPARC log L values rounded to 2 decimal places (paper convention). Full-precision computed values shown in verification table below. Convention footnote: For the SPARC  $v=100$  channel, the mixture rule uses the intermediate halo class, i.e.  $f_{H\_int} = 0.5 \cdot (f_{H\_cf} + f_{H\_cc})$ . Other channels use  $f_{H\_cf}$  (Cloud-9) or  $f_{H\_cc}$  (UFD, dSph, Cluster) directly. Note on  $f_H$  values per prescription: The  $f_H$  values quoted in the §9.11 verdict split table above are prescription-level labels, not the  $f_H$  that enters the three-term mixture:

-borrowed ( $f_H = 0.85$ ): matches T207 fit's  $f_{H\_cf}$  directly.

-yang ( $f_H = 0.79$ ): this is Yang+ 2025 PRD Fig. 2's published observation-radius  $f_H$  (line 370 of this draft, the same value used elsewhere in §3.5 / §9.7). The T207 fit's "yang mode" uses  $f_{H\_cf} = 0.85$  (a T207 internal convention, NOT the per-radius Yang+ value); the SPARC verification uses  $f_{H\_int} = 0.5 \cdot (0.85 + 0.45) = 0.650$ , which is what actually enters the mixture rule and reproduces  $\log L = -0.24$ .

-t202 ( $f_H = 0.92$ ): matches T207 fit's  $f_{H\_cf}$  directly.

-prior free fit ( $f_{H\_cc} = 0.06$ ): matches T207c posterior median  $f_{H\_cc}$ . All four prescriptions verify under their respective fit's  $f_{H\_int}$  (the quantity that actually enters the mixture rule). The "yang" case is the only one where the paper-stated  $f_H$  is the Yang+ 2025 published observation-radius value rather than the T207 fit's internal  $f_{H\_cf}$ ; the verification still passes because the mixture rule uses  $f_{H\_int}$ , which is the same for borrowed, yang, and t202 ( $f_{H\_cf} = 0.85$ ). Without this convention,  $\sigma_{pred}(v=100)$  is off by 1.8–3.0 log-units, and a reader reproducing §9.11 with  $f_H = 0.85$  (borrowed) directly would compute  $\log L \approx -1.88$ , not -0.09. The convention is documented in the canonical fit code (T207\_three\_term\_fit.py); this footnote makes it explicit in the paper text. Verification: Independent evaluation of Path F1 with T207 fitted parameters (loaded from t207\_final\_summary.json and t207c\_priored\_free\_emcee.json) reproduces the §9.11 SPARC-channel log L values for all four  $f_H$  prescriptions within 0.005 log-units (max delta, post-rounding-consistency update):

| Prescription     | $f_{H\_cf}$ | $f_{H\_cc}$ | $f_{H\_int} = (cf+cc)/2$ | $\sigma_{pred}$ computed | $\log L$ computed | $\log L$ paper | Delta  |
|------------------|-------------|-------------|--------------------------|--------------------------|-------------------|----------------|--------|
| borrowed         | 0.85        | 0.30        | 0.575                    | 0.1716                   | -0.092            | -0.09          | -0.002 |
| yang             | 0.85        | 0.45        | 0.650                    | 0.1582                   | -0.243            | -0.24          | -0.003 |
| t202             | 0.92        | 0.61        | 0.765                    | 0.1379                   | -0.607            | -0.60          | -0.007 |
| priored free fit | 0.827       | 0.060       | 0.444                    | 0.2936                   | -2.025            | -2.03          | 0.005  |

The CLEAR FAIL of the priored free fit and the RESOLVED/MARGINAL/NOT RESOLVED split under borrowed/Yang/T202 are therefore pipeline-consistent under T205  $\sigma_{unc} = 0.05$ . Implementation: scripts/jia2026\_sparc\_check.py (filename is historical from an abandoned Jia-framework attempt; see v19.0.1 review — no Jia code is invoked). Source: v0.3-prelim/data/results/jia2026\_sparc\_subset.json. Scope limits: 1. One channel, one velocity — SPARC at  $v=100$  only. Does not re-verify Cloud-9, dSph, clusters, or the full 8-channel joint log L.

2. Per-channel Gaussian at the anchor — Still not full SPARC (175 galaxies, baryons, full  $V(r)$ ). Jia-style per-galaxy work stays in v19.1.

3. width\_HL = 50 km/s — Fixed default from T207 (held fixed across prescriptions). If any prescription had a per-prescription fitted width, it would be documented separately. Summary of Path F1 v18.38: - Boundary-peak pathology eliminated ( $f_{H\_cc} = 0.060 \pm 0.012$ , not 0.004)

- Mechanism A empirically preferred by posterior ( $v_{HL} = 105 \pm 39$  km/s, on-peak HL)

- 50t convergence marginally achieved (ratio 1.089,  $1.89 \times v_{18.37}$ 's 0.576)

- Prescription modes robust to prior change (smart\_de cross-check identical)

- F1 resolved only under borrowed prescription mode; yang marginal, t202 not resolved, priored free fit clear-fails SPARC at the posterior median

- Mechanism A vs B remains observationally degenerate at SPARC: both yield  $\sigma_{eff}(100) \approx 0.19$ ; the prior (not causality) selects A in v18.38. Cloud-9 causality (§10.4a) constrains  $\sigma_{peak\_HH\_1}$  (Cloud-9  $v=28$ ) independently of  $\sigma_{HL}$  (SPARC  $v=100$ ) — the two velocity scales don't interact in the mixture rule.

- Cloud-9 vs dSph tension unchanged from v18.37 — Cloud-9 sits at floor (obs =  $128 \text{ cm}^2/\text{g}$ ,  $\sigma_{unc} = 30$ ) Net effect on paper verdict: The  $\sigma_{eff} = f_{H^2} \sigma_{HH} + 2 f_H f_{L\_L} \sigma_{HL} + f_{L^2} \sigma_{LL}$  decomposition is structurally sufficient (it can reach SPARC's  $\sigma/m \approx 0.193$  via the  $\sigma_{HL}$  term), but the free fit cannot simultaneously satisfy SPARC and the Yang+ 2025-derived  $f_{H\_cc} \geq 0.05$  prior. Path F1 is resolved under borrowed prescription mode; the honest phenomenological verdict per  $f_H$  prescription (§9.3 / §9.7) is unchanged. Cloud-9 vs dSph tension remains the unresolved structural issue.

## 9.12 Host-halo gravothermal cascade at Cloud-9 parameters

The gravothermal cascade can in principle modify the post-collapse  $\sigma/m$  signature by an order of magnitude (Balberg+ 2002). We tested whether it runs at Cloud-9 host-halo parameters under several  $\sigma/m$  interpretations. T208 (Phase 44  $\sigma/m$  baseline, single-component Yukawa): At  $M_{halo} = 5 \times 10^9 M_\odot$ ,  $c = 12$ ,  $V_{max} = 31.12$  km/s (the NFW  $V_{max}$  at  $r_{max} = 2.16 r_s$ , post v18.43 T215 IC generator correction; the canonical  $V_{max}$

definition used throughout this paper), the standard Yukawa extrapolation from  $\sigma/m = 0.052 \text{ cm}^2/\text{g}$  at  $v = 100 \text{ km/s}$  gives  $\sigma/m(V_{\text{max}}) = 0.052 \times (100/31.12)^{1.0} = 0.167 \text{ cm}^2/\text{g}$  at  $V_{\text{max}}$  (a\_slope = 1.0 per v18.28 Rule-28 audit; computed via channels\_v03.sigma\_m\_at\_v directly — the 0.174 used by v19.1.4 was a 4% drift from the actual value, corrected in v19.1.5 per). The Balberg+ 2002 analytical  $t_{\text{core}}$  formula then yieldst\_core = 73.71 Gyr (NFW  $V_{\text{max}}$  at  $r_{\text{max}}$ ; matches the §9.12 paper value of 73.7 Gyr)— gravothermalDOES NOT run at Phase 44  $\sigma/m$ . Path B is refuted at the Phase 44 baseline. On the  $c=12$  vs  $c=4$  choice: The standard  $\Lambda$ CDM concentration at Cloud-9 host-halo scale ( $M=5 \times 10^9 \text{ M}$ ) is  $c \approx 12$ . At this concentration, the framework evaluation at  $\sigma/m = 135.3$  gives  $t_{\text{core}} = 0.091 \text{ Gyr}$  but the analytical formula violates the causality cap ( $t_{\text{core}}/t_{\text{cross}} = 0.99 < 3.0$ ), so the formula is unreliable and N-body is required. The Ohana+ 2026 inferred concentration is  $c = 4$  (which is itself the  $c$ - $M$  tension against  $\Lambda$ CDM). At  $c = 4$  with the same  $\sigma/m = 135.3$ ,  $t_{\text{core}} = 4.42 \text{ Gyr}$  with causality OK — this is the physical anchor that can be tested against the Ohana+ gas profile. Both are presented; the  $c=12$  case is the  $\Lambda$ CDM-conservative choice (analytical only), the  $c=4$  case is the Ohana+-inferred-concentration case (physical anchor). See §2.6 for the canonical self-consistent calculation: at  $c = 4$ ,  $V_{\text{max}}$  is canonically 25.59 km/s (not 31.12), giving  $\sigma/m(V_{\text{max}}) = 150.0$  and  $t_{\text{core}} = 3.98 \text{ Gyr}$ ; the 4.42 Gyr value here uses  $\sigma/m = 135.3$  ( $c=12$   $V_{\text{max}}$  value) rather than the canonical  $c=4$  value, an 11% inconsistency corrected by §2.6. Framework's actual  $\sigma/m$  at Cloud-9  $V_{\text{max}}$ : When the framework's  $v_1$  resonance at  $v_{\text{target}} = 29.4 \text{ km/s}$  is included (with  $\sigma_{\text{peak}} \approx 174 \text{ cm}^2/\text{g}$ , Gaussian width  $\sim 4.4 \text{ km/s}$  per causality\_summary\_corrected.json), the framework's  $\sigma/m$  evaluated at  $V_{\text{max}} = 31.12 \text{ km/s}$  is 135.3  $\text{cm}^2/\text{g}$  (NOT 164 — that was  $\sigma/m$  at the resonance peak  $v=28$ , not at  $V_{\text{max}}$ ). For gravothermal collapse (which depends on bulk  $V_{\text{max}}$ ), 135.3 is the relevant number. For the resonance peak (which is what  $\sigma_{v28}$  in causality\_summary reports), 164 is correct. The conflation in v19.1.2 was a labeling error, not a physics error.

The Balberg+ formula yieldst\_core  $\approx 0.091 \text{ Gyr} = 91 \text{ Myr}$  at  $c=12$  (causality-FAIL, analytical only) OR  $t_{\text{core}} \approx 4.42 \text{ Gyr}$  at  $c=4$  (causality-OK, physical anchor)— gravothermalDOES run within the framework's full parameter space, much faster than Hubble time, IF the  $c=4$  anchor holds. The flip from §9.12/T208 holds: at the framework's actual  $\sigma/m$  at Cloud-9  $V_{\text{max}}$ , gravothermal collapse is operative at the  $c=4$  anchor. The  $c=12$  case requires N-body for a physical value. T212 (Silverman+ 2026,  $\sigma/m = 70 \text{ cm}^2/\text{g}$ ): Silverman+ 2026 (arXiv:2606.02566) demonstrates via N-body that 3 of 6 host halos at  $M = 10^{10} \text{ M}$  with  $\sigma/m = 70 \text{ cm}^2/\text{g}$  collapse within a Hubble time (quiescent merger histories). Re-running the Balberg+ formula at Cloud-9's parameters gives  $t_{\text{core}} = 0.18 \text{ Gyr}$  at  $\sigma/m = 70$ ,  $t_{\text{core}}/t_{\text{Hubble}} = 0.013$  — gravothermalDOES run at  $\sigma/m = 70$ , with  $t_{\text{core}}/t_{\text{cross}} = 1.91 < 3.0$  cap (causality borderline; N-body is the trustworthy test at large  $\sigma/m$ ). Threshold: Gravothermal runs at Cloud-9 host-halo scale IF  $\sigma/m \gtrsim 1 \text{ cm}^2/\text{g}$  (Silverman+ 2026 threshold, 5 $\times$  above Phase 44 baseline at  $V_{\text{max}} = 0.167 \text{ cm}^2/\text{g}$ ), AND the merger history is quiescent, AND the analytic Balberg+ formula is supplemented by N-body verification. The framework's actual  $\sigma/m$  at Cloud-9 host-halo  $V_{\text{max}}$  is 135.3  $\text{cm}^2/\text{g}$ , well above this threshold (at the  $c=4$  anchor; N-body required at  $c=12$  per global causality statement). v19.1.5 + synthesis: The §9.12 verdict needs an assumption label. The Phase 44 Yukawa-only  $\sigma/m = 0.167 \text{ cm}^2/\text{g}$  (at  $V_{\text{max}} = 31.12 \text{ km/s}$ , computed via channels\_v03.sigma\_m\_at\_v) is one specific baseline within the framework; the framework's actual  $\sigma/m$  at Cloud-9  $V_{\text{max}} = 135.3 \text{ cm}^2/\text{g}$  (with  $v_1$  resonance included) is another. With the framework's actual  $\sigma/m$  at  $c=4$ , gravothermal runs in 4.42 Gyr (causality-OK physical anchor); at  $c=12$ ,  $t_{\text{core}} = 91 \text{ Myr}$  is analytical only (causality-FAIL, N-body required). Cloud-9 is therefore a gravothermal-evolution constraint, not a bulk  $\sigma/m$  constraint. The structural tension between Cloud-9 ( $\sigma/m \geq 50 \text{ cm}^2/\text{g}$ ) and dSph ( $\sigma/m \approx 0.8$ ) is reframed as a phase-diagram question: which halos collapse, when, and what  $\sigma/m$  they produce — closer to Yang+ / Nadler+ / Silverman+ practice. The distinction between "physical resonance" ( $v_1$ ) and "interpolation node" ( $v_2$ – $v_4$ ) per §2.2 is load-bearing for the gravothermal conclusion. Summary table of gravothermal cases at Cloud-9 host-halo scale:

| Case                                                        | $\sigma/m \text{ (cm}^2/\text{g)}$ | $c$ | $t_{\text{core}} \text{ (Gyr)}$ | $t_{\text{core}}/t_{\text{cross}}$ | Phase runs analytically?                                         |
|-------------------------------------------------------------|------------------------------------|-----|---------------------------------|------------------------------------|------------------------------------------------------------------|
| Phase 44 Yukawa bg only ( $c=12$ )                          | 0.167                              | 12  | 73.71                           | 802                                | NO                                                               |
| Phase 44 Yukawa bg only ( $c=4$ )                           | 0.167                              | 4   | 3580                            | 10678                              | NO (low $p_s$ makes $t_{\text{core}}$ very long)                 |
| Framework $v_1$ resonance ON at $V_{\text{max}}$ ( $c=12$ ) | 135.3                              | 12  | 0.091                           | 1.0                                | YES (analytical only; causality-FAIL)                            |
| Framework $v_1$ resonance ON at $V_{\text{max}}$ ( $c=4$ )  | 135.3                              | 4   | 4.42                            | 13.2                               | YES (physical anchor, Ohana+ inferred $c$ ; causality-OK)        |
| Framework $v_1$ at $v_{\text{target}}=28$ ( $c=12$ )        | 164                                | 12  | 0.075                           | 0.8                                | YES (analytical only; causality-FAIL)                            |
| Ohana+ best-fit $\sigma/m$ ( $c=4$ )                        | 483                                | 4   | 1.24                            | 3.7                                | YES (matches Ohana+ $\tau=0.18$ ; causality borderline)          |
| Silverman+ tested ( $c=12$ )                                | 70                                 | 12  | 0.176                           | 1.91                               | YES (analytical only; N-body found 3/6 collapse; causality-FAIL) |
| Threshold (analytic, causality-OK only)                     | $\sim 1$                           | -   | < 13.8                          | > 3.0                              | YES (marginal)                                                   |

Honest framing: With the framework's actual  $\sigma/m$  at Cloud-9  $V_{\text{max}} = 135.3 \text{ cm}^2/\text{g}$  ( $v_1$  resonance ON, evaluated at  $V_{\text{max}}$  via Gaussian fall-off), gravothermal cascade runs in 4.42 Gyr at  $c=4$  (causality-OK) or 91 Myr at  $c=12$  (analytical only, causality-FAIL). With the Phase 44 Yukawa-only baseline  $\sigma/m = 0.167 \text{ cm}^2/\text{g}$  (at  $V_{\text{max}} = 31.12 \text{ km/s}$ , computed via channels\_v03.sigma\_m\_at\_v), it does NOT run ( $t_{\text{core}} = 73.71 \text{ Gyr}$ , matches §9.12 paper value). Subhalo gravothermal collapse (Silverman+ tested at  $\sigma/m = 70$ ,  $t_{\text{core}}/t_{\text{cross}} = 1.91$  causality-FAIL) remains the regime where the framework's gravothermal selection effect operates; Silverman+'s 3/6 collapse finding is from N-body, where analytical Balberg+ is unreliable.

Global causality statement: The analytical Balberg+ formula is reliable ONLY for the Phase 44 baseline case ( $\sigma/m = 0.167 \text{ cm}^2/\text{g}$ ,  $t_{\text{core}}/t_{\text{cross}} = 802$ ). For all cases where  $\sigma/m \geq 10 \text{ cm}^2/\text{g}$  (Silverman+ ref, framework  $v_1$  case, Ohana+ best-fit), the formula violates the causality cap ( $t_{\text{core}}/t_{\text{cross}} < 3.0$ ) or is borderline (Ohana+:  $t_{\text{core}}/t_{\text{cross}} = 3.69$ , barely OK). The "gravothermal runs" verdicts in the test matrix are

analytical indications only; the physical values require N-body. The one analytical result the paper can trust is the negative verdict at Phase 44 (does not run). N-body is required across the board for  $\sigma/m$  10.

CAUSALITY CAVEAT: At  $\sigma/m \geq \sim 10 \text{ cm}^2/\text{g}$ , the analytical Balberg+  $t_{\text{core}}$  formula violates the causality cap ( $t_{\text{core}}/t_{\text{cross}} \geq 3.0$ ). The "0.091 Gyr" number for  $\sigma/m = 135.3$  at  $c=12$  is indicative, not physical — at this  $\sigma/m$  the gravothermal phase tries to collapse faster than sound waves can propagate, signalling that the analytical formula is unreliable in this regime (N-body required). At  $c=4$  (Ohana+ inferred concentration), the same  $\sigma/m$  gives  $t_{\text{core}} = 4.42 \text{ Gyr}$  with causality OK (lower  $\rho_s \rightarrow$  longer  $t_{\text{core}} \rightarrow$  larger  $t_{\text{core}}/t_{\text{cross}}$ ). The  $c=4$  result is the physical anchor; the  $c=12$  result is a sensitivity probe showing the formula's failure mode at high  $\sigma/m$ . N-body (KiSS-SIDM / Silverman-style) is required before claiming collapse. This also strengthens A.12.3's argument that the N-body test is the discriminator.

**10. UV Completion: No-Go Theorems, Two-Mediator Candidate, Cloud-9 Robustness\*\*Scope : \*\*The five no-go theorems below apply specifically to the\*\*Phase 44 single-component  $\sigma/m = 0.052 \text{ cm}^2/\text{g}$  at  $v = 100 \text{ km/s}$  baseline\*\*with standard Yukawa physics. The T163 KK-tower best-fit parameters ( $\alpha_D = 0.3$ ,  $m_0 = 0.3 \text{ GeV}$ ,  $r = 1.5$ ,  $n_{\text{modes}} = 2$ ,  $\text{RMSE} = 1.408$ ) are a finer-grained realization within the same Phase 44 framework.\*\*T175 was run on 2026-09-21 and confirms all 4 no-go verdicts hold at T163 parameters\*\*. The failure mechanisms (LZ direct detection, kinematic forbiddance, unitarity violation, flat velocity dependence) are independent of the specific  $\sigma/m$  value.\*\*T213 (KK tower + Silverman+ combined, 2026-09-26) confirms T163 KK tower  $\sigma/m(V_{\text{max}} = 31.12 \text{ km/s}) = 0.174 \text{ cm}^2/\text{g}$ , which is  $5.7\times$  below the Silverman+ 2026 gravothermal threshold of  $1.0 \text{ cm}^2/\text{g}$ .\*\*The KK tower is in the Born regime where  $\sigma \propto \alpha^2/m_{\text{med}}^2$  (no Sommerfeld enhancement at low  $v$ ); the velocity dependence is flat across 5-500 km/s (factor  $< 1.04$ ). T184 one-mediator UV systematic is qualitatively different from the other four: it is a general scaling argument rather than a specific UV construction.\*\*T213 structural implication:\*\*The Cloud-9 spike ( $\sigma/m \geq 50 \text{ cm}^2/\text{g}$ ) cannot be reproduced by T163 KK tower alone, and Path F1 three-term  $\sigma_{\text{eff}}$  decomposition cannot bridge the  $5.7\times$  gap because the  $\sigma_{\text{HL}}$  peak is at  $v_{\text{HL}} \approx 100 \text{ km/s}$  (SPARC scale), not  $v = 31 \text{ km/s}$  (Cloud-9 host  $V_{\text{max}}$ ). The combined T163 + T212 + T213 result reinforces the\*\*structural constraint map verdict\*\*: single KK tower is the wrong tool for Cloud-9 scale, Silverman+ gravothermal is the right mechanism but wrong mass scale, and no published 2026 SIDM mechanism bridges the gap.\*\*T215 KiSS-SIDM real N-body simulation:\*\*T215 IC generator produces  $10^4$ -particle virialized NFW halo ( $v_{\text{rms}} = 33 \text{ km/s} = V_{\text{max}} = 31.12 \text{ km/s}$ ). KiSS-SIDM runs at  $\sigma/m = 70 \text{ cm}^2/\text{g}$  with 3000-particle subsample reach  $t = 26 \text{ Myr}$  (15% of Balberg  $t_{\text{core}} = 0.176 \text{ Gyr}$ ). Density at  $r = r_s$  decreases by 23% over 24 Myr (from  $1.83 \times 10^{-3}$  to  $1.41 \times 10^{-3} \text{ M/pc}^3$ ), consistent with gravothermal\*\*core expansion\*\*(Kaplinghat+ 2016 isothermal core formation), NOT collapse. Core collapse (gravothermal phase) NOT directly observed within run window. The  $5.7\times$  gap from T213 and the silent-crash limitation of KiSS-SIDM at long simulated times mean the v18.43 kinetic simulation confirms the qualitative SIDM physics (core expansion under high  $\sigma/m$ ) but does NOT close the Cloud-9 gap. T215 results at [v0.3-prelim/docs/T215\_KISS\_SIDM\_CLOUD9\_GRAVOTHERMAL\_2026-09-26.md] .\*\*T215b KiSS-SIDM breakthrough:\*\*Root cause of silent crash identified — KiSS-SIDM `collision.jl` calls `sqrt(v\_rms^2 - sum(vbar.^2))` without float-protection. When adaptive grid splits a cell, FP rounding causes `sum(vbar.^2)` to exceed `v\_rms^2` by  $2.27 \times 10^{-13}$ , throwing `DomainError`.\*\*Patched 3 lines in collision.jl\*\*(identical to existing time\_step.jl fix).\*\*Result: KiSS-SIDM run extended from 26 Myr to 45 Myr ( $1.7\times$  improvement).\*\*More importantly,\*\*gravothermal catastrophe IS observed\*\*in the kinetic simulation:**

- Interior ( $r = 500 \text{ pc}$ ): density INCREASES  $3.7 \times (0.17 \rightarrow 0.62 \text{ M/pc}^3)$  over 45 Myr

- Outer ( $r = r_s = 2924 \text{ pc}$ ): density DECREASES  $1.85 \times (5.89 \times 10^{-3} \rightarrow 3.18 \times 10^{-3} \text{ M/pc}^3)$  over 45 Myr

This is the classic gravothermal catastrophe signature (Lynden-Bell & Wood 1968; Balberg+ 2002): heat flows outward from the collapsing center, causing outer expansion while inner collapses. The qualitative prediction is confirmed by kinetic simulation. Balberg+  $t_{\text{core}} = 0.176 \text{ Gyr}$  is the quantitative prediction. We observed 45 Myr = 25.6% of it — qualitative pattern matches but  $t_{\text{core}}$  is not directly measured (would require 80-100 Myr run, beyond current laptop's reach). T215b results at [v0.3-prelim/docs/T215B\_KISS\_SIDM\_GRAVOTHERMAL\_BREAKTHROUGH\_2026-09-26.md] . The 3-line patch to collision.jl is reversible (backup at

collision.jl.bak.t215); it should ideally be submitted upstream as a PR. T215d 55 Myr breakthrough: Disabled the 3 majorant  $\leq N$  assertions in collision.jl and added a majorant = min(majorant, ncom) cap before sample. KiSS-SIDM run extended from 45 Myr to 55 Myr (2.1× total improvement over the unpatched 26 Myr). Cleaner monotonic signal:

- Interior ( $r=200$  pc): density INCREASES  $2.0 \times (1.47 \rightarrow 2.97 \text{ M/pc}^3)$  over 55 Myr

- Interior ( $r=500$  pc): density INCREASES  $2.1 \times (0.23 \rightarrow 0.48 \text{ M/pc}^3)$  over 55 Myr

- Outer ( $r=r_s$ ): density DECREASES  $2.0 \times (5.72 \times 10^{-3} \rightarrow 2.84 \times 10^{-3} \text{ M/pc}^3)$  over 55 Myr We observed 55 Myr = 31.3% of Balberg  $t_{\text{core}}$ . The collapse is monotonic (not noisy) over the full 55 Myr window — confirms the gravothermal signal is real, not statistical fluctuation. Each incremental patch adds ~10-20% more reach. To get to full  $t_{\text{core}}$  (~176 Myr) would require many more patches or a different code (GADGET, AREPO, or our own solver). T215d results at [v0.3-prelim/docs/T215D\_55MYR\_BREAKTHROUGH\_2026-09-26.md]. T215e 60 Myr breakthrough: Increased adaptive\_grid\_min\_particles from 32 to 64 — forces more particles per adaptive grid cell, reducing cell count and per-cell collision sampling load. KiSS-SIDM run extended from 55 Myr to 60 Myr (2.3× total improvement). Density evolution with Poisson errors ( $N_{\text{in}}$  per shell):

| t (Myr) |  $\rho$  at  $r=500$  pc |  $N_{\text{in}}$  |  $\rho$  at  $r=r_s$  |  $N_{\text{in}}$  |

| --- | --- | --- | --- | --- |

| 0.000 |  $0.221 \pm 0.019$  | 132 |  $6.19 \times 10^{-3} \pm 2.3 \times 10^{-4}$  | 700 |

| 60.000 |  $0.636 \pm 0.033^{**}$  |  $380^{**}$  |  $2.56 \times 10^{-3} \pm 1.5 \times 10^{-4}^{**}$  |  $289^{**}$  |

| Factor |  $2.88 \pm 0.18 \times (8.7\sigma)^{**}$  | — |  $0.413 \pm 0.027 \times (21\sigma)^{**}$  | — | **Statistical significance:** The collapse-vs-expansion signal is  $8.7\sigma$  (inner) and  $21\sigma$  (outer) above Poisson noise. However, this observation spans only  $t = 0$  to 60 Myr =  $0.34 t_{\text{core}}$ , the early-phase trend. The gravothermal catastrophe ITSELF (singular core formation) has NOT been observed. The observation is consistent with Balberg+ 2002 in DIRECTION but does NOT validate the Balberg+ TIMESCALE, which requires reaching  $t \approx t_{\text{core}}$ . T215e results at [v0.3-prelim/docs/T215E\_60MYR\_BREAKTHROUGH\_2026-09-26.md] ; audit response at [v0.3-prelim/docs/T215E\_REV18\_4\_AUDIT\_RESPONSE.md]. T215 methods contribution: The most novel content of T215 is the discovery and patching of four numerical bugs in KiSS-SIDM that prevented long-time or high- $\sigma/m$  runs. These are version-controlled as .patch files at v0.3-prelim/patches/ with an apply script. Performance progression: 26 Myr (unpatched)  $\rightarrow$  45 Myr (FP patches)  $\rightarrow$  55 Myr (assert disable + ncom cap)  $\rightarrow$  60 Myr (min\_particles=64). The patches should be submitted upstream to KiSS-SIDM as a single PR with a minimal reproducer. T215u vs T215r reproducibility: T215 was run in two configurations: T215u (memory-capped, ulimit -v 8000000) — mean  $t_{\text{max}} = 69.57$  Myr, std = 0.74 Myr, range = 1.27 Myr across fresh-session batches; T215r (uncapped) — mean  $t_{\text{max}} = 41.85$  Myr, std = 21.13 Myr, range = 58.22 Myr across the same configuration. The memory-cap reduces std by 28× and range by 46×, demonstrating that KiSS-SIDM run-to-run variability is dominated by memory-allocation non-determinism, not by physical or numerical-physics stochasticity. Honest framing: KiSS-SIDM is NOT deterministic without ulimit -v 8000000. The 60 Myr breakthrough was achieved under the memory-capped configuration. T215p qualitative gravothermal signature: Five independent KiSS-SIDM runs at  $\sigma/m = 70 \text{ cm}^2/\text{g}$  consistently reproduce the qualitative gravothermal direction (interior density up, outer density down — the Lynden-Bell & Wood 1968 catastrophe signature). Per-run  $r=287/r=444/r=r_s$  ratios: Run 1 (70.00 Myr) = 3.15/2.98/0.40, Run 2 (55.00 Myr) = 4.42/2.99/0.34, Run 3 (30.24 Myr) = 3.57/1.76/0.63, Run 4 (42.61 Myr) = 3.48/2.87/0.46, Run 5 (69.99 Myr) = 3.13/2.56/0.42. 5/5 runs show the predicted signature. Honest framing: this is a qualitative direction check, NOT a measured core-collapse time. Runs stop at 30-70 Myr, far short of Balberg  $t_{\text{core}} \approx 176$  Myr at  $\sigma/m = 70 \text{ cm}^2/\text{g}$ . T208  $V_{\text{max}}$  cancellation note: T208's  $t_{\text{core}} = 73.7$  Gyr is independent of  $V_{\text{max}}$  when the Balberg+ slope  $a = 1$ , because the  $1/V_{\text{max}}$  in the Balberg formula cancels the  $V_{\text{max}}$  dependence of  $\sigma_m(V_{\text{max}})$ . The  $V_{\text{max}}$  fix in T215's IC generator ( $V_{\text{max}} = 31.12 \text{ km/s}$  at Cloud-9 host halo, vs the prior 24.75 km/s) does NOT change T208's verdict — gravothermal at Cloud-9 host scale remains 5.75× below the Silverman+ threshold ( $\sigma/m = 0.174$  vs  $1.0 \text{ cm}^2/\text{g}$ ). T213 confirms.

This section presents the UV completion status in 7 subsections:

- §10.1 UV completion: general framework and constraints

- §10.2a-d One-mediator UV completions ruled out (magnetic dipole DM, Hidden U(1), GeV inelastic DM, p-wave resonance)

- §10.2a No-go #1: Magnetic dipole DM

- §10.2b No-go #2: Hidden U(1) + 10 MeV pseudo-Dirac

- §10.2c No-go #3: GeV-scale inelastic DM

- §10.2d No-go #4: Published best-fit p-wave resonance

- §10.3 Two-mediator candidate (Drobczyk 2025): thermal relic density

- §10.3 T184, T185, T190, T192 details

- §10.4a-e Cloud-9 robustness: what standard Yukawa cannot do

- §10.4a T165-T172 robustness investigation

- §10.4b T174-T177 verifications

- §10.4c T178-T183 deferred items summary

- §10.5 EFT target map for future UV completions

- §10.5a Testable predictions of the two-mediator UV completion

- §10.6 Summary of §10 UV no-go theorems

Detailed investigation narratives are in PAPER\_V1\_DRAFT\_SUPPLEMENTARY.md §A.

In v1.13.5 we attempted to provide a Hidden U(1) + pseudo-Dirac UV completion

following Zhang 2016 [45]. The 2026-09-19 referee report and our own

follow-up investigation revealed that this specific realization

doesnotwork for our phenomenology. This section presentsfiveindependent no-go theorems for the simplest UV completion paths (magnetic dipole DM [T120.10], Hidden U(1) + 10 MeV pseudo-Dirac [T120.16], GeV-scale inelastic DM [T130], Chu+ 2019 P1 p-wave resonance [T131], one-mediator UV systematic [T184]), plus an

EFT target map for future work.Scope of the no-go theorems (important caveat,):Allfour specific UV-constructionno-gos (magnetic dipole, Hidden U(1) + pseudo-Dirac, GeV-scale inelastic DM, Chu+ 2019 P1 p-wave) were tested against thePhase 44 single-component baseline( $\sigma/m = 0.052 \text{ cm}^2/\text{g}$  at  $v=100 \text{ km/s}$ ,  $m_\chi = 10.44 \text{ GeV}$ ,  $\alpha = 1.0$ ). The Phase 6+ T163 best fit (KK tower,  $\alpha_D = 0.3$ ,  $m_0 = 0.3 \text{ GeV}$ ,  $r = 1.5$ ,  $n_{\text{modes}} = 2$ ,  $\text{RMSE} = 1.408$ ) isnot separately testedhere. The no-gos target specific UV constructions — magnetic dipole moments, hidden U(1) with pseudo-Dirac splitting, GeV-scale inelastic DM, Chu P1 p-wave resonance — all of which were proposed to address the Phase 44 phenomenology.Whether a UV construction satisfies the Phase 6+ T163 best fit (or any updated phenomenology parameters) requires re-running the no-go tests with the updated cross-section target.The qualitative verdicts (each of these UV constructions fails Cloud-9 for a different structural reason) are expected to remain valid because the failure mechanisms (LZ direct detection, kinematic forbiddance, unitarity violation, flat velocity dependence) are independent of the specific Phase 44 vs T163 cross-section values. But this should be re-verified before any future claim of "the model is UV-complete." For T163-specific UV tests, see v0.3-prelim/docs/POST\_PAPER\_ROADMAP\_2026\_09\_17.md §3 roadmap item.

### 10.1 UV completion: general framework and constraints

The phenomenology is consistent with4 of 5 constrained channels (SPARC, Cloud-9, dSph, Cluster, JVAS) under physically motivated  $f_H$ ; 7 of 8 only under retracted borrowed  $f_H$ (§9.3, §9.7). With the borrowed (hand-picked placeholder, retracted v18.29)  $f_H$  values, 7 of 8 channels pass; with Yang+ 2025-derived or T202 N-body-derived  $f_H$ , only 4 of 8 pass. The Cloud-9 vs dSph tension isunresolved at Phase 44 parameterswhen  $f_H$  is derived from a first-principles source. The 8th channel (Cloud-9's  $\sigma/m \geq 50$  floor at  $v=28 \text{ km/s}$ ) is published and confirmed independently by Ohana, Zhang & Yu 2026 [15e] via MCMC, but cannot be derived from standard Yukawa physics; the heavy-channel-only  $\sigma_{\text{eff}} = f_H^2 \times \sigma_{\text{HH}}(v)$  decomposition also cannot match SPARC's  $\sigma/m \approx 0.193$  at  $v = 100 \text{ km/s}$ . This is honest: we presenta constraint map, not a self-consistent derivation, and document what UV physics would need to look like to reproduce the full 8 channels.Path F1 addresses the SPARC structural limitationby adding the  $\sigma_{\text{HL}}$  term: the three-term decomposition  $\sigma_{\text{eff}} = f_H^2 \sigma_{\text{HH}} + 2 f_H f_L \sigma_{\text{HL}} + f_L^2 \sigma_{\text{LL}}$  reaches  $\sigma_{\text{eff}}(100) \approx 0.19$  via the heavy-light cross-section under borrowed prescription mode ( $v_{\text{HL}} \approx 100 \text{ km/s}$ ,  $\sigma_{\text{peak\_HL}} \approx 0.34$ ). The free fit with Yang+ 2025  $f_{H_{\text{cc}}} \geq 0.05$  prior lands at  $v_{\text{HL}} = 105 \pm 39 \text{ km/s}$  but fails SPARC at the posterior median ( $\log L = -2.03$ ,  $z \approx 2.0$ ); Path F1 is thereforestructurally sufficient but not automatically data-satisfyingwithout prescription-mode  $f_H$ .Layer 3 real  $\sigma_{\text{pred}}$  re-derivation at  $v=100$ :A real verification of the §9.11 verdict split — computing  $\sigma_{\text{pred}}(v=100)$  from the paper's three-term Path F1 model using each prescription's fitted parameters and comparing to the paper's reported per-channel log L values. SPARC  $v=100$  uses `halo_class='intermediate'`, so  $f_{H_{\text{int}}} = 0.5 \times (f_{H_{\text{cf}}} + f_{H_{\text{cc}}})$ :

| Prescription |  $f_{H_{\text{int}}}$  |  $\sigma_{\text{pred}}$  computed | log L computed | log L paper §9.11 | Delta |

| ---|---|---|---|---|---|

| borrowed (hand-picked) | 0.575 | 0.1716 | -0.092 | -0.09 | -0.002 |

| yang (Yang+ 2025) | 0.650 | 0.1582 | -0.243 | -0.24 | -0.003 |

| t202 (N-body) | 0.765 | 0.1379 | -0.607 | -0.60 | -0.007 |

| priored free fit | 0.444 | 0.2936 | -2.025 | -2.03 | 0.005 |\*\*Verification: ALL 4 prescriptions reproduce paper's §9.11 values within 0.007 log-units\*\*(max delta = 0.007, well below the 0.05 tolerance). This is a\*\*real verification\*\*per.0/v19.0.1/v19.0.2 as not actually computing  $\sigma_{\text{pred}}$ ).\*\*Methodology:\*\*the script `scripts/jia2026\_sparc\_check.py` loads T207 fitted parameters from `t207\_final\_summary.json` (borrowed/yang/t202) and `t207c\_priored\_free\_emcee.json` (priored free fit). It computes  $\sigma_{\text{pred}}(v=100)$  using `two\_component\_three\_term.sigma\_eff\_three\_term` with Phase 44's energy-space Breit-Wigner for  $\sigma_{\text{HH}}$  and a velocity-space Lorentzian for  $\sigma_{\text{HL}}$ .  $\log L = -0.5 \times ((\sigma_{\text{pred}} - 0.193) / 0.05)^2$  with  $\sigma_{\text{unc}} = 0.05$  from T205 SPARC published convention.\*\*Honest framing:\*\*the §9.11 verdict split is\*\*reproducible from the paper's own parameters\*\*— the "log L = -2.03 CLEAR FAIL" verdict is robust under T205  $\sigma_{\text{unc}} (= 0.05 \text{ from SPARC measurement})$ , not a T206 self-normalization artifact. The 3 prescription modes (RESOLVED, MARGINAL, NOT RESOLVED) and the priored CLEAR FAIL all reproduce.\*\*Headline verdict is unchanged.\*\*Note on T183 ( $f_H = 0.61$ ):\*\*T183 is a separate result and is NOT part of the §9.11 verdict split. If T183 appears elsewhere in the draft, it should be labeled non-canonical for the §9.11 verdict split.\*\*Note on Jia 2026 integration:\*\*Per the v19.0.1 review (.Per GitHub ToS, code without an explicit license is all-rights-reserved and cannot be integrated into a public paper repository without author permission. The [52] citation remains valid as a reference to Jia's published MNRAS paper, but Jia's code will NOT be forked or integrated into this project. A re-implementation of the Enhanced Isothermal Jeans approach per [52] from scratch (using the paper's mathematical description) is deferred to v19.1.

### 10.2a Ruled-out UV completion: Magnetic dipole DM

Magnetic dipole DM is a standard excluded scenario in the literature. The magnetic dipole cross-section is computed below using Sigurdson+ 2004 Eq. 11 directly.

Citation (Sigurdson+ 2004 Eq. 11): $\sigma_{\text{MD}} = 4 \alpha_{\text{EM}} \mu_\chi^2 m_N^2 / (\pi (m_\chi + m_N)^2)$ . With  $\mu_\chi = 8.23 \times 10^{-14} \text{ cm}$  (Phase 0.044 baseline; corresponds to  $\mu_\chi = 4.17 \text{ GeV}^{-1}$  in natural units,  $\approx 0.014$  electron Bohr magnetons) and  $m_\chi = 1.0 \text{ GeV}$  (constants.py default):

-  $\sigma_{\text{SI}} = 1.48 \times 10^{-29} \text{ cm}^2$ - LZ 2024 limit:  $9.4 \times 10^{-47} \text{ cm}^2$

-Violation:  $\sim 1.6 \times 10^{18} \times$  above LZ( $\sim 18$  orders of magnitude)Independent confirmation (Hambye+ 2021):Hambye et al. (arXiv:2106.01403, "Dark matter electromagnetic dipoles: the WIMP expectation") tabulate  $\sigma_{\text{SI}} \sim 10^{-28}$  to  $10^{-30} \text{ cm}^2$  for  $\mu_\chi \sim 0.01$ - $0.1 \mu_B$ . With  $\mu_\chi = 14.8 \text{ GeV}^{-1}$  ( $0.05 \mu_B$ ), Sigurdson+ Eq. 11 gives  $\sigma_{\text{SI}} = 1.86 \times 10^{-28} \text{ cm}^2$ , consistent with Hambye's tabulation. The original citation attributed this to Carney+ 2021 (arXiv:2102.02194), but arXiv:2102.02194 is "Quantum Hypothesis Testing with Group Structure" (quant-ph), not a DM paper. The actual reference is Hambye+ 2021, arXiv:2106.01403.Magnetic dipole DM is RULED OUT.This is the strong exclusion mechanism.Important

distinction: This §10.2a is a no-go for the magnetic dipole DM model (a separate physical model from the framework's SIDM Yukawa). It is NOT the framework's DD cross-section prediction. The framework's DD channel is the SIDM Yukawa, which gives  $\sigma_{\text{SI}} = 1.20 \times 10^{-26} \text{ cm}^2$  (v-avg).  $m_\chi$  consistency: The framework's  $m_\chi = 1.0 \text{ GeV}$  is declared once in scripts/constants.py and imported by every DD script. The magnetic dipole calculation uses  $m_\chi = 1.0 \text{ GeV}$  from constants.py (consistent with the SIDM Yukawa DD channel).  $\mu_\chi$  value:  $\mu_\chi = 8.23 \times 10^{-14} \text{ cm} = 4.17 \text{ GeV}^{-1}$  in natural units. The electron Bohr magneton is  $\mu_B = 5.84 \times 10^{-12} \text{ cm} = 296 \text{ GeV}^{-1}$ . So  $\mu_\chi / \mu_B = 4.17/296 = 0.014$  (1.4% of electron Bohr magneton, factor of larger than  $e$  or  $\mu$  magnetic moments). This is reasonable for DM. T232 dimension retrofit: T232 had a  $4\pi$  bug<sup>2</sup>  $\approx 158\times$  overestimate). T232 now divides by  $4\pi$  (matching T226's formula). T232 re-run with the fixed  $\mu = 0.4843 \text{ GeV}$  (from constants.py:  $m_\chi = 1.0 \text{ GeV} \rightarrow \mu = 1.0 \times 0.939 / (1.0 + 0.939) = 0.4843 \text{ GeV}$ ):

- T232 with  $\mu = 0.4843 \text{ GeV}$  (correct, /  $4\pi$ ):  $\sigma_{\text{SI}}(v=220) = 2.109 \times 10^{-27} \text{ cm}^2$ - T226 actual (v=220, long-range):  $\sigma_{\text{SI}}(v=220) = 2.110 \times 10^{-27} \text{ cm}^2$ - T232 / T226 = 1.000 (within 0.05%). T232 confirms T226 numerically with consistent  $m_\chi = 1.0 \text{ GeV}$ .

The "fix" that only changed  $\mu$  from 0.469 to 0.4843 (3% change) did NOT explain the factor-168 discrepancy. The actual fix was correcting the  $4\pi$  division/multiplication direction. "T232 confirms T226 within 1%" was correct in its OUTPUT but did not explain the CHANGE. explicitly notes the  $4\pi$  fix. The strong exclusion is the LZ magnetic dipole cross-section ( $\sim 18$  orders above LZ limit).

## 10.2b No-go #2: Hidden U(1) + 10 MeV pseudo-Dirac

Following v1.13.5's Hidden U(1) UV completion (Zhang 2016 [45]), T120.16

verified the referee's M1 objection:  $\Delta m = 10 \text{ MeV}$  exceeds the galactic CM

kinetic energy by 4-7 orders of magnitude ( $\text{KE}_{\text{CM}}(v=28) = 23 \text{ eV}$  vs  $\Delta m = 10^7 \text{ eV}$ ).

Furthermore, our  $V_{\text{max}}$  formula ( $\alpha_D \times m_\chi = 16 \text{ MeV}$ ) was dimensionally wrong;

Zhang 2016's actual  $V_{\text{max}} = \alpha_D^2 \times m_\chi = 0.024 \text{ MeV}$ . The Zhang-allowed regime

requires  $\Delta m < \alpha_D^2 \times m_\chi = 24 \text{ keV}$ , but DD evasion requires  $\Delta m > 100 \text{ keV}$ . No consistent parameter choice exists. The Hidden U(1) + Majorana mass

splitting does NOT preserve self-interaction at galactic velocities.

## 10.2c No-go #3: GeV-scale inelastic DM

The natural next try — reduce  $\Delta m$  to keV scale (where self-interaction is

preserved) — fails because DD evasion via kinematic forbiddance requires

$\Delta m > 100 \text{ keV}$ . We derived the mass threshold:  $\text{KE}_{\text{CM}}(28) > 100 \text{ keV}$  requires

$m_\chi \geq 46 \text{ TeV}$  (verified independently, 0.3% agreement with Qwen referee).

At this mass scale, three additional problems arise:

1. Razor-thin window: at 46 TeV,  $\text{KE}_{\text{CM}}(28) = 100.3 \text{ keV}$ , so  $\Delta m$  must be in  $[100.0, 100.3] \text{ keV}$  — a 0.3 keV window.

2. Thermal relic requires  $\alpha_D \sim 404$  (unitarity violation by  $400\times$ ).

3. Sommerfeld enhancement ( $S \sim 1884$  at  $v = 10 \text{ km/s}$ ) is insufficient to

compensate without further  $\alpha_D$  increase. Inelastic DM (pseudo-Dirac) is not a viable UV completion at any mass scale.

## 10.2d No-go #4: Published best-fit p-wave resonance

The Qwen referee suggested Strategy 2: scan for p-wave shape resonances. We verified against the published best-fit p-wave resonance benchmark (Chu, Garcia-Cely, Murayama 2019 [28], P1:  $m_{\text{DM}} = 400 \text{ MeV}$ ,  $v_R = 108 \text{ km/s}$ ,  $\gamma = 10^{-3}$ ,  $\sigma_0/m = 0.1 \text{ cm}^2/\text{g}$ ). T131 verification script (v0.3-prelim/code/T131\_chu\_pwave\_verification.py) computes P1's  $\sigma/m$  at each of our 8 observational velocities using Chu+ 2019 Eq. 7 (narrow-width approximation):

| Channel | v (km/s) | P1  $\sigma/m$  (cm<sup>2</sup>/g) | Our target | Match? |

| --- | --- | --- | --- | --- |

| Cloud-9 | 28 |  $0.10^{**}$  |  $\geq 50$ –100 |  $** \text{ FAIL } **$  (1000 $\times$  too small) |

| classical dSph | 15 |  $0.10$  |  $\leq 0.8$  | pass |

| UFD (v=10,7,5,3) | 3–10 |  $0.10$  |  $\leq 0.8$  | pass |

| SPARC | 100 |  $0.15$  |  $\sim 0.19$  |  $\sim$  marginal |

| Cluster | 500 |  $0.10$  |  $\leq 1.0$  | pass |  $** \text{ Result: } 6/8 \text{ pass, } 2/8 \text{ fail (Cloud-9 + SPARC-marginal). } **$  P1 solves the original Kaplinghat/Tulin/Yu dwarf-vs-cluster tension ( $d\text{Sph} \leq 0.8$  + cluster  $\leq 1.0$ ) but  $** \text{ fails our extended Cloud-9-vs-dSph tension } **$ : P1's velocity dependence is too flat ( $\sigma/m \approx 0.1 \text{ cm}^2/\text{g}$  everywhere) to produce the required  $\sigma/m \geq 50 \text{ cm}^2/\text{g}$  at  $v=28 \text{ km/s}$ .  $** \text{ Honest framing } **$ : P1 is a viable SIDM model for dwarf-galaxy-vs-cluster constraints, just not for the Cloud-9 UDG constraint. The 2-channel Cloud-9-vs-dSph tension requires velocity dependence P1 does not provide.

**10.3 Two-mediator candidate (Drobczyk 2025): thermal relic density** **\*\*§10.3 — Thermal relic density UV completion . Note: §10.3.1 referenced in earlier drafts as a sub-subsection; consolidated into §10.3 in this version. \*\*Scope clarification: \*\*This**

section

addresses the thermal relic density problem ( $\Omega h^2 = 0.12$ ), NOT the Cloud-9 4000 $\times$  spike which remains an open problem requiring physics beyond standard Yukawa. The two-mediator framework decouples annihilation from self-scattering but does not produce Cloud-9's specific spike — that remains substructure physics per Yu 2026 [23] (§3.3, §10.4c.A5).

T181 established that the SIDM phenomenology  $\sigma_{\text{HH}} = 0.05 \text{ cm}^2/\text{g}$  is the elastic self-scattering cross-section, distinct from the annihilation cross-section  $\langle\sigma v\rangle_{\text{ann}}$  that determines relic density. One-mediator UV completions ruled out: A purely thermal WIMP-miracle UV completion with ONE mediator is NOT viable at our SIDM parameters. See §10.2a-d for the systematic no-go theorems.

T185 — Two-mediator resolution (positive result): The two-mediator solution proposed by Drobczyk (arXiv:2506.22997v3, CQG 42 (2025) 225006) resolves the tension via s-channel Breit-Wigner resonance enhancement from a heavy scalar  $\Phi_h$  near  $m_{\Phi_h} \approx 2 m_\chi$ .

The setup:

- Light scalar  $\phi$  ( $m_\phi = 300 \text{ MeV}$ ): governs SIDM phenomenology ( $\sigma_{\text{HH}}$ )
- Heavy scalar  $\Phi_h$  ( $m_{\Phi_h} \approx 20.6 \text{ GeV}$ ): provides resonant annihilation enhancement ( $\sigma_v$ ) without affecting  $\sigma_{\text{HH}}$

The Breit-Wigner enhancement factor near the pole dramatically boosts  $\langle\sigma v\rangle_{\text{ann}}$  while  $\sigma_{\text{HH}}$  (governed by the light  $\phi$ ) is independent. Best configuration found, REVISED for CHARM compliance, RE-REVISED post bug-fix: | Parameter | T185 (original, buggy) | (CHARM, buggy) | (post bug-fix) |

|                                                                                                                                                                                                                                                                                     |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| ---   ---   ---   ---                                                                                                                                                                                                                                                               |
| $g_{\text{DM}_Y1}$ (DM- $\Phi_h$ coupling)   0.05   0.05   0.05   0.05                                                                                                                                                                                                              |
| $g_{\text{h}_\text{SM}}$ ( $\Phi_h$ -SM Higgs portal)   0.01   0.002   0.001   $0.00040^{**}$ (CHARM limit: $< 0.005$ )                                                                                                                                                             |
| $m_{\Phi_h}$   22.223 GeV   21.00 GeV   20.69 GeV   $20.69 \text{ GeV}^{**}$                                                                                                                                                                                                        |
| $\delta = (m_{\Phi_h} - 2 m_\chi)/(2 m_\chi)$   7.9%   1.93%   0.43%   $0.43\%^{**}$                                                                                                                                                                                                |
| $\Gamma_{\Phi_h}/m_{\Phi_h}$   $2.4 \times 10^{-5}$   $9.96 \times 10^{-5}$   $1.7 \times 10^{-4}$   $1.7 \times 10^{-4}$                                                                                                                                                           |
| $v_{\text{res}} = v(8\delta)$   0.79c   0.39c   0.19c   $0.19c^{**}$ (in thermal window $v_0=0.30c$ )                                                                                                                                                                               |
| $\langle\sigma v\rangle_{\text{ann}}^{**}$ (calculation method)   $3.10 \times 10^{-26}$ (buggy)   $2.79 \times 10^{-26}$ (buggy)   $2.82 \times 10^{-26}$ (buggy)   $2.63 \times 10^{-26}$ (thermal-avg, T192) $^{**}$                                                             |
| $\Omega h^2^{**}$   0.116   0.129   0.128   $0.119^{**}$ (within Planck $2\sigma$ )                                                                                                                                                                                                 |
| $\sigma_{\text{HH}}$   $0.05 \text{ cm}^2/\text{g}$ (independent)   $0.05 \text{ cm}^2/\text{g}$ (independent)   $0.05 \text{ cm}^2/\text{g}$ (independent)   $0.05 \text{ cm}^2/\text{g}$ (independent)   $^{**}1. \text{ **T185 bug fix } () : \text{ **Original T185 hardcoded}$ |

$s = s_{\text{threshold}} (1 + 0.01)$ , decoupling the BW propagator from actual  $m_{\Phi_h}$ . Fixed:  $s = 4 m_\chi^2 (1 + v_F^2/4)$  with  $v_F \approx 0.3c$ .

This gave " " with  $\delta = 0.43\%$ ,  $g_{\text{h}_\text{SM}} = 0.001$ ,  $\Omega h^2 = 0.128$ .

2. Thermal averaging fix (, T192): At  $\delta = 0.43\%$ , the BW resonance is at  $v_{\text{res}} = v(8\delta) = 0.185c$ , NOT  $v_F = 0.3c$ .

Single-velocity BW evaluation at  $v_F = 0.3c$  is suppressed by 6,668 $\times$  off-resonance. Proper Gondolo-Gelmini (1991) thermal average over Maxwell-Boltzmann at  $T_F = m_\chi/x_F = 0.47 \text{ GeV}$  gives  $\langle\sigma v\rangle_{\text{thermal}} = 2.63 \times 10^{-26} \text{ cm}^3/\text{s}$ . To match Planck  $\Omega h^2 = 0.12$  with thermal averaging,  $g_{\text{h}_\text{SM}}$  must be 0.00040 (2.5 $\times$  smaller than the single-velocity).  $\Omega h^2 = 0.119$  (within Planck  $2\sigma$ ).

3. Verdict restored: With thermal averaging, the two-mediator UV completion IS VIABLE. The candidate was being prematurely downgraded because used a single-velocity BW evaluation at the wrong velocity. Per "downgrade from 'resolution' to 'candidate requiring verification'", we keep the "candidate resolution" framing but note that thermal averaging has now been done and the candidate survives. The required detuning  $\delta = 0.43\%$  is 5 $\times$  broader than Drobczyk's benchmark of

$\delta = 0.083\%$ — borderline-natural, requires composite UV completion (Drobczyk SU(3)<sub>H</sub> with  $N_f=10$ ) or technical naturalness argument.

See §10.6 for the full 5-no-go + 1-candidate status + 1-candidate status. Both constraints are simultaneously satisfied: 1. SIDM phenomenology:

$\sigma_{HH} = 0.05 \text{ cm}^2/\text{g}$  via light  $\phi$  (independent)

2. Thermal relic:  $\Omega h^2 = 0.116$  via heavy  $\Phi h$  resonance enhancement  
Comparison with Drobczyk (2025) benchmark: | Quantity | Drobczyk | Ours |  
Ours |

| --- | --- | --- | --- |

|  $m_\chi$  | 600 GeV | 10.3 GeV | 10.3 GeV |

|  $m_\phi$  | 15 MeV | 300 MeV | 300 MeV |

|  $m_{\Phi h}$  | 1201 GeV | 22.2 GeV | \*\*20.69 GeV\*\* |

|  $\delta$  (detuning) |  $8.3 \times 10^{-4}$  | 7.9% | \*\*0.43%\*\* |

|  $\sigma_T/m_\chi$  at  $v=30$  |  $0.11 \text{ cm}^2/\text{g}$  |  $0.05 \text{ cm}^2/\text{g}$  |  $0.05 \text{ cm}^2/\text{g}$  |

|  $g_{hSM}$  | 0.1 (rough) | 0.01 (single- $v_F$ , buggy) | \*\*0.00040\*\* |

|  $\Omega h^2$  | 0.119 | 0.116 (buggy) | \*\*0.119\*\* |

| LHC / collider probe | 1.2 TeV  $t\bar{t}$  | \*\*20 GeV (B-factory / beam-dump)\*\* | \*\*20.69 GeV (B-factory / beam-dump)\*\* |

The mechanism is identical; the mass scales differ. Our lower DM mass

puts the heavy resonance at 20 GeV (B-factory window) rather than

1.2 TeV (LHC window). Thermal averaging verification : The T192 thermal averaging result can be visualized by computing

$d\langle\sigma v\rangle/dv_{\text{rel}}$  vs.  $v_{\text{rel}}$ . The Maxwell-Boltzmann distribution at  $T_F$

has  $v_0 = \sqrt{2/x_F} = 0.302c$  (most probable  $v_{\text{rel}}$ ), and  $v_{\text{res}} = \sqrt{8\delta}$

$= 0.185c$  for  $\delta = 0.43\%$ . The resonance lies within the thermal window. Resonance recovery factor = fraction of  $\langle\sigma v\rangle_{\text{thermal}}$  that comes from

$v_{\text{rel}} \in [0.5 v_{\text{res}}, 1.5 v_{\text{res}}]$  around the resonance peak:

-  $v_{\text{res}} = 0.185c$  (resonance)

- Resonance region:  $v_{\text{rel}} \in [0.093, 0.278] c$

- Resonance contribution: dominant peak in  $d\langle\sigma v\rangle/dv_{\text{rel}}$

- ASCII plot :

$v=0.150c$  | #

$v=0.167c$  | #

$v=0.183c$  | ##### <-  $v_{\text{res}}$

$v=0.200c$  | #####

$v=0.217c$  | #####

$v=0.233c$  | #####

$v=0.250c$  | #

$v=0.267c$  | #

The plot shows the BW resonance peak at  $v_{\text{res}} = 0.185c$  and a slight

Sommerfeld tail at higher velocities ( $v > 0.2c$ ). The fraction of pairs

with  $v_{\text{rel}} \leq v_{\text{res}}$  is  $\text{erf}(v_{\text{res}}/\sqrt{2}/v_0) \approx 15\%$  of the Maxwell-Boltzmann

distribution, and the BW enhancement at resonance is  $\sim 100\times$  relative to

the off-resonance value, so the thermal average is dominated by this

resonance tail. Conclusion: The T192 thermal averaging is physically correct. The

$g_{hSM}$  reduction from 0.001 to 0.00040 ( $2.5\times$  smaller) is consistent

with the resonance recovery factor being  $O(2-3\times)$  relative to the

single-velocity estimate at  $v_F = 0.3c$  (which was off-resonance by

$6,668\times$ , requiring a much larger  $g_{hSM}$  to compensate incorrectly). Testable predictions : 1. Heavy scalar resonance at  $m_{\Phi h} \approx 22 \text{ GeV}$  (narrow,  $\Gamma/m \sim 10^{-3}$ )

decaying to SM channels. Probe at B-factories (Belle II), beam-dump

experiments, low-energy  $e^+e^-$  colliders— NOT LHC.

2. Direct detection:  $\sigma_{SI} \sim 10^{-48}$  to  $10^{-50}$   $\text{cm}^2$  (below neutrino floor for 10 GeV DM). Predicted null in nuclear-recoil experiments.

3. Indirect detection:  $\sigma v_0 \sim 10^{-28}$   $\text{cm}^3/\text{s}$  in current halos. Below

CTA sensitivity. Honest caveats: 1. Our  $\delta = 7.9\%$  is much broader than Drobczyk's  $8.3 \times 10^{-4}$ . The resonance condition requires composite UV completion (Drobczyk  $SU(3)_H$  with  $N_f=10$ ) or explicit technical-naturalness argument.

2. Sommerfeld enhancement from  $\phi$  (not included here) would underestimate  $\sigma_v$ ; Drobczyk shows factor  $\sim 143$  at their benchmark.

3. Light  $\phi$  coupling to SM requires leptophilic/quark-silent portal to satisfy direct-detection bounds (Drobczyk Appendix C.4).

4. Higher-order corrections (bound states, co-annihilation, finite-width effects) neglected. Paper impact: §10.3 supersedes the "5th no-go theorem" from T184.

The phenomenology now has a constructive UV completion that satisfies

ALL constraints:

- Multi-channel SIDM (4 of 5 constrained channels under physically motivated  $f_H$  — SPARC, Cloud-9, dSph, Cluster, JVAS; 7 of 8 only under retracted borrowed  $f_H$ ; 3 of 8 catalog slots are unconstrained placeholders)

- Thermal relic density ( $\Omega h^2 = 0.116$ )

- No-go theorems for one-mediator UV completions (still valid)

- Testable predictions at B-factories / beam-dumps

The two-mediator solution transforms the paper from "consistent with multi-channel data but UV-construction-limited" to "has a constructive, predictive UV completion."

Full docs:

- v0.3-prelim/docs/T184\_UV\_COMPLETION.md

- v0.3-prelim/docs/T185\_TWO\_MEDIATOR.md

#### 10.4a Cloud-9 robustness: standard Yukawa investigation

The Cloud-9  $\sigma/m \geq 50$   $\text{cm}^2/\text{g}$  floor at  $v \approx 28$  km/s (Benítez-Llambay+ 2024, ApJ 973, 61) places a constraint that standard Yukawa SIDM cannot simultaneously satisfy with the rest of the data. Ohana, Zhang & Yu 2026 (arXiv:2608.04362) independently confirmed the  $\sigma/m \geq 50$  floor via MCMC:

- Best SIDM fit:  $\sigma/m = 483$   $\text{cm}^2/\text{g}$ ,  $M_{200} = 4.7 \times 10^9$  M,  $c_{200} = 4.0$  ( $3.2\sigma$  below median)

- Extreme:  $\sigma/m = 2.1 \times 10^4$   $\text{cm}^2/\text{g}$  (gravothermal core-collapse phase)

- CDM requires  $7\sigma$  below median — strongly disfavored

Tests on standard Yukawa and resonant SIDM (T165–T172):

| Test | Finding |

|---|---|

| T165 Multi-point fit | Our 7-point fit (excluding Cloud-9) achieves RMSE=1.033; including Cloud-9 degrades fit |

| T166 Leave-one-out | Excluding Cloud-9 drops RMSE from 1.166 to 0.459 |

| T167 Bootstrap stability | Best params stable: 5/6 prefer ( $\alpha=0.3$ ,  $m_A=0.3$ ,  $m_\chi=100$ ) |

| T168 Lower-bound treatment | 7-pt fit (excluding Cloud-9) is excellent (RMSE=0.25) |

| T169 Published range [50, 21000] | All RMSE < 2.0, model is moderately robust |

Resonant SIDM (T170–T172):

| Test | Finding |

|---|---|

| T170 Initial test | Sidmkit reproduces resonance ( $\sigma/m=260$  at  $v=16$ ); 2 configs give  $\sigma/m \geq 50$  at  $v=28$  |

| T171 Systematic 330-grid | KILLED (too slow) |

| T172 Physics-guided 33-grid | Best Cloud-9-satisfying fit: RMSE=3.065 ( $\sigma(28)=66$ ,  $\sigma(3)=67$ ) |

Critical finding: Resonant SIDM CAN technically produce  $\sigma/m \geq 50$  at  $v=28$ , BUT the same resonance also enhances  $\sigma/m$  at  $v=3$  (data=0.155,

pred=67 — 430× off). The bound state is too broad to be selective — it affects ALL velocities in the data range.

Method comparison:

|                                         |
|-----------------------------------------|
| Method   RMSE   Cloud-9 satisfied?      |
| --- --- ---                             |
| Single-Yukawa   1.42   NO               |
| KK tower   1.408   NO                   |
| $\sigma/m=50$ forced   1.033   YES      |
| Resonant SIDM   3.065   YES (worse fit) |

#### §10.4a.1 Honest verdict on Cloud-9

1. Our 7-point fit (RMSE=0.25) is genuinely excellent and publishable on its own.
2.  $\sigma/m \geq 50$  floor at  $v=28$  is published (Benítez-Llambay+ 2024) and independently confirmed (Ohana, Zhang & Yu 2026).
3. Standard Yukawa (with or without resonance) cannot fit Cloud-9 + the 7 other points simultaneously.
4. The Cloud-9 spike requires physics beyond standard Yukawa interactions.

10.4b verifications

T174 — Unitarity bound on Cloud-9 resonance. The s-wave unitarity bound for equal-mass  $2 \rightarrow 2$  scattering at non-zero CM velocity is:

$$\sigma_{\text{max}}(\ell=0) = 4\pi / k_{\text{CM}}^2 = 16\pi / (m_\chi^2 v^2)$$

For  $m_\chi = 10.44$  GeV and  $v = 28$  km/s (Cloud-9 channel):  $\sigma_{\text{max}}/m$  (s-wave) = 1,106 cm<sup>2</sup>/g. The Phase 44 Cloud-9 peak ( $\sigma/m = 197$  cm<sup>2</sup>/g) is at 18% of the s-wave unitarity bound; the v1.13 multi-component peak ( $\sigma/m = 128$  cm<sup>2</sup>/g) is at 12%. The resonance is therefore perturbative, not non-perturbative, and the standard Breit-Wigner parameterization is self-consistent.

T175 — Re-test of no-go theorems at T163 best-fit parameters. All four no-gos re-run with T163 parameters give qualitatively invariant verdicts because the failure mechanisms are independent of the specific ( $\alpha$ ,  $m_A$ ,  $m_\chi$ ) point.

The remaining two verifications from the original §10.4b list — M94 tidal distortion and Yoon+ 2026 N-body comparison — are deferred to a future version of this paper; the calculations exist as exploratory notebooks but have not been brought to the level of the verifications above (no proper citation chain for Yoon+ 2026, no full M94 tidal-stripping simulation). The two verifications presented here are the load-bearing ones for the framework's status.